

Software Language Engineering with ART

Adrian Johnstone

a.johnstone@rhul.ac.uk

February 18, 2025

Department of Computer Science
Egham, Surrey TW20 0EX, England

Contents

1	Introduction	1
1.1	The classical approach to translation	1
1.2	Early compilers and modern challenges	6
1.3	Specification styles for syntax	7
1.4	Specification styles for semantics	8
1.5	Ambiguity	9
1.6	Our approach – Ambiguity Retained Translation	11
2	Models of program execution	15
2.1	The fixed-code-and-program-counter interpretation	16
2.2	The problem with assignment	17
2.3	The problem with the program counter	17
2.4	The reduction interpretation	18
2.5	The problem with loops	19
2.6	A reduction evaluation of GCD in MiniGCD	19
2.7	Executable semantics and automation	20
2.8	Specifying reduction semantics: a preview	22
2.9	Specifying attribute-action semantics: a preview	23
3	Rewriting	25
3.1	Semantic equivalence of programs	25
3.2	Static and dynamic properties	26

3.3	Mechanised rewriting	27
3.4	String rewriting	28
3.5	Term rewriting	31
3.6	Internal syntax style	32
3.7	Terms	33
3.8	Rewriting in ART	34
4	Context Free Grammars and parsing	37
4.1	Outer and inner syntax	38
4.2	Chapter overview	40
4.3	CFG definitions and examples	41
4.4	Parsing	49
4.5	Derivation selection using choosers	50
4.6	GIFT annotations	50
4.7	GIFT applications	54
4.8	RDSOB: Implementing a parser toolchain	58
4.9	Recursive Descent with Singleton Ordered Backtrack parsing	59
4.10	Engineering a complete Java parser	62
4.11	Using built in matchers	65
4.12	Using attributes and inline semantics	69
4.13	Implementing inline semantics	72
5	Reduction semantics	75
5.1	The basic idea	75
5.2	Execution via substitution	76
5.3	Avoiding empty terms—the special value <code>_done</code>	79
5.4	Term variables are metavariables	80
5.5	Pretty-printed rules	81
5.6	Pattern matching of terms	81

5.7	Pattern substitution	83
5.8	Rules and rule schemas	84
5.9	The interpreting function F_{SOS}	87
5.10	Structural Operational Semantics and execution traces	89
5.11	An eSOS specification for the GCD language	99
6	Attribute interpreters	103
6.1	Attribute Grammars	104
6.2	Attribute Grammars	104
6.3	Semantic actions in ART	107
6.4	Syntax of attributes in ART	108
6.5	Accessing user written code from actions in ART generated parsers	109
6.6	A naïve model of attribute evaluation	110
6.7	The representation of attributes within ART generated parsers	111
6.8	The ART RD attribute evaluator	111
6.9	Higher order attributes	112
6.10	An attribute-action specification for MiniGCD	114
6.11	Exercises	114
7	Software language design and pragmatics	115
7.1	The semantic facets of programming languages	115
A	ART user manual	117
A.1	Downloading and running ART for the first time	117
A.2	IDE and graphical component installation	118
A.3	Command line interface	119
A.4	Integrated Development Environment	119
A.5	ART script language	119
A.6	Lexical structure	119

A.7 Abbreviations	119
A.8 Rewrite rules	121
A.9 Context free grammar rules	121
A.10 Choose rules	121
A.11 Directives	121
A.12 Lexical builtins	121
A.13 The ART value system	122
A.14 ART value plugins	122
B The Royal Holloway course	127
B.1 Learning outcomes	128
B.2 Assessment	128
B.3 Teaching week by week	131
B.4 Protocol for asking questions	131
B.5 Project work week by week	131
C Lab exercises	133
C.1 Solid modelling with OpenSCAD	134
C.2 Simple rewrites	147
C.3 Starting a language	159
C.4 Making a full language, and project demos	172
C.5 Context Free Grammar Rules and GIFT rewrites	173
C.6 SOBRD parsing and attribute evaluation	174
C.7 Control flow with delayed attributes	175

Chapter 1

Introduction

Software Language Engineering concerns the design and implementation of compilers, interpreters, source-to-source translators and other kinds of programming language processors.

All forms of engineering are a mixture of creative insight and disciplined implementation. For instance, the architect of a bridge relies on structural engineers who can take a high level design and perform detailed calculations on that structure to test whether it will withstand daily use.

Ideally, this would be true for software too: our creativity would be expressed only through sound and principled techniques; that is techniques that have been found to be safe and efficient using mathematical and other forms of analysis.

In practice we mostly write software in a *hopeful* way, and then use testing to try and find the gaps in our understanding. Unfortunately, programming languages are inherently difficult to test since they are designed to be flexible notations with very many combinations of interacting features and in any case, as Edsger Dijkstra famously noted, *Program testing can be used to show the presence of bugs, but never to show their absence.* [DDH72, p6].

Our goal is the construction of *maintainable* tools which have concise specifications that are amenable to automated checking, and which allow automatic generation of the tool itself. By working at a high level, we hope to reduce implementation errors, just as high level programming languages with static type checking can catch many of the errors that arise when programming at machine level.

We emphasise maintainability because language processors are typically complex with many internal dependencies, and are thus fragile in the face of attempts to extend or modify them. High level programming language specifications in widely understood notations would make it much easier to extend and modify those specifications so that other engineers could both maintain and reuse our work into the future.

1.1 The classical approach to translation

Most programming language processors are built around the notion of these five classical *phases*: Lex – Parse – Analyse – Rework – Perform which together check the input source text, and then perform the actions specified therein.

The source program text to be translated into actions is simply a string of characters. In the **Lex** phase, this string is partitioned into a sequence of substrings called *lexemes*, and the resulting sequence of lexemes is passed to the parser. Typically these lexemes comprise a single identifier, keyword or constant. In free-format languages, whitespace and comments are usually discarded by the lexer and do not appear in the lexeme sequence.

The purpose of the **Parse** phase is to (a) check that the lexemes appear in an order that is allowed by the rules of the language, and (b) to build a *derivation* which shows how the lexeme string can be constructed from the language rules. For programming languages these rules are usually specified by string rewrite rules which together form a *grammar*. Together, the **Lex** and **Parse** phases are often referred to as the *front end*.

The **Analyse** phase constructs an internal representation of the source program in a form that supports the later phases, and checks certain long range properties, such as whether the type declaration of a variable matches its subsequent usages. Typically a *symbol table* is also constructed which lists identifiers and their role in the source program.

The resulting representation is often called an *Abstract Syntax Tree* (AST) but there is no general agreement on what should be in such a tree and what supporting structures should be provided: indeed complex compilers such as the GNU suite often have a variety of different internal representations which are used to support different facets of the compilation process, and not all of these representations are tree-like. As a result, in our work we avoid the term AST and instead refer to internal representations as *internal syntax* as opposed to the original form of the program which we call the *external syntax*.

The **Rework** phase iteratively modifies the internal representation, usually in an attempt to improve performance. For instance constant expressions inside loops may be moved so that they are evaluated once before the loop is entered, rather than being recomputed on every loop iteration. To be correct, the rework must not change the meaning of a program; establishing that correctness is hard. This phase is traditionally called an *optimiser* but we avoid that term because in general the resulting code is rarely optimal: indeed different kinds of rework can cancel each other out and in extreme cases actually reduce performance.

The **Perform** phase traverses the final form of the internal representation and performs the specified actions. In a *compiler* the actions output a machine level program which can be subsequently executed on some processor architecture. In an *interpreter* the **Perform** phase directly executes the actions itself. Compiled code is usually faster than interpreted code, but a good compiler requires much more engineering effort than an interpreter. Some systems blur the line between the two by, perhaps, starting execution in an interpreted mode but then pausing to compile frequently-executed code to native machine language which can then execute at full speed. The **Perform** phase is often called the *back end* and we may refer to the combined **Analyse** and **Rework** phases as the *middle end*.

We should note that although these five independent phases are a very useful way to *think* about the various tasks a production-quality compiler must perform, in a real translator they may be intertwined, or even absent. Simple translators may not have a Rework phase, and some parsing techniques are effective when the lexemes are just the individual characters in the source program: such *character level parsers* do not have a Lex phase.

Classical front end techniques

The years 1960–75 represent a golden age of research into front end techniques: Donald Knuth noted that ‘*compiler research was certainly intensive, representing roughly one third of all computer science in the 1960s*’[KD14, p.46].

The goal of that research was to balance utility with efficiency: Alfred Aho characterised part of the work as *Searching for a class of grammars that was big enough to describe the syntactic constructs that you were interested in, yet restricted enough that you could construct efficient parsers from it.*’[KD14, p.42].

This work coalesced around two classical approaches: (i) limited bottom-up parsing, in particular the YACC parser generator and its descendants such as Bison, and (ii) limited top-down parsing implemented using recursive descent. Typical bottom up parser generators are implementations of the theory of shift-reduce parsing based on LR tables; top-down parsers embody the theory of predictive LL parsing. There are many alternative and hybrid approaches, and a vast research literature.

We call these ‘limited’ parsing techniques to highlight the constraints that they demand of the language designer, who must massage their language specification into forms that are acceptable to these non-general algorithms. Once that is achieved, both approaches offer linear processing times (that is the processing time is simply proportional to the length of the input) and both approaches are sufficiently frugal in their use of memory that they were practical on 1970s computers with storage limited to a few tens of thousands of bytes.

Our approach emphasises *general* parsing which allows the language designer much more freedom. We shall return to this topic in section 1.6.

Classical approaches to describing languages

The convergence of theoretical analysis and engineering practice in the classical Lex and Parse phases represents a major (possibly *the* major) achievement of the first thirty years of Computer Science. Fifty years later, the lack of an agreed way to concisely and precisely specify the actions of a programming language in the Analyse, Rework and Perform phases is a continuing concern.

How are programming languages defined in current practice? Most widely used

programming languages have a ‘standard’: a document that describes the effects that should be induced by the phrases of the programming language. For instance, here is an extract from the Pascal report

9.2.2.1. If statements.

The if statement specifies that the statement following the symbol then be executed only if the Boolean expression yields true.

IfStatement = “if” BooleanExpression “then” Statement

Here is an extract from one of the *Java Language Specification* documents:

14.9. The if Statement

The if statement allows conditional execution of a statement.

IfThenStatement:

if (Expression) Statement

The Expression must have type boolean or Boolean, or a compile-time error occurs.

14.9.1. The if-then Statement

An if-then statement is executed by first evaluating the Expression.

If evaluation of the Expression completes abruptly for some reason, the if-then statement completes abruptly for the same reason.

Otherwise, execution continues by making a choice based on the resulting value:

(a) If the value is true, then the contained Statement is executed; the if-then statement completes normally if and only if execution of the Statement completes normally.

(b) If the value is false, no further action is taken and the if-then statement completes normally.

And here is a corresponding extract from one of the draft ANSI-C standards:

6.8.4 Selection statements

Syntax

selection-statement:

if (expression) statement

Semantics

A selection statement selects among a set of statements depending on the value of a controlling expression.

A selection statement is a block whose scope is a strict subset of the scope of its enclosing block. Each associated substatement is also a block whose scope is a strict subset of the scope of the selection statement.

6.8.4.1 The if statement

Constraints

The controlling expression of an if statement shall have scalar type.

Semantics

The substatement is executed if the expression compares unequal to 0.

All three extracts describe the same language feature: the simple conditional

statement.

If we write in C or Java the program fragment

```
z = 0; if (x == y) z = 3;
```

then we expect that after execution the variable `z` will hold the value 3 only if the variables `x` and `y` hold the same value. We get the same effect in Pascal by writing

```
z := 0; if x = y then z := 3;
```

From these examples we can deduce the following.

- ◊ All three languages have conditional statements that ‘do the same thing’ in some sense;
- ◊ The textual form of the program fragments for Java and C are the same, but the Pascal fragment has a different form, even though the effect is the same for all three languages.

In programming languages, the meaning of a fragment is best thought of as its effect on the *state* of the computer, where the state is the set of values maintained by a program, which could simply be the contents of the computer’s memory along with any changes to the computer’s input and output devices.

state

The written form of a programming language fragment is called its *syntax* and the effect on the system’s state its *semantics*. The meaning of a program is the accumulated semantics of its phrases; the semantics of a programming language is the accumulated semantics of every phrase that could ever be written in that language.

syntax

semantics

Informal and formal semantics

The extracts above from Pascal, C and Java language standards are all trying to explain the syntax and semantics of a conditional statement. In each case the semantics is described in careful English prose (what we might call ‘legalistic’ English) but with differing levels of detail. Nearly all programming language standards adopt this approach, but that is problematic because it is hard to check that a prose specification is complete (that is, there aren’t any special cases that have been left undefined) and consistent (that is, that there aren’t any conflicting statements). It is even harder to check that a programming language processor, such as the Java compiler, correctly implements all aspects of the standard. Since so much of modern life is mediated by software written in programming languages, this vagueness in the underlying specification of languages and their implementations is worrying.

In an ideal world, we would have a commonly understood concise notation for describing the semantics of a language fragment with which we could construct arguments for the completeness and consistency of a programming language standard, and which we could use to check the correctness of compilers, interpreters and other language processors perhaps using a computer itself to do the checking. A semantics described this way would be called a *formal semantics*.

formal semantics

In fact such notations do exist, but they have not been widely adopted by programming language designers, for perhaps two reasons: firstly they are conventionally presented in a mathematical style that deters many software practitioners; and secondly complete language descriptions are very dense and can be quite long. Now, the second reason is not a very good one, because legalistic English prose descriptions of semantics are also very dense and long: the Java Language Specification for Java 22 runs to 876 pages.

The defining quality of a formal semantics is that it should in some sense be *mechanical*, that is: it should be amenable to implementation as a manual procedure that could be followed without insight or thought, or as a computer program. The legalistic English in the extracts above does not meet that criterion because English is itself open to interpretation. We do not want to read the semantics for a language feature and then have to argue about what the semantics itself means – that is simply to move the problem on one level!

1.2 Early compilers and modern challenges

single-pass compiler

For many languages, classical parser generators can automatically produce front ends. Most of these tools incorporate some facility for triggering *ad hoc* side effects during parsing which allows simple translators to be constructed in a style called *Syntax Directed Translation*. For instance, many early compilers for the Pascal programming language were little more than a `Parse` phase that extended a symbol table in response to each declaration, and emitted machine-level code in response to expressions and control flow constructs. A particular advantage of this approach is that the source program is read only once, and all of the translation work is performed during that read so we do not need to hold the whole program in memory and can work line-by-line. However, this *single-pass compiler* approach constrains the kind of languages that can be translated because at the time that expressions are processed, we must know the type of the operands so as to decide whether to emit, say, a floating point or an integer addition. As a result, languages such as Pascal and C originally required the types of all identifiers to be declared before use. In addition, in a single pass compiler we cannot perform any long range analyses, or realistically re-order the code to allow more efficient execution.

In multi-pass compilers the details of the `Analyse`, `Rework` and `Effect` phases vary widely between systems, and are often poorly documented, so it can be hard to establish the completeness, consistency and correctness of real compilers. Compiler errors are not rare, at least in the early stages of a language's development: as the user base for a language grows, confidence in the asso-

ciated translators naturally increases because the users are effectively testers. Language processors with few users should perhaps be treated with caution.

Just as software applications have a life-cycle, typically starting small and acquiring extra features over time, languages themselves display a life cycle, and modern versions of languages such as C and Java provide features that present significant challenges to all phases. The C language began as a small systems-oriented language but has had significant new functionality grafted on to it over the years, especially the extension to object orientation with C++.

As new features were added with necessary extensions to the syntax, the task of producing an appropriate grammar that was admissible by classical parsing algorithms became so difficult that around the turn of the twentieth century, the GNU C++ compiler abandoned parser generators in favour of manually crafted parsers. Without a generator, we must carefully analyse the hand written parser code and reassure ourselves that the front end is complete, consistent and correct.

The challenges multiply as we move through the phases. Modern versions of Java support limited *type inference* (that is the ability to deduce the type of an identifier at compile time without the programmer needing to explicitly specify it) and *pattern matching* (the use of patterns involving variables and constants in selection statements). Designing these sorts of features in a way that provides both backwards compatibility with earlier versions of the language *and* supports the sometimes quirky special cases that arise is very difficult indeed, and can lead to curious and unexpected behaviours.

1.3 Specification styles for syntax

Although formal semantics has not achieved much traction with language designers, there is almost complete agreement that syntax should be specified using a particular style of rewrite rule called a context free string-rewrite rule and that is evident in the extracts above; they all give a context free rewrite rule for the syntax, although the notation used for the rule varies slightly.

For Pascal we have

```
IfStatement = "if" BooleanExpression "then" Statement
```

The doubly-quoted symbols are Pascal keywords. The other symbols are placeholders for language fragments. For instance, a BooleanExpression could simply be the constant `true` or an expression like `x > y`.

In the C and Java standards, the placeholder symbols are written in an italic font with the literal keywords in an upright font. Clearly the placeholders could represent arbitrarily long pieces of program but when focussing on the syntax and semantics of the conditional statement, we don't want to have to specify their exact form. This is an example of *abstraction*, that is the hiding of unnecessary detail so that we can focus on the matter in hand.

abstraction

The purpose of a rule is to give a template for one feature of the language: the Pascal rule tells us that a conditional statement starts with the keyword `if` which must be followed by an expression yielding a boolean, then the keyword `then` followed by an arbitrary statement. Somewhere else in the grammar we expect to see definitions for the placeholders `BooleanExpression` and `Statement`. The C and Java versions tell us that a conditional statement starts with the keyword `if` which must be followed by an expression yielding boolean that *must* be enclosed in parentheses – in C and Java the keyword `if` is always followed by an open parenthesis.

A complete set of context free rules with no missing definitions and a nominated *Context Free Grammar* *start rule* is called a *Context Free Grammar* (CFG).

1.4 Specification styles for semantics

We do not use English to write programs because the meaning of English phrases is too fuzzy. Ambiguity, allusion, and occasional precision are exactly what poets need, but our topic is not poetry: we are trying to build reliable computer systems. Hence, our programs are written in *formal languages* which we hope have a well defined syntax and semantics resulting in the same behaviour on all implementations. Nevertheless, current practice is to describe the semantics of the programming language itself in English.

Clearly we *ought* to be writing the specifications for programming languages formally too so that we can use software tools to help us show that our syntax and semantics are complete, consistent and correct, and to generate the phases of our translators, just as compilers check our programs and generate executable programs. As we've seen with GNU C++, even in the front end implementers have retreated from that ideal over time due to (in that case) the weakness of classical parser generator tools.

It seems that non-trivial languages always have ‘dark corners’ where the interaction of useful features can cause surprising effects. Java’s design benefited greatly from previous generations of programming languages, but the *Java Puzzlers* book [BG05] show a wide variety of small programs with surprising effects – and that book was published in 2005 at the time of Java 5. The language has been significantly extended since then and further quirks will have resulted. Of course, these are confusions that arise at the level of the *user* of a language: the understanding of the language by *implementers* must (if their implementations are to be correct) subsume all of these details and more. A specification written in English is unlikely to answer all implementers’ questions.

The most significant attempt to define a general purpose language using formal rules (rather than English-language descriptions) is Standard ML. SML began as a scripting language for a theorem prover. A key design goal was to provide an external syntax that was comfortable for mathematicians via the use of type inference which (almost) eliminates the need for type declarations and pattern matching on *Algebraic Data Types* to directly support the kinds of case analysis

that naturally arise when writing down proofs. The formal definition of the language was published in 1997 [MHMT97] and *in principle* allows automatic construction of an interpreter from the rules given in the book. We are going to use a related specification style, coupled to new algorithms which remove the weaknesses of classical front end tools, and interpreters which can directly execute formal semantic specifications.

1.5 Ambiguity

An ambiguous statement is one that can be interpreted in more than one way.

Ambiguity is a fertile source of jokes:

How do you make a sausage roll? Release the sausage at the top of a ramp.

How do you make a professor fast? Take their lunch away.

And so on.

Allegedly a medicine was once advertised with this slogan: *Try our headache cure: you won't get better.* That is probably too obvious to be true, but from the mid-1950s an aspirin-based analgesic really was promoted with the line *Nothing works faster than Anadin*; one interpretation being that taking nothing would offer faster relief than using the product.

We do not want ambiguity in our programming languages since that would suggest that different implementations might make different interpretations, and so a program might behave differently on different machines (or even different runs on the same machine). However, Java and other languages are littered with phrases that have multiple *possible* meanings. Here are three example questions: answers below.

1. In C and Java, does the entirely valid phrase $z = x---y$ mean the same as $z = (x--) - y$ or $z = x - (--y)$ or $z = x - (-(-y))$?
2. In everyday arithmetic (never mind programming languages) does $5 - 4 - 3$ evaluate to -2 or to 4 , that is, should we interpret the expression as $((5 - 4) - 3)$ or $(5 - (4 - 3))$?
3. In this Java program fragment $y = 6; \text{if } (x > 3) \text{ if } (x > 5) y = 1; \text{else } y = 0;$ what should the final value of y be when x is 4? If we think that the fragment has the same meaning as

```

1 y = 6;
2 if (x > 3) {
3   if (x > 5)
4     y = 1;
5 }
6 else

```

7 | $y = 0;$

then the final value of y is 0.

If we use this interpretation

```

1 y = 6;
2 if (x > 3) {
3   if (x > 5)
4     y = 1;
5   else
6     y = 0;
7 }
```

then the final value of y is 6. The difference between the two interpretations is essentially whether the **else** clause belongs to the outer or the inner **if** statement.

In Section ?? you will find a Java listing that illustrates all of the different interpretations which you can compile and run to see the differing outcomes.

disambiguation rule

For these three rather simple examples, it turns out that there is an agreed *disambiguation rule* (but beware: deciding how to resolve ambiguities in general can be very challenging).

1. The lex phase partitions the input using a so-called *longest match* strategy so the input is broken down into $z = (x--)$ – y and the resulting values are $x=3$ $y=6$ $z=-2$.
2. We are taught in elementary school that, by convention, subtraction binds more tightly to the left so the expression is interpreted as $((5 - 4) - 3)$ resulting in -2. Programming languages implement this convention.
3. The rule is that the **else** clause binds to the most recent **if** statement, so the phrase is interpreted as $y = 6; \text{if } (x > 3) \{ \text{if } (x > 5) y = 1; \text{else } y = 0; \}$ and the final value of y is unchanged by the **if** statements.

There are well established techniques that may be used to ensure that classical lexers and parsers always pick the agreed interpretation when handling these simple cases, but some ambiguities are harder to manage.

Consider the Java declaration `Set<Set<Integer>> setOfSets;` The intrinsic arguments are bracketed using `<...>` and so the nesting finishes with `>>`. But those two characters together also represent the right-shift operator in Java, so how does the lex phase know whether to partition `>>` as two closing bracket lexemes `>`, `>` or as a single operator `>>`? Once we have a phrase level analysis from the parser we can resolve this ambiguity, but in the classical pipeline the

lexer runs first, and classical lexers can only return a single sequence of lexemes. Users of classical tools have to adopt a variety of complex mechanisms to overcome these difficulties, and that makes the resulting translators hard to understand and hard to completeness, consistency and correctness. We shall show how to use *general lexer* and a *general multiparser* which allows all

We should stress that this kind of ambiguity does not result from using English to express the semantics: we have to be able to manage ambiguity in all translation stages even if they are fully formal.

1.6 Our approach – Ambiguity Retained Translation

A fundamental flaw in the classical front end pipeline is that each phase must commit to a *single* interpretation before moving on the next, but as we have seen with the `>>` ambiguity, the information needed to make that decision may not become available until later phases have done their work. Similar problems arise in the `Parse` phase when we may need type information for identifiers before it has been analysed.

What we need is a pipeline in which we can keep our options open, resolving ambiguities only when the necessary information becomes available. We call this strategy *Ambiguity Retained Translation* (ART) which is also the name of the translation tool we designed to illustrate the approach. You will find detailed documentation on ART in Appendix A along with installation instructions in Section A.1.

The ART front end In ART, the classical LR and LL style deterministic parsers are replaced by a *multiparser* and a *multilexer*. As the name suggests, multiparsers are capable of returning multiple derivations (interpretations) of a string, and in fact our MGLL algorithm will return *all* derivations, even when there are infinitely many of them [SJW23]. MGLL achieves this in worst case cubic time and cubic space; in practice the worst case bound is elusive and the algorithms only require linear time for typical current programming language phrases, degrading gracefully towards the cubic bound when processing difficult parts of the grammar.

multiparser
multilexer

This capability does not come for free though: the data structures required during parsing and lexing grow rapidly and require many megabytes for realistic applications, and the amount of processing required for each input character is also significantly greater than for classical algorithms. At the time that the classical approach was being developed, MGLL would have been entirely impractical as it required far more memory than would have been available and would have run very slowly. However, modern machines typically run around a thousand times faster and have a thousand times more memory than those 1970s computers and so we can use these more advanced algorithms to free the language designer from the constraints of the classical algorithms.

Semantics in ART There are many formalisms that have been developed for

capturing the semantics of programming languages. We could, for instance, establish a relationship between phrases of a language and well-understood mathematical objects such as sets and functions, and then use traditional mathematical proof techniques to investigate program properties. That approach requires solid mathematical training to be fruitful.

In ART we do something much simpler: take the derivation tree that emerges from the front end and progressively transform it under the control of *term rewriting* rules until the program has all been rewritten away, accumulating the program's side effects as we go. The particular form of rewrite rules that we use is called a *Structural Operational Semantics (SOS)*; this approach was introduced by Gordon Plotkin in 1980 [Plo04].

ART also offers a more concise specification style called an *attribute grammar* from which can automatically generate rewrite rules. That allows us to 'explain' attribute grammars as a constrained form of our rewrite rules. Importantly, attribute grammars may be interpreted rather efficiently, and so a rewrite system limited to attribute grammar style rules will run faster with an attribute-evaluator style interpreter than with a full rewrite interpreter. It turns out that placing extra constraints in the style of attribute grammar will allow even faster interpretation.

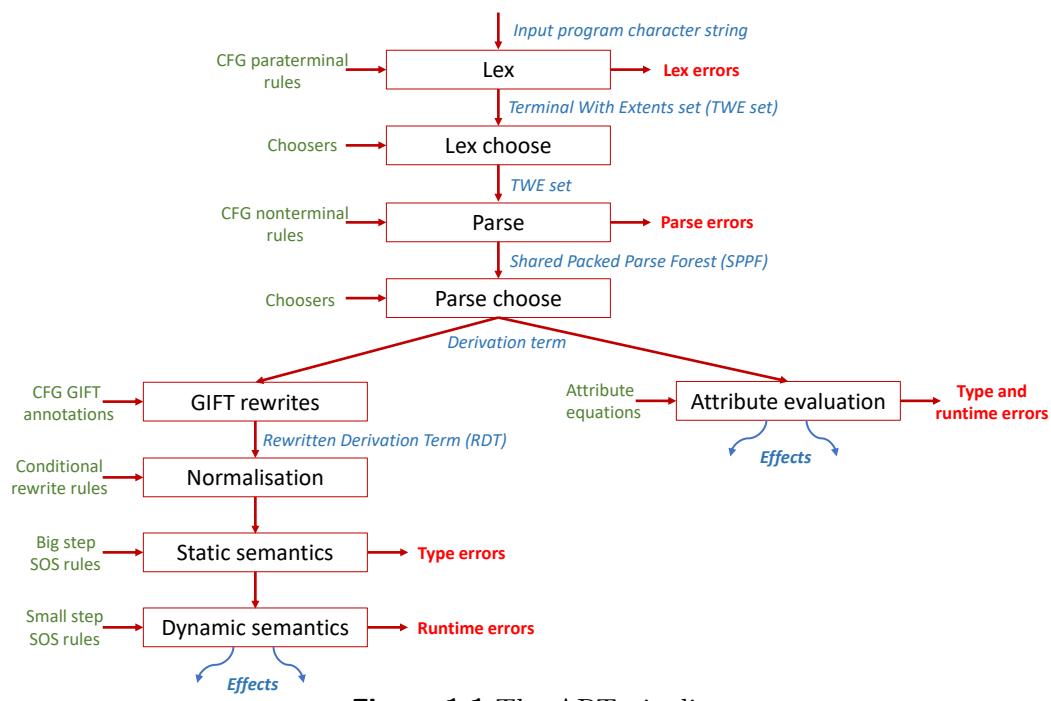
This refinement process rewrite rules to from illustrates a general principle: we want to proceed from formal specifications to implementations using, as automation as far as possible, because whenever we intervene manually we risk introducing errors.

The ART pipeline An ART specification is a collection of three different kinds of rules: **Context Free Grammar** rules specifying string rewrites which define the syntax; **Chooser** rules used to selectively discard interpretations; and **Conditional Term Rewrite** rules specifying tree rewrites which define the semantics

As well as these declarative rules, there may be *directives* which specify, for instance, which parts of the grammar is to be processed by the lex phase and which give test cases.

Conceptually ART follows the classical pipeline phases Lex – Parse – Analyse – Rework – Perform except that ambiguity management phases are inserted after Lex and Parse, and the other phases are merged together into sets of rewrite rules which may work on the internal form of the program in an intertwined way. Figure 1.1 shows this internal structure.

The ART front end (comprising the Lex – Lex choose – Parse – Parse choose phases) delivers a derivation term, unless lexical or parser syntax errors are detected.



Chapter 2

Models of program execution

We shall use as a running example a tiny language which illustrates the core procedural concepts of variables, assignment, arithmetic and control flow in the form of conditionals and loops.

The inspiration for our language is Euclid's integer Greatest Common Divisor algorithm, described in the second proposition of *Elements VII* some 2,300 years ago. It is worth looking up the original description which is written in quite verbose prose. Here is a version written in Java.

```
1 public class GCD {
2     public static void main(String[] args) {
3         int a = 6;
4         int b = 9;
5
6         while (a != b) {
7             if (a > b)
8                 a = a - b;
9             else
10                b = b - a;
11            }
12        int gcd = a;
13    }
14}
```

Java programs need quite a lot of setting up and anyway this is a course on language design, so let us construct our own, more compact programming notation to express the same program.

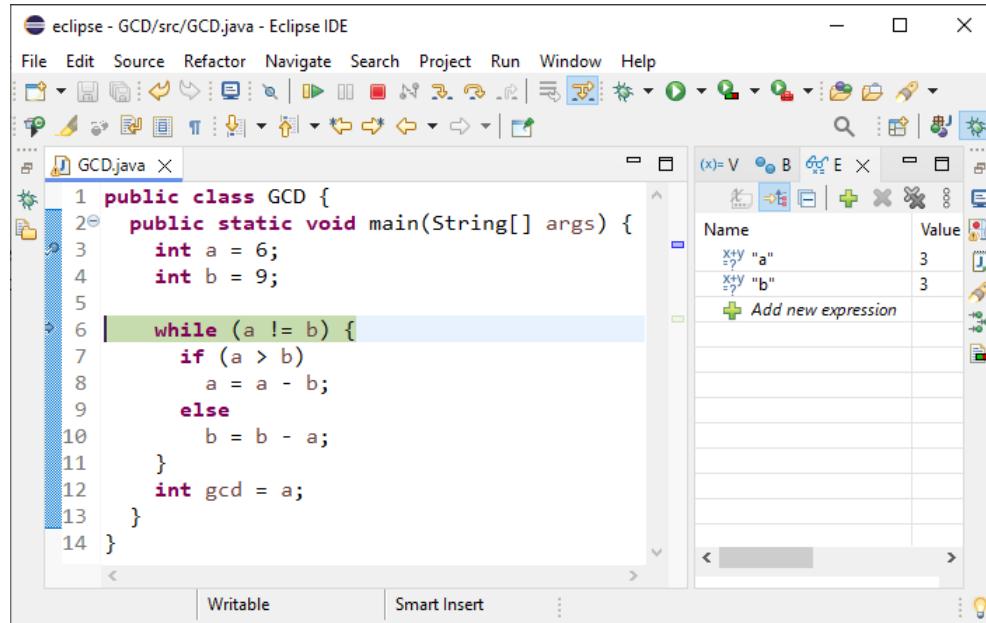
```
1 a := 6;
2 b := 9;
3
4 while a != b {
5     if a > b
6         a := a - b;
7     else
8         b := b - a;
9    }
10 gcd := a;
```

We shall call this notation the MiniGCD language. Note that assignment to a variable is denoted by `:=`, not by `=` as it would be in FORTRAN, C and Java. As in C and Java, statements are terminated with (not separated by) a semi-colon and can be grouped within braces. Variable names are not pre-declared and are assumed to be of type integer. MiniGCD only contains the features used here: it does not even provide addition (though we will extend it later).

2.1 The fixed-code-and-program-counter interpretation

In almost all modern computing devices most programs are lists of instructions that reside in the memory or *store*. The instructions for a particular program do not change as it is being executed. A special register called the *Program Counter* (PC) which points to the next piece of code to be executed, and is usually simply incremented as each instruction is executed which induces sequential execution of the instructions. At a branch point we may test a condition and update the PC with a new values depending on that outcome; that causes the processor to start execution at some new location.

The sequence of values displayed by the program counter during a program's execution records the *control flow* for this particular input. The easiest way to visualise the control flow for a program is to load it into a development environment such as Eclipse and then run it under the controller of the debugger, which can execute it one line at a time. Here is a screen shot from Eclipse showing our Java GCD which is currently at line 6 on the final iteration.



The screenshot shows the Eclipse IDE interface with the following components:

- Top Bar:** File, Edit, Source, Refactor, Navigate, Search, Project, Run, Window, Help.
- Toolbar:** Various icons for file operations, search, and navigation.
- Left Panel:** Shows the Java code for GCD.java. Line 6 is highlighted with a blue background, indicating the current execution point. The code is as follows:

```

1 public class GCD {
2     public static void main(String[] args) {
3         int a = 6;
4         int b = 9;
5
6         while (a != b) {
7             if (a > b)
8                 a = a - b;
9             else
10                b = b - a;
11        }
12        int gcd = a;
13    }
14 }
```
- Right Panel:** Shows the state of the store. It has two columns: Name and Value. There are two entries: "a" and "b", both with a value of 3.

Name	Value
X+Y "a"	3
X+Y "b"	3
Add new expression	

We can see the program's Java code on the left and the state of the store with variables `a` and `b`, both presently mapped to the value 3. The program counter value is represented by the small arrow in the margin at line 6.

2.2 The problem with assignment

The *substitution model* for variables that is used when thinking mathematically is much simpler to reason about than the *assignment model* for variables used in procedural programming languages. In mathematics, if I say $x = 3$ then, whilst that version of x remains in scope, x and 3 will be synonyms and so anywhere that x appears subsequently in that scope I could cross it out and write 3. In procedural programming languages like Java, if I write $x = 3$ I may subsequently write $x = 4$ in the same scope region, and so the relationship between x and its value depends on the most recent assignment to x according to the execution history for a particular input. Hence assignment in languages like Java is fundamentally different to mathematical equality (and that is why programming languages use different symbols to denote assignment and equality).

A physical store is a fixed set of cells, each with a fixed address but containing a value which may be changed. A useful mathematical model of a store is as a set of *bindings* where each binding associates an identifier with a value.

Evaluating a *declaration* in the program term has the effect of creating a new store S' from S which has all of the bindings in S and the new binding required by the declaration. Assigning a new value to a variable has the effect of changing the mapping of one variable in the store, and using a variable in an expression requires us to look up the value mapped to the variable's identifier.

2.3 The problem with the program counter

Our hardware works with (mostly) static code and a program counter but that does not mean that a formal model of program execution must take the same view. Just as aviation pioneers had to learn that wing-flapping was not a useful way to get humans airborne (propellers, jet engines and aerofoils being a better engineering proposition) the pioneers of formal approaches to programming language semantics had to find a way of dispensing with both assignment and the program counter. Why is this?

The substitution model is simple and easy to reason about but the assignment model has the great advantage of being efficient in that identifiers may be re-used within a scope rather than having to have fixed content throughout the runtime of a program and that saves both memory and allocation time. The use of assignment to variables presents a challenge to formal analyses of program semantics though it is manageable as we shall see. However it is *particularly* problematic that the program counter itself works by assignment because that obscures the control flow within a program, and that makes it difficult to decide whether two programs states are the same.

2.4 The reduction interpretation

There is an alternative way of thinking about program execution that does not require the use of a program counter. The trick is to think of the program code itself as something that can be progressively rewritten until all that we have left is a result, coupled to a set of bindings that record changes to variables.

Consider this program

```

1 x := 3;
2 y := 10+2+4;
```

The first thing the program does is assign **3** to variable **x**, that is create the store binding $x \mapsto 3$.

Now, there is a sense in which the program fragment **x := 3; y := 10+2+4;** coupled to the empty store $\{\}$ is equivalent to the program fragment **y := 10+2+4;** with the store $\{x \mapsto 3\}$ because they both lead to the same final result. We could start with either and get the same result.

It is helpful to think of the store as representing the computer's state, and the program fragment as representing 'that which is left to do', so we can represent the execution of a program as a sequence of pairs comprising a program that represents only what remains to be done and a store:

1. **x := 3; y := 10+2+4; {},**
2. **y := 10+2+4; {x \mapsto 3}**

Next we need to evaluate expression **10+2+4** before we can assign the result to **y**. In detail, the computer can only execute one arithmetic operator at a time so we must pick a sub-expression to evaluate first; let us choose to execute **10+2** and rewrite it to the result **12**.

3. **y := 12+4; {x \mapsto 3}**

Now we do the other arithmetic operation: **12+4** is rewritten to **16**.

4. **y := 16; {x \mapsto 3}**

Finally we can assign to **y** and set the program fragment to **_done** which is a special value indicating that there is no more computation required.

5. **_done, {x \mapsto 3, y \mapsto 16}**

Execution is now complete. Note that we could start in any of the five states above and end up with the same output.

<i>reduction trace</i>
<i>reduction step</i>

We call this kind of display of machine states the *reduction trace* for our program, and each line represents a *reduction step*—so called because usually the program fragment reduces in size at each step (though not always, as we shall

see in the next section). The steps match up rather well with the individual machine level instructions that would be executed by a real computer, and at every point we have a complete record of the state of the machine as well as being able to see what else we have to do.

2.5 The problem with loops

A reduction semantics for linear code and conditional code is straightforward, but we need to think carefully about loops. The approach we use here is to make use of a *program identity*, that is a program transformation that does not change the semantics of a program term, but does change the syntax, and thus the reduction trace. If we have a loop of the form

```
while booleanExpression do statement;
```

then we can *always* transform it into

```
if booleanExpression { statement; while booleanExpression do statement; }
```

We have effectively unpacked the first iteration of the loop and are handling it directly with an **if** statement followed by a new copy of the **while** loop which will compute any further iterations. When we have completed all of the iterations we shall encounter a term like

```
if false { statement; while booleanExpression do statement; }
```

which can then be rewritten away. This device, then, allows us to treat **while** loops using only **if** statements.

2.6 A reduction evaluation of GCD in MiniGCD

Figure 2.1 on page 21 presents a reduction semantics trace for the GCD algorithm written in MiniGCD program shown on page 16.

There are a large number of steps in this trace, which make for intimidating reading, but bear in mind that each step (very roughly) corresponds to a machine operation such as fetching an operand or adding two numbers. Useful programs entail the execution of a *lot* of operations: some of the programs we run on modern processors take an appreciable amount of time to execute even though a 4GHz processor will, in just two seconds, execute one instruction for every person on the planet—a number well beyond our abilities to directly comprehend. This is just a roundabout way of saying that machine operations are fine grained, and we need an awful lot of them to do useful work. Any attempt to list all of the steps that are gone through by a non-trivial running program is going to generate a long list.

We shall use a slightly more compact form to display the steps. First, we write the entire program term on a single line: rather than the nicely laid out version shown on page 16, we say

```
a=6; b=9; while a!=b if a>b a:=a-b; else b:=b-a; gcd=a;
```

As before, our starting point is the whole program term coupled to an empty store:

```
a=6; b=9; while a!=b if a>b a:=a-b; else b:=b-a; gcd=a;, {}
```

Each step of our trace involves identifying a part of the program term that we shall execute next, and then rewriting the program term to represent what is left to do of the original term. We call the subterm that is to be replaced a *reducible expression* or *redex* for short. In the trace below, we have highlighted the chosen redex in red at each stage. Sometimes there is a choice of redexes available: for instance when processing the GCD program initialisations of **a** and **b**, it does not matter which order we process them in. We have chosen to do **a** first.

redex

When reading the reduction trace below, the bold headings should simply be treated as comments: they are there to break up the reductions into related blocks as an aid to comprehension and have no part in the formal description of program execution. At each step look for the highlighted redex: the following step should contain a term which has all of the blue parts from its predecessor, and a replacement for the redex. The black part of each reduction step will record any changes arising from side effects of the reduction, which for this program are limited to creating or updating bindings but in general might also include changes to the input and output, the raising of exceptions, and other so-called *semantic entities*

semantic entities

normal forms

The execution terminates when we get to a term for which no further reductions are available, that is, a term that contains no redexes. We call such terms *normal forms*. In this case, the final term is empty, which naturally has no redexes.

Upon termination, the variable **gcd** is bound to 3 which is indeed the greatest common divisor of 6 and 9.

2.7 Executable semantics and automation

Manually constructing the execution trace shown in Figure 2.1 would be time consuming, although it does have the benefit of showing very clearly and concisely what each step of the computation does (as opposed to the legalise English commentaries that we showed at the start of this chapter).

The whole point of this way of thinking about programming languages is to allow *automation*. We need descriptions of programming languages that facilitate the mechanical construction of language processors. Ideally, we should like to be able to specify both the syntax and the semantics of a language like GCD in a few pages, and then have the computer run GCD programs for us so that we can test the specification and satisfy ourselves that it works the way we want it to. We call this prototyping style of language execution an *Executable*

Start of trace

```
a=6; b=9; while a!=b if a>b a:=a-b; else b:=b-a; gcd=a; {}
```

```
b=9; while a!=b if a>b a:=a-b; else b:=b-a; gcd=a; {a ↠ 6}
```

Rewrite using while p s → if p { s ; while p s }

```
while a!=b if a>b a:=a-b; else b:=b-a;gcd=a; {a ↠ 6, b ↠ 9}
```

Evaluate $a \neq b$ **with store** { $a \mapsto 6, b \mapsto 9$ }

```
if a!=b { if a>b a:=a-b; else b:=b-a; while a!=b if a>b a:=a-b; else b:=b-a; } gcd=a; {a ↠ 6, b ↠ 9}
```

```
if 6!=b{ if a>b a:=a-b; else b:=b-a; while a!=b if a>b a:=a-b; else b:=b-a; } gcd=a; {a ↠ 6, b ↠ 9}
```

```
if 6!=9{ if a>b a:=a-b; else b:=b-a; while a!=b if a>b a:=a-b; else b:=b-a; } gcd=a; {a ↠ 6, b ↠ 9}
```

```
if true { if a>b a:=a-b; else b:=b-a; while a!=b if a>b a:=a-b; else b:=b-a; }gcd=a; {a ↠ 6, b ↠ 9}
```

Evaluate $a > b$ **with store** { $a \mapsto 6, b \mapsto 9$ }

```
if a>b a:=a-b; else b:=b-a; while a!=b if a>b a:=a-b; else b:=b-a; gcd=a; {a ↠ 6, b ↠ 9}
```

```
if 6>b a:=a-b; else b:=b-a; while a!=b if a>b a:=a-b; else b:=b-a; gcd=a; {a ↠ 6, b ↠ 9}
```

```
if 6>9 a:=a-b; else b:=b-a; while a!=b if a>b a:=a-b; else b:=b-a; gcd=a; {a ↠ 6, b ↠ 9}
```

```
if false a:=a-b; else b:=b-a; while a!=b if a>b a:=a-b; else b:=b-a; gcd=a; {a ↠ 6, b ↠ 9}
```

Evaluate $b - a$ **with store** { $a \mapsto 6, b \mapsto 9$ }

```
b:=b-a; while a!=b if a>b a:=a-b; else b:=b-a; gcd=a; {a ↠ 6, b ↠ 9}
```

```
b:=b-6; while a!=b if a>b a:=a-b; else b:=b-a; gcd=a; {a ↠ 6, b ↠ 9}
```

```
b:=9-6; while a!=b if a>b a:=a-b; else b:=b-a; gcd=a; {a ↠ 6, b ↠ 9}
```

```
b:=3; while a!=b if a>b a:=a-b; else b:=b-a; gcd=a; {a ↠ 6, b ↠ 9}
```

Rewrite using while p s → if p { s ; while p s }

```
while a!=b if a>b a:=a-b; else b:=b-a;gcd=a; {a ↠ 6, b ↠ 3}
```

Evaluate $a \neq b$ **with store** { $a \mapsto 6, b \mapsto 3$ }

```
if a!=b { if a>b a:=a-b; else b:=b-a; while a!=b if a>b a:=a-b; else b:=b-a; } gcd=a; {a ↠ 6, b ↠ 3}
```

```
if 6!=b{ if a>b a:=a-b; else b:=b-a; while a!=b if a>b a:=a-b; else b:=b-a; } gcd=a; {a ↠ 6, b ↠ 3}
```

```
if 6!=3{ if a>b a:=a-b; else b:=b-a; while a!=b if a>b a:=a-b; else b:=b-a; } gcd=a; {a ↠ 6, b ↠ 3}
```

```
if true { if a>b a:=a-b; else b:=b-a; while a!=b if a>b a:=a-b; else b:=b-a; } gcd=a; {a ↠ 6, b ↠ 3}
```

Evaluate $a > b$ **with store** { $a \mapsto 6, b \mapsto 3$ }

```
if a>b a:=a-b; else b:=b-a; while a!=b if a>b a:=a-b; else b:=b-a; gcd=a; {a ↠ 6, b ↠ 3}
```

```
if 6>b a:=a-b; else b:=b-a; while a!=b if a>b a:=a-b; else b:=b-a; gcd=a; {a ↠ 6, b ↠ 3}
```

```
if 6>3 a:=a-b; else b:=b-a; while a!=b if a>b a:=a-b; else b:=b-a; gcd=a; {a ↠ 6, b ↠ 3}
```

```
if true a:=a-b; else b:=b-a; while a!=b if a>b a:=a-b; else b:=b-a; gcd=a; {a ↠ 6, b ↠ 3}
```

Evaluate $a - b$ **with store** { $a \mapsto 6, b \mapsto 3$ }

```
a:=a-b; while a!=b if a>b a:=a-b; else b:=b-a; gcd=a; {a ↠ 6, b ↠ 3}
```

```
a:=a-3; while a!=b if a>b a:=a-b; else b:=b-a; gcd=a; {a ↠ 6, b ↠ 3}
```

```
a:=6-3; while a!=b if a>b a:=a-b; else b:=b-a; gcd=a; {a ↠ 6, b ↠ 3}
```

```
a:=3; while a!=b if a>b a:=a-b; else b:=b-a; gcd=a; {a ↠ 6, b ↠ 3}
```

Rewrite using while p s → if p { s ; while p s }

```
while a!=b if a>b a:=a-b; else b:=b-a;gcd=a; {a ↠ 3, b ↠ 3}
```

Evaluate $a \neq b$ **with store** { $a \mapsto 3, b \mapsto 3$ }

```
if a!=b { if a>b a:=a-b; else b:=b-a; while a!=b if a>b a:=a-b; else b:=b-a; } gcd=a; {a ↠ 3, b ↠ 3}
```

```
if 3!=b{ if a>b a:=a-b; else b:=b-a; while a!=b if a>b a:=a-b; else b:=b-a; } gcd=a; {a ↠ 3, b ↠ 3}
```

```
if 3!=3{ if a>b a:=a-b; else b:=b-a; while a!=b if a>b a:=a-b; else b:=b-a; } gcd=a; {a ↠ 3, b ↠ 3}
```

```
if false { if a>b a:=a-b; else b:=b-a; while a!=b if a>b a:=a-b; else b:=b-a; } gcd=a; {a ↠ 3, b ↠ 3}
```

Assign result

```
gcd=a; {a ↠ 3, b ↠ 3}
```

```
,{a ↠ 3, b ↠ 3 gcd ↠ 3}
```

Figure 2.1 Reduction trace for the GCD algorithm with inputs 6,9

Executable semantics semantics to distinguish it from an efficient production-quality compiler.

In this testing phase, we would be prepared to accept quite poor performance because our test programs would be small. If we were building a general purpose language we would probably need to then move on to a more efficient style of implementation, though still guided by the specification which would be the ultimate arbiter of correctness. However, it turns out that on modern hardware, for many applications needing a small language, an executable semantics might be fast enough on its own. We shall revisit this topic in Chapter 7.

2.8 Specifying reduction semantics: a preview

As a preview—all of which will make sense by the time you have read all of the subsequent chapter—let us now look at an executable formal *reduction* semantics for the GCD language. It comes in three parts (i) a context free grammar that specifies the external, human friendly, syntax of the GCD language; (ii) a set of reduction rules and (iii) some test inputs.

Here are the context free rules, written in ART's notation. We shall explain these in detail in Chapter 4.

```

1 seq ::= statement^* | statement seq
2 statement ::= assign^* | while^* | if^*
3 assign ::= &ID ':=' expression ';'
4 while ::= 'while'^* expression 'do'^* statement
5 if ::= 'if'^* expression statement | 'if'^* expression statement 'else'^* statement
6 expression ::= rels^*
7 rels ::= adds^* | gt^* | ne^*
8   gt ::= adds '>' adds
9   ne ::= adds '!=>' adds
10  adds ::= operand^* | sub^* | add^*
11    add ::= adds '+' operand
12    sub ::= adds '-' operand
13  operand ::= _int32^* | deref^*
14  _int32 ::= &INTEGER
15  deref ::= &ID

```

Here are the reduction rules, which are explained in Chapter 5.

```

1 !configuration ->, _sig
2
3 --sequenceDone --- seq(_done, _C), _sig -> _C, _sig
4 --sequence _C1, _sig -> _C1P, _sigP --- seq(_C1, _C2), _sig -> seq(_C1P, _C2), _sigP
5
6 --ifTrue --- if(true, _C1, _C2), _sig -> _C1, _sig
7 --ifFalse --- if(false, _C1, _C2), _sig -> _C2, _sig
8 --ifResolve _E, _sig -> _EP, _sigP --- if(_E, _C1, _C2), _sig -> if(_EP, _C1, _C2), _sigP

```

```

9
10 -while --- while(_E, _C),_sig -> if(_E, seq(_C, while(_E,_C)), __done), _sig
11
12 -assign _n |> __int32(_) --- assign(_X, _n), _sig -> __done, __put(_sig, _X, _n)
13 -assignR _E, _sig -> _I, _sigP --- assign(_X, _E), _sig -> assign(_X, _I), _sigP
14
15 -gt _n1 |> __int32(_) _n2 |> __int32(_) --- gt(_n1, _n2),_sig -> __gt(_n1, _n2),_sig
16 -gtR _n |> __int32(_) _E2, _sig -> _I2,_sigP --- gt(_n, _E2),_sig -> gt(_n, _I2), _sigP
17 -gtL _E1, _sig -> _I1, _sigP --- gt(_E1, _E2),_sig -> gt(_I1, _E2), _sigP
18
19 -ne _n1 |> __int32(_) _n2 |> __int32(_) --- ne(_n1, _n2),_sig -> __ne(_n1, _n2),_sig
20 -neR _n |> __int32(_) _E2, _sig -> _I2,_sigP --- ne(_n, _E2) ,_sig -> ne(_n, _I2), _sigP
21 -ne _E1, _sig -> _I1, _sigP --- ne(_E1, _E2),_sig -> ne(_I1, _E2), _sigP
22
23 -sub _n1 |> __int32(_) _n2 |> __int32(_) --- sub(_n1, _n2),_sig -> __sub(_n1, _n2),_sig
24 -subR _n |> __int32(_) _E2, _sig -> _I2,_sigP --- sub(_n, _E2),_sig -> sub(_n, _I2), _sigP
25 -subL _E1,_sig -> _I1,_sigP --- sub(_E1, _E2),_sig -> sub(_I1, _E2), _sigP
26
27 -deref __get(_sig, _R) |> _Z --- deref(_R),_sig -> _Z, _sig

```

Finally, here are the tests: they start with the directive **!try** which is followed by the whole program to be executed, written as a string. (You can also give the name of a file for large input programs).

It turns out that the context free rules act so as to convert this string form of the program into a *term* which is a tree-like structure. It is possible, and actually recommended, to write the reduction rules before deciding on the external human-friendly syntax, and so ART allows you to execute a term directly, bypassing the parsing stage. The second **!try** here contains the term that the parser generates when running on the string in the first try, so these two generate the same execution trace.

```

1 !try "a := 6; b := 9; while a != b do if a > b then a := a - b; else b := b - a;" 
2
3 !try seq(assign(a, 6), seq(assign(b, 9), while(ne(deref(a), deref(b)), if(gt(deref(a), deref(b)),
4 assign(a, sub(deref(a), deref(b))), assign(b, sub(deref(b), deref(a))))))), __map

```

2.9 Specifying attribute-action semantics: a preview

A reduction semantics is ideal for prototyping: it separates out the syntax rules from the reduction rules, and the reduction rules are amenable to further formal analysis. However, there is a cost: the repeated rewriting of the tree is expensive, both in terms of space and computation time. There is an alternative approach, variants of which are widely used in production translators. The idea is to take the tree that the parser produces and associate data fields called attributes with each node, and also to associate actions with each tree node

which are allowed to look at nearby attributes, and compute new values. We then walk the tree executing actions without changing the tree itself. We call these kinds of specifications action-attribute systems.

One form of an attribute-action system embeds into the Context Free Grammar rules actions written as Java fragments. Here is such a specification: the actions are delimited by !! ... !! pairs. We shall examine this approach in Chapter 6.

```

1 !prelude !! import java.util.Map; import java.util.HashMap; !!
2 !support !! Map<String, Integer> variables = new HashMap<>(); !!
3 statements ::= statement statements
4 | statement !! System.out.println("Final variable map " + variables); !!
5
6 statement ::= 
7   ID ':='! e0 ';'! !! variables.put(ID1.v, e01.v); !! // assignment
8 | 'if' e0 statement!< 'else' statement!< // if statement
9   !! if (e01.v != 0) interpret(statement1); else interpret(statement2); !!
10 | 'while' e0!< 'do' statement!< // while statement
11   !! interpret(e01); while (e01.v != 0) { interpret(statement1); interpret(e01); } !!
12
13 e0 ::= e1 !! e0.v = e11.v; !!
14 | e1 '>' e1 !! e0.v = e11.v > e12.v ? 1 : 0; !! // Greater than
15 | e1 '!=' e1 !! e0.v = e11.v != e12.v ? 1 : 0; !! // Not equal to
16
17 e1 ::= e2 !! e1.v = e21.v; !!
18 | e1 '-' e2 !! e1.v = e11.v - e21.v; !! // Subtract
19
20 e2 ::= INTEGER !! e2.v = INTEGER1.v; !! // Integer literal
21 | ID !! e2.v = variables.get(ID1.v); !! // Variable access
22 | '('! e1 !! e2.v = e11.v; !! ')!' // Parenthesised expression
23
24 ID <v:String> ::= &ID !! ID.v = lexeme(); !!
25 STRING_SQ <v:String> ::= &STRING_SQ !! STRING_SQ.v = lexemeCore().translateEscapes(); !!
26 INTEGER ::= &INTEGER !! INTEGER.v = Integer.parseInt(lexeme()); !!

```

Our main goal is conciseness, and both reduction and attribute-action specifications achieve that. The reduction specification contains only 15 syntax rules, and 18 rewrite rules. The attribute-action specification contains only eight context free rules and 15 actions. Each specification gives us variables, expressions, assignment, selection and loops (albeit with only a few operators, one type and no procedural abstraction).

Now let us explain how to read these rules, and how to design with them.

Chapter 3

Rewriting

Sometimes things look different but mean the same thing. For instance the program fragment `3+4` evaluates to the same result as `4+3`. If we are only interested in the result of an expression, then we say they are *equal*, and we can write $3+4 = 4+3 = 7$.

If we are being very careful, then we would say that the expressions are equal *up to evaluation*. In some contexts, these expressions would not be thought of as equal. For instance the expression `3+4` comprises three characters, and the expression `7` only one, so if we are interested in how much storage we need in a computer to hold the string representation of the expression, then `7` is not equal to `3+4`.

An *equation* is two expressions separated by the *equality symbol* `=`. At a fundamental level, this tells us that the two expressions either side are interchangable because they evaluate to the same object, and that means that we can freely replace one by the other. It turns out that we can do a great deal of useful program translation just by using equations.

3.1 Semantic equivalence of programs

In programming languages we are used to the idea of ‘equivalent’ programs. For instance, this Java loop:

```
1 for (int i = 1; i < 10; i++) System.out.print(i + " ");
```

generates the same output as

```
1 int i = 1; while (i < 10) { System.out.print(i + " "); i++ }
```

If all we are interested in is output of a program, we might say that these two fragments are *equal* up to output, or just output-equal. More loosely, we often say that two programs are *semantically equivalent* if they produce the same effects. In this example, the iteration bounds are constant, and to get the same output effect we could just have written

```
1 System.out.println("1 2 3 4 5 6 7 8 9 ")
```

These three fragments are all semantically equivalent, but the third one will almost certainly run faster as it does not have the overhead of the loop counter and

only makes one call to `println()`. Usually, our notion of semantic equivalence does *not* include performance, but only the values computed by a program.

Establishing semantic equivalence of fragments may be very easy, or very hard indeed. For instance, there is a simple recipe which will convert **for** loops to the form in our second fragment:

$$\text{for}(I ; P ; U) S \rightarrow I ; \text{while}(P)\{ S ; U \}$$

In this recipe, I stands for any valid initialisation expression, P for a predicate (an expression yielding a boolean), U is an update expression and S stands for the statement which is controlled by **for** construct.

In this example, I is `int i = 1`, P is `i < 10`, U is `i++` and S is `System.out.print(i + " ")`

rewrite rule

rewrite schema

This recipe is our first example of a *rewrite rule*. Strictly speaking, this is a *rewrite schema* because the placeholders I, P, U, S are *variables* that can stand for any expression of arbitrary complexity and so the schema represents an infinite set of individual rewrite rules, each involving explicit expressions that do not contain placeholder variables.

Note how the structure of the Java **for** construct neatly brings together all four of the elements that we need to perform the transformation to a **while** construct.

If we reverse the direction of the arrow, the resulting rule is not so useful:

$$I ; \text{while}(P)\{ S ; U \} \rightarrow \text{for}(I ; P ; U) S$$

We could use this reverse-rule to undo our particular transformation from **for** to **while**, but for most **while** loops it would be quite hard to figure out where the four placeholder expressions are—for instance the initialisation expression might be a long way before the **while** keyword. The initialisation expression might even be inside an **if** statement, and so the start value for the loop might be data-dependent. To correctly rework general **while** statements to **for** would require a detailed analysis of the whole program.

refactoring

Now, there are many useful transformations that can be implemented: and they for the basis of the *refactoring* transformations that many Integrated Development Environments provide, as shown in Figure 3.1 for Eclipse.

The recipes for performing refactorings can be much more complex than the single rewrite rule above, and some aspects of programming languages (such as unrestricted **void*** pointers in ANSI-C) make it very hard to be sure that refactorings are *safe*, by which we mean that a refactoring transformation has not changed the meaning of the program.

3.2 Static and dynamic properties

static property

dynamic property

A *static property* of a program is one that depends only on the program text. A *dynamic property* depends both on the program text and some particular input;

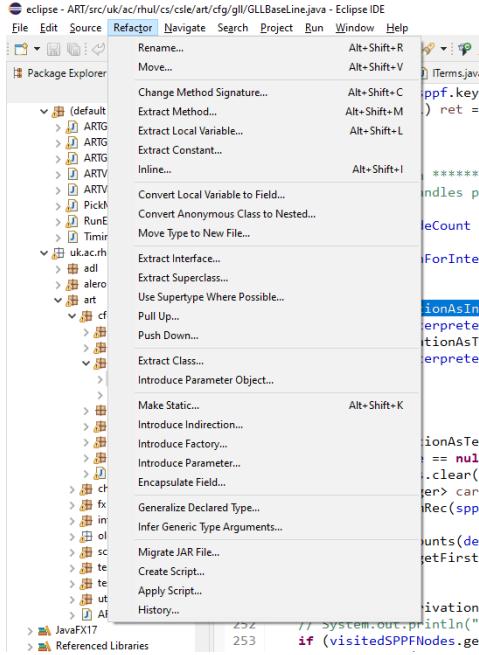


Figure 3.1 Refactoring in Eclipse

hence if we want to evaluate a dynamic property of a program we shall need to run the program, and that run will only tell us about the behaviour in response to a particular input. Dynamic properties demand that we use a testing-style strategy to check our assumptions.

Static properties are powerful, in the sense that they tell us things that will be true for *any* run of the program. The rewrite rule above exploits a statically visible relationship between the **for** and **while** constructs in Java to establish semantic equivalence between the first and second versions.

Establishing the equivalence to our third version (the print statement that simply outputs the result of the program) is *much* harder. Although the equivalence is strictly speaking, statically visible, we would need considerable reasoning to establish the equivalence and in practice we would probably just execute the fragments and compare the results: there are no useful structural relationships between the fragments that we can exploit to give a general rewrite rule.

3.3 Mechanised rewriting

We want to *mechanise* our rewriting mechanisms so that a computer can perform rewrites without human intervention. We want to write a program which will take a set of rewrite rules and some input, and then transform the input for us according to those rules. Such a program is called a *rewriter*. An IDE with refactorings (such as Eclipse) is a specialised rewriter that offers a fixed set

rewriter

of transformations (though there may be extension mechanisms).

ART is a rewriter that provides two kinds of rewriting: a specialised *string rewriter* which parses input strings—usually user programs written in some human-friendly external syntax—against a set of context-free string rewrite rules to produce a derivation tree, and a *term rewriter* which reworks that tree, perhaps reducing the whole tree to a value with some side effects (reduction semantics) or leaving the tree intact and rewriting data attributes associated with each node (attribute semantics).

The major challenge in implementing a rewriter is in locating a rule whose constants and variables line up with the structure that we are rewriting. The major challenge in using such a rewriter is to write rules which are well-behaved in ways that we shall discuss: primarily we want rules which yield the same outcome regardless of the order in which they are applied.

3.4 String rewriting

Anybody who has used a text editor is familiar with the principle of string rewriting which is really just cut and paste. Let us consider a simple subproblem, and think about it formally.

3.4.1 Rules for upper casing

Upper-casing a piece of text involves locating all instances of lower case letters in the range $[a]$ to $[z]$ and rewriting them into their corresponding upper case character in the range $[A]$ to $[Z]$. We could capture the uppercasing operation as 26 separate rules:

$$a \rightarrow A, \quad b \rightarrow B, \quad \dots, \quad z \rightarrow Z$$

This is a good start, but we want to change all of the characters in a string of arbitrary length, and this rewrite relation only modifies strings of length 1 containing a single lower case character.

We cannot reasonably be expected to write out a separate rewrite rule for every possible case: that would entail writing something of the form $x \rightarrow y$ for every string containing lower case alphabetic characters. Even if we limited ourselves to strings of length 5, there would be 26^5 such rewrite rules, which is 11,881,376 separate rewrite rules.

rule schema

A *rule schema* is a shorthand for a set of rewrite rules in which *variables* are used as placeholders for arbitrary bits of string. We shall use variables α and β to represent substrings of any length, including the empty string which has length zero.

This allows us to write out the more useful relation \rightarrow as follows

$$\alpha a \beta \rightarrow \alpha A \beta, \quad \alpha b \beta \rightarrow \alpha B \beta, \quad \dots, \quad \alpha z \beta \rightarrow \alpha Z \beta$$

This is just an abbreviated recipe for the *infinite* set of rewrite rules that handle the individual strings.

3.4.2 Confluence

Consider this rewrite scheme which is superficially similar to the upper-casing schema

$$\alpha a \beta \rightarrow \beta A \alpha, \quad \alpha b \beta \rightarrow \beta B \alpha, \quad \dots, \quad \alpha z \beta \rightarrow \beta Z \alpha$$

The modification here is that α and β are swapped over on the right hand sides, so the effect of the rewrite is to both to convert lower to upper case letters, and to swap over the ends of the string using the position of the lower case letter as a pivot.

This is a rather tricky rewrite system to reason about. Consider the string $AxyZ$. This has two lower case letters, and so two rewrites will be needed to achieve a normal form. Depending on the order in which we perform the rewrites, we get different results:

$$AxyZ \rightarrow yZX \rightarrow ZXAY$$

$$AxyZ \rightarrow ZYAx \rightarrow XZYA$$

We say that these rules are *non-confluent*. The term comes from the study of rivers: sometimes rivers split into multiple channels called distributaries might spread out into a delta, or might come back together into a single channel at a confluence.

We want programs to give same results whenever we run them with particular inputs, so we need our reduction semantics rules to be confluent, but that can be quite challenging. The core approach is to attach conditions to each rule such that at most one rule can be applied at each rewrite step. We might also attach a priority to each rule, and if more than one is applicable at some step, we choose the highest priority one. We only use this prioritisation as a last resort because ordered prioritised rules are harder to reason about, and this more error prone. On the other hand, ordering the rules can allow more concise specifications.

3.4.3 Avoiding implicit equality tests

This innocuous looking rewrite scheme has a hidden constraint:

$$\alpha a \alpha \rightarrow \alpha A \alpha$$

The left hand side of the rewrite $\alpha a \alpha$ will only match strings in which the letter a has identical strings on each side, such as $xyaxy$ and $zzazz$. The point is that by using α twice on the left hand side of a rewrite rule, we implicitly demand that both instances match the same substring.

In general this extra equality test can create a fixed overhead on real implementations of term rewriting. One way of avoiding the overhead is to require each variable in an open expression to be only used once, and that is the approach we take.

3.4.4 Avoiding contiguous variables

Here is another rewrite schema that is perfectly acceptable but which can cause real term rewriting implementations indigestion (and thus should perhaps be avoided):

$$\alpha\beta a \rightarrow \alpha A\beta$$

Now this is interesting because the left hand expression $\alpha\beta a$ will match any string ending with an a , and for strings of length greater than one there are *multiple ways to do the matching*. Consider the string xya . The concatenation of α and β can match the prefix xy in these three ways, leading to three rewrites (so the rewrite relation is not a function):

α	β	Rewrite
xy	ϵ	xyA
x	y	aAy
ϵ	xy	Axy

Here we have used *epsilon* ϵ to represent the empty string in the variable columns.

If were to restrict our variables to only matching a single character then this problem would go away, but that is quite a strict constraint. In ART, we distinguish between simple variables that match a single input element and ‘star’ variables that can match sequences of elements. We only allow a single star variable on the left hand side of a rule.

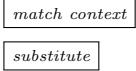
3.4.5 Matching, binding and substituting

<i>open expression</i>
<i>substitution</i>
<i>closed</i>
<i>pattern</i>

An expression with variables in it is called an *open expression*. A *substitution* is used to ‘fill in’ the variables. An expression whose variables have all been substituted (that is an expression with no variables in it) is called a *closed expression*. A *pattern* can be either a closed or open expression.

In rewriting, the expression to the left of the rewrite symbol is in a *match context* and the expression to the right is in a *substitute context*:

match-pattern → substitute-pattern



One way to think about this is that we are given some input string which we then attempt to fit the match-pattern to. Literals must match exactly, and pattern variables are then fitted into the remaining substrings to create a set of *binding*. We then create the rewritten results by taking the substitute-pattern and generating all of the strings that we can make by replacing variables with their bound values.

binding

3.5 Term rewriting

Programs often contain *expressions* such as

$$17/(4 + (x/2))$$

They have a well-defined syntax: for instance $4) * (x + 2)$ is not a syntactically well formed expression because of the orphaned opening parenthesis.

This particular way of writing expressions follows the style that we learn in school which makes use of *infix operators* like $+$ and $/$ to represent the operations of addition and division; they are called infix because are written in between the things they operate on. Expressions can nest and we understand that evaluation of an expression proceeds from the innermost bracket: to compute $17/(4 + (x/2))$ we first need to divide the value of x by 2, then add 4, and then divide the result into 17.

The choice of infix notation is just that: an arbitrary choice, and we could have decided to use a different syntax to specify the same sequence of operations, such as

```
divide(17, add(4, divide(x, 2)))
```

We call this form a *prefix* syntax because each operation is written in front of the (parenthesized) list of arguments that it is to operate on.

Although infix notation is familiar from everyday use it does not extend very comfortably to operations with more than two arguments. As a rare example: Java and C both provide the `p ? et : ef` notation for an expression in which predicate `p` is evaluated and then either expression `et` or expression `ef` is evaluated depending on whether the result of `p` was true or false.

In practice most programming languages provide infix notation for commonly understood operations such as addition, less-than and logical-AND, but use prefix notation for other operations. Usually we can define procedures which are then called using a prefix notation. So, for instance, in Java we might write

```
System.out.println(Math.max(x,y))
```

If you are interested in the design of external language syntax then there are some alternatives to this approach that you might like to investigate. For instance Scheme and other LISP-like language use an exclusively prefix style; the printer control language PostScript uses Reverse Polish Notation; the Smalltalk language effectively uses an infix notation to activate all methods; the C++ language allows the dyadic operator symbols like `+` to have their meanings extended to include new datatypes, and the Algol-68 language allowed completely new dyadic operator symbols to be defined.

3.6 Internal syntax style

As language *implementors* and specifiers, we are mostly concerned with *internal* syntax—that is, how to represent programs compactly within the computer. We would like a general notation which is quite regular and thus does not require us to switch between different styles of writing what are essentially similar things. We should like to be able to easily transform programs so that if we chose, we could rewrite an expression such as $3 + (5 - (10/2))$ into $3 + (5 - 5)$ or even 3.

The *prefix* style is both familiar from mathematics and programming, and easy to manipulate inside the computer so we shall use that style almost exclusively to describe entire programs, and not just expressions. For instance the program

```
x = 2;
while (x < 5) { y = y * y; x++;}
```

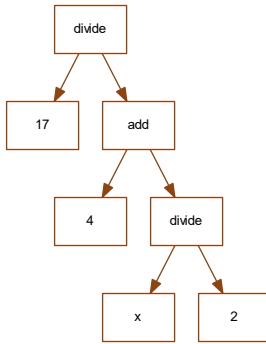
might be written

```
sequence(assign(x,2),
         while(greaterThan(x,5),
                sequence(assign(y, mul(y, y)),
                        assign(x, add(x, 1)))))
```

Here, the concatenation of two statements X and Y in Java is represented by `sequence(X, Y)` and an assignment such as `x = 2;` by `assign(x,2)`.

This notation has the great merit of uniformity: the wide variety of syntactic styles which are used in high level languages to improve program readability for humans is replaced by a single notation that requires us to firstly specify what we are going to do, say `add` and then give a comma-delimited parenthesised list of arguments that we are going to operate on.

The heavily nested parentheses can make this a rather hard-to-read notation although careful use of indentation is helpful. Sometimes, for small expressions at least, it can be helpful to use a tree diagram to see the expression. For instance $17/(4+(x/2))$, which we would write `divide(17,add(4,divide(x,2)))` can be drawn as



3.7 Terms

We call the components of a prefix expression *terms*. Syntactically, we can define terms using an inductive (recursive) set of rules like this.

1. A symbol such `[1]`, `[π]` or `[:=]` is a term.
2. A symbol followed by a parenthesized comma-delimited list of terms is a term.

Rule one defines terms made up of single symbols. Rule 2 is recursive, and this allows us to construct terms of arbitrary depth by building one upon another.

The *arity* of a term is the number of terms within its parentheses. Terms from rule 1 have no parentheses: they are arity-zero. Equivalently, the arity is the number of children a term symbol has in its tree representation. Rule 1 terms have no children and so are the leaves of a term tree.

Quite often, all instances of a symbol will have the same arity. For instance, addition is usually thought of as a binary (arity-two) operation, and an expression $3 + 4 + 5$ could be represented by the term `add(add(3, 4), 5)`. However, we could instead decide to have *variable* arity addition, in which case $4 + 4 + 5$ could be represented as `add(3,4,5)`.

3.7.1 Denoting term symbols

In term rewriting, we are very permissive about what constitutes a symbol: they are not just limited to the kinds of alphanumeric identifiers that we use in conventional programming languages.

Now, great care is needed when reasoning about and writing down terms. Rule 2 above make comma and parentheses special: how would we go about writing a symbol that contained parentheses or command? We call these special characters *metacharacters* because they are used in the denotation of terms.

If we do want a parenthesis or a comma within a symbol, we usually write it with a preceding back-slash `\() \,`. Of course, we have now added another meta-symbol, so if we want a back-slash in a symbol name we have to write it as `\\`.

3.8 Rewriting in ART

ART supports both limited string rewriting and general conditional term rewriting. String rewriting is limited in the sense that we only define rules for a parser which can convert external human-friendly syntax into prefix form that we need for our rewrite rules. ART does not actually do any string rewriting; instead the parser works out the ways in which a particular input string could be constructed from the rules. If we want to experiment with rewriting itself, we use term rewrite rules.

3.8.1 String rewriting in ART

String rewrite rules must be context free, and are specified using rules that look like this:

```

1 X ::= 'a' Y 'b'
2
3 Y ::= 'y'
4
5 Y ::= #

```

nonterminal

Each rule comprises: a *nonterminal* symbol (`X` and `Y` above) followed by the `::=` symbol (which you should read as ‘expands-to’) followed a sequence of one or more symbols, each of which may be (a) a nonterminal or (b) a stropped ‘`string`’ representing a *context-free terminal*.

context-free terminal

As a special case, we may also have rules whose right hand side is a single `#` character. This specifies that a nonterminal may be rewritten to nothing, i.e. it disappears when we rewrite it. (We usually use ϵ for the empty string when

writing in a mathematical style, but few keyboards provide that character so we use $\#$ instead.

We shall explore the use of these rules to specify syntax in detail in Chapter 4. As a preview, the way to read these rules is to start with a string comprising the LHS of the first rule: we shall call that the working string. Now select the leftmost nonterminal in the working string, look at the rule(s) that have on their left hand side and rewrite the working string, replacing the leftmost nonterminal with one of its right hand sides. When the working string has no nonterminals left in it, the rewriting process is complete.

If we apply that process to these rules we get these two derivations

$$X \Rightarrow aYb \Rightarrow ayb$$

and

$$X \Rightarrow aYb \Rightarrow ab$$

Note how in the second derivation, Y has been rewritten to $\#$ which does not appear in the actual string: $\#$ is a meta-symbol representing ‘nothing’.

As a notational convenience, we are allowed to group rules which have the same left hand side together, separating them with the $|$ character which you should read as ‘or’. Thus we can write these rules more compactly as

1	X ::= 'a' Y 'b'
2	
3	Y ::= 'y' #

3.8.2 Term rewriting in ART

Term rewrite rules look like this:

1	— label
2	optionalCondition1 optionalCondition2 ...
3	---
4	conclusion

The label is optional, and serves merely to give the rule a name which can be reported as a rewrite rule is being applied. If absent, the system will generate a unique name of the form Rn .

Conditions are also optional. They can be used to check, for instance, the values and also to invoke recursive calls to the term rewriter (a technique we shall cover in detail in Chapter 5) but plain unconditional rewrites will not have conditions.

The separating bar `---` is required.

`transition`

`relation symbol`

The conclusion is written as a *transition* which has a left hand term, a *relation symbol* (usually some sort of arrow) and a right hand term. At any given point the rewriter has a ‘current term’ which we call Θ (read as capital Theta). The rewriter looks through its rules to find one where the left hand side of the conclusion matches Θ . It then evaluates the conditions above the line, and if they are fulfilled then Θ is rewritten to the right hand side of the conclusion.

Importantly, these terms can have placeholder variables in them, which allow subterms to be carried over unchanged. Variable names always start with a *single* underscore.

The ART script language is free-format so we can space things out vertically, or for simple rules write them all on one line.

So, for example, if we represent the external syntax

`'while'expression 'do'statement`

using the internal syntax term

`while(expression, statement)`

then our **while** loop conversion trick can be implemented as

`1 | --- while(_E, _S) -> if(_E, seq(_S, while(_E,_S)), __done)`

Informally, when we match the rule to Θ , the variables `_E` and `_S` are bound to the expression and statement terms respectively. They are then substituted in to the right hand side term as part of the rewrite.

Chapter 4

Context Free Grammars and parsing

A context free grammar (CFG) is a string rewriting system which gives rules for composing sentences in a language. Each sentence is a string of characters or of words separated by whitespace. The complete collection of valid sentences specified by a grammar is called the *language* of that grammar.

Nearly all programming language standards documents use CFG rules to specify the syntactically valid forms of programs. In this chapter we look at the basic idea of a CFG; show how to specify common language idioms using CFG rules; and show how to add so-called GIFT annotations to CFG rules so as to generate terms suitable for the reduction-semantics rewriter we shall discuss in the next chapter.

There are three different ways that we might use a CFG:

1. As a language *generator* which takes a CFG Γ and lists valid sentences from $L(\Gamma)$ (the language of Γ) that could be used to test our language processors.
2. As a language *recogniser* which takes a CFG Γ and a string σ and returns boolean TRUE if σ is a sentence in $L(\Gamma)$ and FALSE otherwise.
3. As a language *parser* which takes a CFG Γ and a string σ and returns *derivations* of σ in Γ .

A derivation is a sequence of rewrite steps showing how σ can be constructed using the rules in Γ .

A *derivation tree* composes those rewrite steps into a tree or term,

There maybe more than one derivation tree for some σ in $L(\Gamma)$ in which case we say that Γ is *ambiguous*.

Γ may have cycles, in which case there will be an infinite set of derivation trees for some $\sigma \in L(\Gamma)$.

A *general CFG parser* is one which can

1. process any CFG;

2. return a representation of *all* of the derivation trees of any $\sigma \in L(\Gamma)$.

ART includes a general parser which uses the MGLL algorithm

4.1 Outer and inner syntax

Language processing tools usually maintain internal representations that are optimised for the task in hand. In the context of language based software tools, we often call this internal form the *abstract* or *inner syntax* of a language. The equivalent technical terms *intermediate form*, *internal representation* and *model* also appear in the literature. The term abstract syntax suggests some sort of relationship to the ‘real’ syntax of the user language, which is then called the *outer* or *concrete syntax*.

Transformation from outer to inner syntax typically involves (a) throwing away of redundant detail in the outer syntax (such as the parentheses that surround the predicate following an `if` keyword in Java, (b) ‘normalisation’ of forms such as converting `for` and `do...while` loops in Java to `while` loops so that only one style of looping has to be treated internally, and (c) rendering the result in the chosen internal syntax style.

Real systems are often layered, with the internal ‘abstract’ syntax of one layer becoming the starting point for an inner translation. As a result, we find the terms abstract and concrete unhelpful and prefer to use the terms *inner* and *outer* syntax for each layer; the inner syntax for a layer is the outer syntax for its enclosed layer. So for Java, for instance, the outermost syntax is the human-friendly form documented in the Java Language Specification. An implementation of `javac` (the Java compiler) will have its own internal form which is suitable for analysing and checking Java programs, and then some sort of back end which can be thought of as translating from that compiler-specific internal form to Java Virtual Machine Byte code which is then written into the `.class` files. This code then forms the external syntax for the Java Virtual Machine interpreter `java`.

4.1.1 Inner syntax for reduction semantics

The reduction interpreters that we shall describe in the next chapter use term rewriting to successively rewrite program terms into simpler forms, accumulating side effects during the process. For these kinds of language processors, we have an external syntax which is comfortable for human programmer such as

```

1 a:=15;
2 b:=9;
3 while a != b do
4   if a > b then
5     a:=a-b else

```

```

6|      b:=b-a;
7|
8| gcd := a

```

and a corresponding *inner* syntax which might be

```

1 seq(seq(seq(assign(a, 15), assign(b, 9)),
2   while(ne(deref(a), deref(b)),
3     if(gt(deref(a), deref(b)),
4       assign(a, sub(deref(a), deref(b))),
5       assign(b, sub(deref(b), deref(a))))),
6 assign(gcd, deref(a)))

```

This internal syntax is a nested expression over terms such as `seq(_X, _Y)` which sequences operations together, `assign(_N, _v)` which binds a value to a name and `sub(_v, _w)` which evaluates $v - w$.

4.1.2 Syntactic sugar, redundancy and syntactic ‘noise’ in human-friendly external syntax

An outer syntax designed for humans often contains elements which protect against common error patterns without adding any semantics. For instance, we could design a concise Java conditional expression which allowed us to write expressions such as

```

1| x = a > b ? y + 2 z * 3

```

The real Java conditional operator requires a colon between the two expressions

```

1| x = a > b ? y + 2 : z * 3

```

Why is this? Well, it allows the concrete syntax analyser to detect the situation where the user mistypes the second expression, omitting the variable

```

1| x = a > b ? y + 2 : * 3

```

which would be rejected, because there is no monadic * operator in Java. In our reduced syntax, this would be

```

1| x = a > b ? y + 2 * 3

```

which is a valid expression (though not the one the user intended) and so would be accepted by the parser.

This use of syntactic elements to catch common errors also explains why in Java and C the predicates of `if`, `switch` and `while` statements must be surrounded with parentheses, even though they carry no semantic information.

Another aspect of concrete syntax that is redundant in the derivation term is the use of parentheses to enforce operator execution order in expressions. We shall see how to write grammar productions that enforce associativity and priority rules for operators in the absence of parentheses, and in a fully parenthesized expression requires no such rules. In a tree, we use the depth of a node to encode its execution priority under the rules that the tree will be traversed top down, left to right with operators being executed in post-order. It is clear, then, that parentheses in the user expression may be omitted from the tree, and by the same argument other grouping elements such as braces around compound statements may be suppressed without losing fidelity.

4.1.3 The legacy of non-general parsing

A further source of redundancy in outer syntax as typically found in language standard documents such as that for ANSI-C is that they have been written so as to be admissible by classical deterministic parsing algorithms, and as such they can contain complicated BNF constructs which could be simplified for use with a general parser. This problem also affects language *exposition* for human readers. For instance, the first version of the Java Language Specification contains two grammars which we call the *pedagogic* and the *near-deterministic* grammars. In the main body of the document, individual language constructs are introduced with a grammar fragment that describes their syntax, accompanied by an informal English-language description of the semantics. The union of all these grammar fragments specifies the language, but unfortunately simply concatenating the pedagogic grammar fragments does not yield a grammar that is admissible by traditional parser generators. As a result, the JLS authors provide a second grammar which would be admissible, and describe its relationship with the pedagogic grammar so as to convince the reader that they generate the same language. With a more powerful parsing technology, it might have been directly use the pedagogic grammar, reducing the scope for errors.

4.2 Chapter overview

The main goal of this chapter is to show how to write CFGs that both describe the valid sentences in your programming language *and* have derivations that are in the right form to be fed directly into an executable reduction semantics based on rewrite rules in the style described Chapter 5.

By writing the CFG carefully we can use a CFG parser to automatically translate from external to internal syntax. In this we are aided by so-called GIFT annotations which specify local folding of nodes in the derivation tree. A complete reduction interpreter is then specified as a set of CFG rewrite rules to convert a program written in external syntax to an initial term written in internal syntax, and a set of rewrite rules that successively reduces that term to a value, accumulating side-effects along the way.

A secondary goal is to give an overview of the parsing process using the simplest algorithm I know of which can do at least *some* useful work: a non-general algorithm called RDSOB (known elsewhere as OSBRD!). The intention is to give insights into the process of parsing, without discussing details of the general MGLL algorithm. We shall then see how to enhance the basic parser with attributes and actions so as to make a rudimentary translation tool. This acts as a preview of the action-attribute technique described in detail in Chapter 6.

4.3 CFG definitions and examples

Context Free Grammars have a finite set of literals or *terminal symbols* T , a finite set of *nonterminals* N and a set of rewrite rules or *productions* written, for instance, as $X ::= Y \text{ 't' } Z$ where X, Y, Z are nonterminals and $'t'$ is a terminal. One of the nonterminals is called the *start symbol*.

More formally, we write:

a CFG Γ is a 4-tuple (N, T, S, P) with $N \cap T = \emptyset, S \in N, P \subseteq N \times (N \cup T)^*$

which we read as

A Context Free Grammar is a set of nonterminal symbols N , a set of terminal symbols T (N and T disjoint), a special start nonterminal symbol S and a set of productions the left hand side of which is a nonterminal, and the right hand side of which is a (possibly empty) sequence of terminals and nonterminals.

4.3.1 CFGs in ART

In ART, terminal symbols are always stropped '`like`' '`this`' and nonterminals never are, which ensures that we can immediately see if a symbol is a terminal or a nonterminal. As noted earlier, we prefer to have an explicit symbol for 'nothing', so we represent the empty righthand side of a rules with a hash symbol `#`. Unless otherwise specified, the first start symbol is the left hand side of the first rule in the specification.

Here is a first example.

¹ $S ::= 'a' \text{ R T}$
² $R ::= 'r'$
³ $T ::= 't'$

This CFG has three nonterminals and three terminals, with $N = \{S, R, T\}$ and $T = \{'a', 'r', 't'\}$. The start symbol S is just S because the first rule is a production for S . Our conventions are useful because in reality all we need to do is list the rewrite rules in P — we can infer N , T and S from that list.

4.3.2 Sentential forms and derivations

A sentence is a (possibly empty) string of terminals. A *sentential form* is a (possibly empty) string of terminals and nonterminals. All sentences are sentential forms (but, of course, not all sentential forms are sentences!).

A derivation step rewrites a sentential form Σ into another sentential form Σ' by (a) picking one of the nonterminals X in Σ and then (b) replacing it by the right hand side of one of the productions of X . Sentential forms that contain no nonterminals cannot be rewritten: they are sentences.

Where a sentential form has more than one nonterminal, which should we pick for rewriting? To use the language of rewriting, how do we select a redex? It turns out that it does not matter since CFG rewriting is always confluent. Hence, by convention we always choose the left-most nonterminal to expand. When we are being careful, we call this the left-derivation.

4.3.3 Generating a language

We generate strings by maintaining a current sentential form which is initialised with just the start symbol. We then repeatedly rewrite the left-most nonterminal, generating new sentential forms and sentences.

We can ask ART to do this for us but watch out because interesting grammars have infinite languages, so we need to specify an upper bound on the number of sentential forms we want or the process will not terminate! Using the example above

```

1 S ::= 'a' R T
2 R ::= 'r'
3 T ::= 't'
4 !generate 10 sententialforms

```

gives

```

1 1 _S
2 2 a _R _T
3 3 a r _T
4 "4 a r t

```

ART has listed four sentential forms. In each, nonterminals are preceded with an underscore to show that they are really *variables* which will in due course be replaced. Any sentences (sentential forms without nonterminals) are preceded with a " character.

We have asked for 10 sentential forms, but in fact this very simple grammar only has four sentential forms, so the output stops quickly.

The first sentential form is just the start symbol, as it always is. We then expand the rightmost nonterminal in the only way we can by replacing S with its only right hand side ' a ' R T . Next we must rewrite R as it is now the rightmost symbol, and then we rewrite T at which point the sentential form is a sentence, and we are finished.

We sometimes write these derivations in a more compact form as

$$S \Rightarrow aRT \Rightarrow arT \Rightarrow art$$

Here, the convention is that lower case letters represent terminals and upper case nonterminals.

Sometimes we want to miss out intermediate steps, so we write $\stackrel{*}{\Rightarrow}$ to mean ‘zero or more derivation steps’. So, for instance, we could say

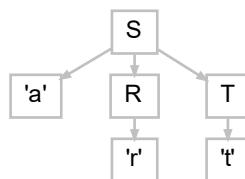
$$S \stackrel{*}{\Rightarrow} arT \Rightarrow art$$

or just

$$S \stackrel{*}{\Rightarrow} art$$

We can now define the language of Γ as $L(\Gamma) = \{u | S \stackrel{*}{\Rightarrow} u, u \in T^*\}$

We can also display the derivation as a *derivation tree*:



The root of the derivation tree is always labelled with the start symbol: it corresponds to the first sentential form in the list above. The leaves of the tree are the terminals in the resulting sentence, and terminals can only appear on the leaves. The internal nodes are labelled with nonterminals, and the children of an internal node are labelled with the elements of the production used to expand that nonterminal instance in the derivation.

As a special case, the production $X ::= \#$ is represented as a node labelled X with a single leaf child labelled $\#$. The empty string $\#$ is neither a terminal nor a nonterminal but a denotation for the absence of either; when writing out a sentence from a derivation tree we list the leaf node terminal labels in order, omitting any $\#$ labels.

As usual, we can represent the derivation tree textually as a prefix form term: for this tree the corresponding term is $S(a, R(r), T(t))$

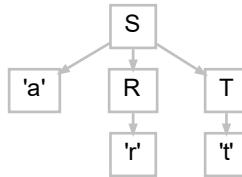
4.3.4 Four representations of derivations

We have seen four representations of the same derivation:

```

1 1 _S
2 2 a _R _T
3 3 a r _T
4 "4 a r t

```



$S \Rightarrow aRT \Rightarrow arT \Rightarrow art$

$S(a, R(r), T(t))$

These representations are all equivalent in the sense that they contain the same information, and that given one we can immediately write out the other three. They are in fact three different external syntaxes for the same internal object — a derivation. All but the bottom left one are outputs from ART; we shall usually use the term style at bottom right.

4.3.5 Adding a rule

Let us add another rewrite rule for nonterminal R

```

1 S ::= 'a' R T
2 R ::= 'r' 'r'
3 R ::= 'r'
4 T ::= 't'
5 !generate 10 sententialforms

```

This yields

```

1 1 _S
2 2 a _R _T
3 3 a r _T
4 4 a r r _T
5 "5 a r t
6 "6 a r r t

```

Sentential forms 3 and 4 show the R in form 2 being replaced by two different productions, and their further expansion to two different sentences at 5 and 6.

ART allows us to group together CFG rules that have the same left hand side, writing the alternate right hand sides out separated by | characters (which we

read as ‘OR’), so we can rewrite the rule for **R** as **R** ::= 'r' 'r' | 'r' giving

```

1 S ::= 'a' R T
2 R ::= 'r' 'r' | 'r'
3 T ::= 't'
4 !generate 10 sententialforms

```

As you would expect, running this script gives the same output as above.

4.3.6 Using recursion to specify infinite languages

There are only two sentences in our language so far. How would we specify the language $\{art, arrt, arrrt, arrrrt, \dots\}$ that is, the set of strings which start with an *a*, finish with a *t* and have one or more *r* characters between? The trick is to use an inductive (recursive) specification:

```

1 S ::= 'a' R T
2 R ::= 'r' | 'r' R
3 T ::= 't'
4 !generate 10 sententialforms

```

The rule for **R** now expands an instance of **R** either to the terminal 'r' or to 'r' R. In the second case, we effectively just insert an **r** into the sentential form, followed by a new instance of **R** so we can go round again over and over until we finally rewrite the **R** to just 'r'.

ARTs output shows this in action:

```

1 1 _S
2 2 a _R _T
3 3 a r _T
4 4 a r _R _T
5 "5 a r t
6 6 a r r _T
7 7 a r r _R _T
8 "8 a r r t
9 9 a r r r _T
10 10 a r r r _R _T

```

ART is performing *breadth first expansion* here in which all of the possible left most expansions of the current sentential form are loaded onto the end of queue after which the start of the queue is unloaded to the current sentential form.

If we just want to see sentences, we can request them with **!generate 10 sentences** which gives

```

1 "1 a r t
2 "2 a r r t
3 "3 a r r r t
4 "4 a r r r r t
5 "5 a r r r r r t
6 "6 a r r r r r r t
7 "7 a r r r r r r r t
8 "8 a r r r r r r r r t
9 "9 a r r r r r r r r r t
10 "10 a r r r r r r r r r r t

```

4.3.7 Left and right recursion

The previous rule $R ::= 'r' \mid 'r' R$ is an example of *right* recursion because the recursive self-reference is at the right hand end. We can generate the same language with the left recursive rule $R ::= 'r' \mid R 'r'$

```

1 S ::= 'a' R T
2 R ::= 'r' | 'r' R
3 T ::= 't'
4 !generate 10 sententialforms

```

yields

```

1 1 _S
2 2 a _R _T
3 3 a r _T
4 4 a _R r _T
5 "5 a r t
6 6 a r r _T
7 7 a _R r r _T
8 "8 a r r t
9 9 a r r r _T
10 10 a _R r r r _T

```

In this case we again insert an r into the sentential form, but this time preceded by a new instance of R .

In future examples we shall be interested in exploiting these variant growth patterns, both of which are useful. However some parsing techniques, including the RDSOB algorithm described in this chapter, will go into an infinite loop when processing a left recursive grammar rule. ART's default MGLL parser does not have this defect and will handle any CFG specification.

4.3.8 Ambiguity

There may more than one way to generate a sentence; that is there may be some sentence σ which can be derived in more than one way. A CFG with this property is called *ambiguous*. Here is an example of an ambiguous CFG which generates only one sentence, but in two ways.

```

1 S ::= 'a' R | A 't'
2 A ::= 'a' 'r'
3 R ::= 'r' 't'
4 !generate 10 sententialforms

```

yields

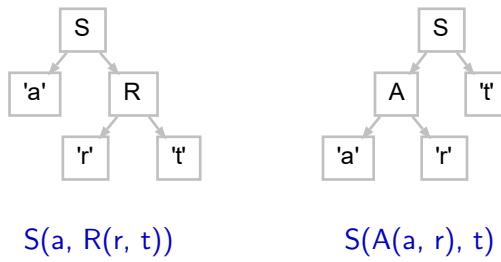
```

1 1 _S
2 2 a _R
3 3 _A t
4 "4 a r t
5 "5 a r t

```

There are two derivations here: $S \Rightarrow aR \Rightarrow art$ and $S \Rightarrow At \Rightarrow art$ and as a result the sentence *art* appears twice in the list at positions 4 and 5.

The two derivation trees and their terms are



When parsing, ART will find both derivations. *Usually* we want to select only one, and we can do that with a *chooser* relation which we shall discuss in section 4.5. If no choosers are specified, ART parsers use a default derivation selection strategy called *longest match*: wherever there is a choice of derivation steps we select the one that matches the longer part of the input. In this case, ART would select $S \Rightarrow At \Rightarrow art$ corresponding to $S(A(a, r), t)$ since the **A** child of **S** in the right hand derivation matches two elements of the input, whereas the '**a**' child of **S** in the left hand derivation only matches one element.

4.3.9 Cycles

A particularly unpleasant situation arises when a CFG has a nonterminal that can derive just itself. Such a rule is called *cyclic* because it will cause a generator

to loop forever without producing any sentences.

```

1 S ::= 'a' R T
2 R ::= 'r'| R // cycle!
3 T ::= 't'
4 !generate 10 sententialforms

```

yields

```

1 _S
2 a _R _T
3 a r _T
4 a _R _T
5 "5 a r t
6 a r _T
7 a _R _T
8 "8 a r t
9 a r _T
10 a _R _T

```

The problem here is that every time we use the rule $R ::= R$ we are simply replacing an instance of R with an instance of R , so the sentential form does not change and we shall keep going forever.

This particular example is easy to spot, as it is *directly* cyclic, but sometimes cycles arise within a chain of productions.

Although they occasionally appear in programming language standards, cyclic grammars are best thought of as laboratory curiosities to be avoided in practical applications of CFGs. ART will happily handle cycles, returning an encoding of the derivations which uses a loop to show where the cycles are, but derivation selection is challenging in the presence of cycles.

4.3.10 A useful letter generator

We round off this section with a useful (maybe) specification that generates thank you letters:

```

1 letter ::= salutation person ',' 'thank' 'you' 'for' 'my' gift '.' 'It' 'is' reaction '!' 'I' 'will' action 'it' '!'
2 gift ::= occasion 'present'
3 occasion ::= 'birthday' | 'house' 'warming'
4 salutation ::= 'Dearest' | 'Hey'
5 person ::= 'you' | 'Madam' | 'Sir'
6 reaction ::= 'incredible' | 'in' 'the' 'bin'
7 action ::= 'love' | 'hide' | 'try' 'to' 'forget'

```

ART tells us that this grammar has exactly 72 sentences. We could work that out without actually generating them by noting that

1. There are no recursive rules, so the language is finite.
2. Rules `occasion`, `salutation` and `reaction` have two alternate productions each
3. Rules `person` and `action` each have three alternates.
4. Rule `gift` has one instance of `occision` which has two productions, so there only two ways to rewrite `gift`
5. The start rule has only one production which has one instance of each of `salutation`, `person`, `gift`, `reaction` and `action`. To get the total number of sentences that can be generated from `letter`, we multiply together the number of sentences in each of their sublanguages to get

$$2 \times 3 \times 2 \times 2 \times 3 = 72$$

My favourite sentence is number 12:

`1 "12 Dearest you, thank you for my house warming present. It is in the bin. I will try to forget it.`

4.4 Parsing

Language generation is helpful way of understanding CFGs and can be useful if we want to construct a random set of test sentences for our language processors, but by far the most common application of CFGs is to *parsing*: given a CFG Γ and a putative sentence σ , a general parser should either return the full set of derivations if $\sigma \in L(\Gamma)$ or a helpful error message if $\sigma \notin L(\Gamma)$.

In ART, we activate the parser with a `!try` directive: for instance

```

1 S ::= 'a' R T
2 R ::= 'r'
3 T ::= 't'
4
5 !try "art"
6 !print outcome
7 !print derivation
8 !show derivation

```

which yields

```

1 Accept
2 current derivation term: [1480]
3 S(a, R(r), T(t))

```

The `!try "art"` directive sets the current input to `art` and runs the parser.

`!print outcome` will print `Accept` or `Reject` depending on whether the string can be generated by those rules, and `!print derivation` will print out the derivation term. The `!show derivation` directive produces a *visualisation* of the derivation tree, which maybe a pop up window in the IDE or a file `derivation.dot` which may be displayed using the GraphViz tools.

If we ask ART to parse a string that is not in the language of the rules, we get an error message:

```

1 !try "ara"
2 !print outcome
3 !print derivation

```

yields

```

1 1,2 GLLBL syntax error
2   1: ara
3   -----
4 Reject
5 current derivation term: [0]
6 null term

```

The message `1,2 GLLBL syntax error` means that ART was unable to find a derivation, and the furthest it got whilst parsing was line 1,column 2. ART will then print out the line in question with a pointer to the column position. The string is rejected, and the resulting derivation term is the null term.

We shall have much more to say about parsers and parsing later, but for now all we need to know is that a `!try` directive will parse an input string if it can, and then automatically pass the resulting term on to an interpreter if interpreter rules have been defined.

4.5 Derivation selection using choosers

Not in 2025 curriculum

4.6 GIFT annotations

It is useful to be able to compress derivation trees into trees which carry only such information from the derivation that we wish to carry forward into other stages of the translation process. Common transformations include:

- ◊ the suppression of recursion-scaffolding nodes,
- ◊ the construction of expression trees made up solely of nodes labeled with terminals,

- ◊ the suppression of entire sub-trees,
- ◊ the local reordering of sub-trees,
- ◊ the insertion of new pieces of tree.

The GIFT formalism provides a small set of operations with postfix annotations that specify their application to the tree nodes associated with grammar elements. We specify them by writing them into the grammar, but it is helpful to think of them being attached to tree nodes.

GIFT stands for Gather-Insert-Fold-Tear. The ART tool presently only implements the two Fold operations, but we shall discuss applications of the other operators. Collectively, the GIFT operations may be viewed as special cases of a more general approach called term rewriting, which allows tree to be rewritten using tree-to-tree rewrite rules.

The best way to think about the GIFT operators is that they are annotations that are loaded into the derivation tree, and that a GIFT rewriting phase then rewrites the derivation tree under the control of those operators into a Rewritten Derivation Tree (RDT).

4.6.1 Fold operators

The fold operators can only be applied to a node which has a parent: that is the root node may not be folded.

There are two kinds of fold: fold-under (^) and fold-over (^^).

A rule such as

```
1 | X ::= `a `b `c^ `d
```

will generate an (as-yet-unrewritten) derivation subtree of the form

```
1 |      X
2 |      / / \ \
3 |      a b c^ d
```

and rule such as

```
1 | Y ::= `a `b `c^^ `d
```

will generate derivation subtree of the form

```
1 |      X
2 |      / / \ \
3 |      a b c^^ d
```

The idea of the fold operators is that the edge joining the annotated node to its parent is folded in half so that the child node and the parent node are coincident. If we fold under (^) then the child goes under the parent; if we fold over then the child goes over the parent. Alternatively, you can see that for a fold under we delete the child node and keep the parent node; if we fold over then we delete the parent node and replace it with the child node.

For fold-under, then, we have

X ::= a b c^ d

gives

1	X => X
2	/ / \ \ / / \
3	a b c^ d a b d

and

X ::= a b c^ d

gives

1	X => c
2	/ / \ \ / / \
3	a b c^ d a b d

Note that this allows us to build trees which have *terminals* as internal nodes.

So far, we have only considered fold operators on terminal nodes, which have no children. If we apply a fold operator to a nonterminal instance, then we must explain how the children are to be treated. The metaphor of edge folding helps here: the children of the annotated node are inserted as a group into the siblings of the annotated node. We can think of this as the children being dragged up a level in the tree.

For fold-under

1	X ::= a b Y^ d
2	Y ::= y z

1	X => X
2	/ / \ \ / / \ \
3	a b Y^ d a b y z d
4	/\
5	y z

For fold-over

```

1 X ::= a b Y^^ d
2 Y ::= y z

```

```

1 X => Y
2 / / \ \ / / | \ \
3 a b Y^^ d a b y z d
4   /\
5     y z

```

4.6.2 The Tear operator

We can suppress an entire subtree by attaching the Tear ($^{^{\wedge\wedge\wedge}}$) annotation.

```

1 X ::= a b Y^^^ d
2 Y ::= y z

```

```

1 X => X
2 / / | \ / | \
3 a b Y^^^ d a b d
4   /\
5     y z

```

4.6.3 Insertions

Nodes can be named by appending a colon and an identifier, and named tear nodes can be inserted elsewhere in the tree:

```

1 X ::= a b Y:t^^ d [t]
2 Y ::= y z

```

```

1 X => X
2 / / | \ / | | \
3 a b Y^^^ d a b d Y
4   / \ / \
5     y z y z

```

4.6.4 The Gather operator

Sometimes we want to bring together nonterminal subtrees under a new parent.

4.7 GIFT applications

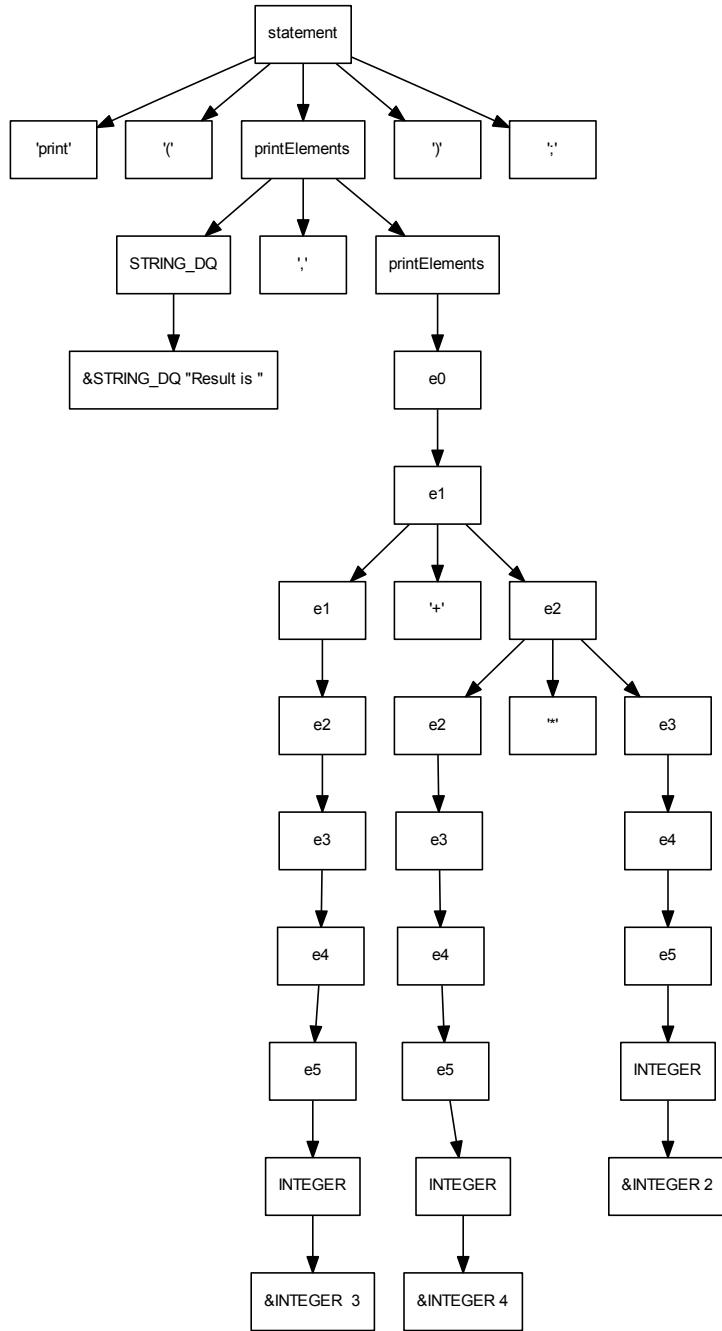
It is often convenient to be able to represent expression trees as being made up of operators and operands. Operators such as + and × are grammar terminals, and we can achieve this affect by promotin the operator symbols over their parent nonterminals.

Here is a first attempt, in which the `minisyntax` grammar has been modified so that every operator symbol has had a `^` annotation applied to it.

```

1 statement ::= 'print' '(' printElements ')'; (* print statement *)
2
3 printElements ::= STRING_DQ |
4                 STRING_DQ '^' printElements |
5                 e0 | e0 '^' printElements ;
6
7 e0 ::= e1 |
8     e1 '>'^ e1 | (* Greater than *)
9     e1 '<'^ e1 | (* Less than *)
10    e1 '>=''^ e1 | (* Greater than or equals*)
11    e1 '<=''^ e1 | (* Less than or equals *)
12    e1 '==''^ e1 | (* Equal to *)
13    e1 '!='^ e1 ; (* Not equal to *)
14
15 e1 ::= e2 |
16     e1 '+'^ e2 | (* Add *)
17     e1 '-'^ e2 ; (* Subtract *)
18
19 e2 ::= e3 |
20     e2 '*'^ e3 | (* Multiply *)
21     e2 '/'^ e3 | (* Divide *)
22     e2 '%'^ e3 ; (* Mod *)
23
24 e3 ::= e4 |
25     '+'^ e3 | (* Posite *)
26     '-'^ e3 ; (* Negate *)
27
28 e4 ::= e5 |
29     e5 '**'^ e4 ; (* exponentiate *)
30
31 e5 ::= INTEGER | (* Integer literal *)
32     '(' e1 ')'; (* Parenthesised expression *)
33
34 STRING_DQ ::= &STRING_DQ ;
35
36 INTEGER ::= &INTEGER ;

```



```

1 statement ::= 'print'^^ '('^ printElements^ ')'^ ','^ ';' (* print statement *)
2
3 printElements ::= STRING_DQ |
4   STRING_DQ ','^ printElements^ |
5   e0 | e0 ','^ printElements^ ;
  
```

```

6
7 e0 ::= e1^^ |
8     e1 '>'^^ e1 | (* Greater than *)
9     e1 '<'^^ e1 | (* Less than *)
10    e1 '>=''^^ e1 | (* Greater than or equals*)
11    e1 '<=''^^ e1 | (* Less than or equals *)
12    e1 '==''^^ e1 | (* Equal to *)
13    e1 '!='^^ e1 ; (* Not equal to *)
14
15 e1 ::= e2^^ |
16     e1 '+'^^ e2 | (* Add *)
17     e1 '-'^^ e2 ; (* Subtract *)
18
19 e2 ::= e3^^ |
20     e2 '*'^^ e3 | (* Multiply *)
21     e2 '/'^^ e3 | (* Divide *)
22     e2 '%'^^ e3 ; (* Mod *)
23
24 e3 ::= e4^^ |
25     '+'^^ e3 | (* Posite *)
26     '-'^^ e3 ; (* Negate *)
27
28 e4 ::= e5^^ |
29     e5 '**'^^ e4 ; (* exponentiate *)
30
31 e5 ::= INTEGER^^ | (* Integer literal *)
32     '(!'^^ e1^^ '!'^^); (* Parenthesised expression *)
33
34 STRING_DQ ::= &STRING_DQ ;
35
36 INTEGER ::= &INTEGER ;

```

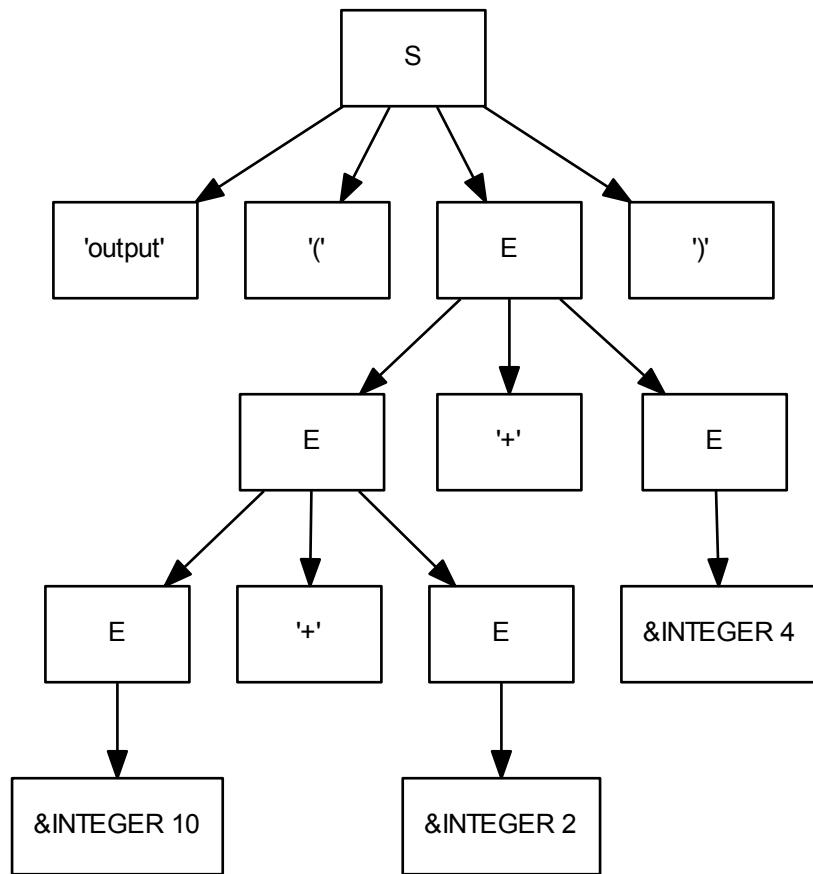
A *term* is a well formed formula (wff) in some formal language Γ . Our formal view of semantics will be a game in which terms representing program fragments and semantic entities such as the store and output will evolve through transitions. We shall move between three different ways of representing terms: (a) as the formula itself, as a derivation tree of the formula in an abstract syntax Γ' , and as a fully parenthesized representation of the that derivation tree. So, for instance, if Γ has these productions

```

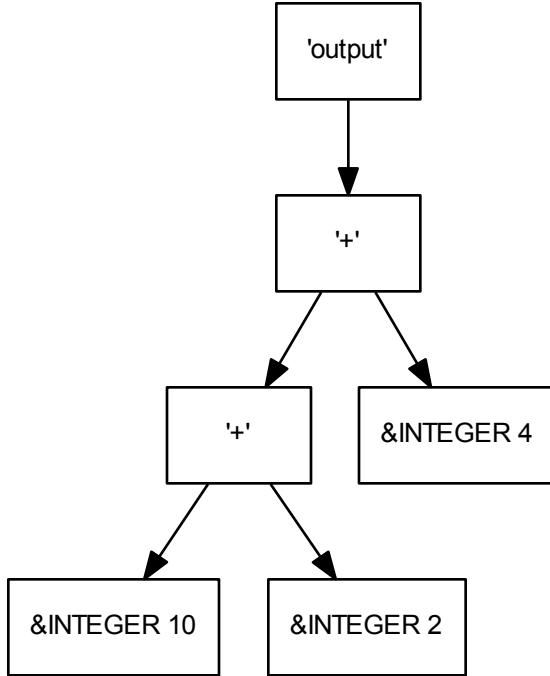
S ::= 'output' '(' E ')'
E ::= E '+' E
E ::= E '*' E
E ::= INTEGER

```

then the well formed formula `output(10+2*4)` has this derivation tree



Using GIFT annotations, we could map this to a more compact trees:



and we would then represent the derivation tree term as `output(+(+10, 2), 4)`

4.8 RDSOB: Implementing a parser toolchain

When we implement a translator, we parse the source language into an intermediate form, and then traverse the intermediate form outputting the object language.

A *parser generator* is a program which reads specifications for a grammar Γ written in BNF (or EBNF) and outputs the source code for a parser. When compiled the parser will test strings to see if they are in the language $L(\Gamma)$, and perhaps build a derivation tree.

Parser generators, then, can be thought of as processors for a DSL (BNF) which translates to, say, Java. Embedded within each parser will be a *parsing algorithm*. We shall illustrate this process using a parsing algorithm called Recursive Descent with Singleton Ordered Backtracking (RDSOB). It is a rather limited algorithm: its advantage is that it is easy to understand and easy to generate. This means that it is possible to fully explain the internals not just of the parsers but of parser generators: that is programs that write out the parsers.

4.9 Recursive Descent with Singleton Ordered Backtrack parsing

RDSOB is a long acronym for a very simple parsing technique. The parsers may be written by hand, and there is no requirement to compute properties of the grammar: in fact an RDSOB parser can be produced as a syntax-directed translation from BNF; it is in effect a pretty-printed version of the grammar.

RDSOB is not used for production parsers because (a) the performance of RDSOB parsers is exponential in the length of the input string for some ‘nasty’ grammars and (b) because RDSOB parsers fail to recognise some strings that are in the language of the grammar being parsed.

(b) sounds disastrous, but in fact the commonly-used parsing techniques such as LALR(1), SLR(1) and LL(1) all suffer from the same problem. In fact any non-general parsing technology will fail to accept some strings for some grammars. However, it is possible to compute in advance whether a grammar is LALR(1) (or SLR(1) or LL(1)...) and so the user is at least told that their grammar will not behave as they expect. Although we could do some processing to help the user (and in our RDSOB toolchain we do test for one obvious error condition as we shall see) a basic syntax driven translation produces a parser which silently misbehave. The user can write what appears to be a perfectly reasonable grammar, have a parser generated and then find that it does not work as it should: we call these situations *nasty surprises*.

4.9.1 The RDSOB algorithm

Consider a grammar $\Gamma = (N, T, X_S, P)$ where, as usual, N is a set of nonterminals $\{X_1, X_2, X_3, \dots, X_k\}$, T is a set of terminals, X_S is the start nonterminal and P is a set of productions $\{X \rightarrow \rho, \rho \in (N \cup T)^*\}$.

The RDSOB algorithm works on *ordered* grammar. In an ordered grammar the subset of productions $\{X_i \rightarrow \rho_1, X_i \rightarrow \rho_2, \dots\}$ are ordered, and are tested in that order. There is no ordering associated with the nonterminals or the terminals; it is just the order of productions *within* a particular nonterminal X_i that is significant.

This seemingly innocuous change has a big impact on the languages that can be successfully parsed by RDSOB compared to a truly general technique such as GLL. We are highlighting this difference here because occasionally one encounters parsing tools which use algorithms based on ordered grammars, and in my experience the authors often do not adequately explain the limitations of the technique.

Informally, an RDSOB parser is a set of (possibly recursive) functions, one per nonterminal. The functions take no parameters, and return a boolean. The input string is held in a buffer `String input` and there is a global variable `int`

`cc` which holds the index of the *current character*.

At the start of the parse function for nonterminal X_i , the value of `cc` on entry is remembered in a local variable `int rc` which holds the index of the *restart character*. Each alternate production $X_i \rightarrow \rho_j$ is then laid out as a nest of `if` statements: for a terminal we test against a direct match; for a nonterminal we call the appropriate parse function and for ϵ we do nothing. If the nest evaluates true, then we have found a match against that alternate, and so the parse function returns true. If not, we proceed to alternate $X_i \rightarrow \rho_{j+1}$. If all alternate productions fail, the parse function returns false.

Each parse function also remembers which alternate (if any) succeeded. The running parser maintains a global array of integers called the *oracle* and a global variable `int co` which holds the index of the next free slot in the oracle.

An RDSOB parser explores the grammar by recursively calling the parse functions. Sometimes these exploration fail after several layers of function call, and in that case the parser *backtracks* to the next level up so as to continue testing its alternate productions. In fact the backtracking can recursively unwind an arbitrary number of levels. As a result, we need to remember where we were in the oracle when we entered the parse function so that we can rest `co` at the start of each alternate; the local variable `int ro` remembers this restart oracle index.

4.9.2 An RDSOB example in Java

Here is a small grammar.

```

1 S ::= 'b' | 'a' X '@'
2 X ::= 'x' X | #

```

We specify terminals within single quotes, and ϵ is written as `#`. Alternate productions are separated by a vertical bar, and each rule is terminated with a period. Repeated nonterminals are not allowed.

The language of this grammar is `{ b, a@, ax@, axx@ axxx@, ... }`.

When processed by the RDSOB parser generator, the following two Java parse functions are produced:

```

1 boolean parse_S() {
2     int rc = cc, ro = co;
3
4     /* Nonterminal S, alternate 1 */
5     cc = rc; co = ro; oracleSet(1);
6     if (match("b")) { return true; }
7

```

```

8  /* Nonterminal S, alternate 2 */
9  cc = rc; co = ro; oracleSet(2);
10 if (match("a")) {
11   if (parse_X()) {
12     if (match("@")) { return true; } }
13
14   return false;
15 }
16
17 boolean parse_X() {
18   int rc = cc, ro = co;
19
20   /* Nonterminal X, alternate 1 */
21   cc = rc; co = ro; oracleSet(1);
22   if (match("x")) {
23     if (parse_X()) { return true; } }
24
25   /* Nonterminal X, alternate 2 */
26   cc = rc; co = ro; oracleSet(2);
27   /* epsilon */ return true;
28 }
```

Let us follow this code through whilst parsing the string `axx@`.

The parser initially loads `input` with the string `axx@§` where `§` is not the dollar character but rather stands for some special end-of-string marker. (In Java, we use the character containing zero or '`\0`'; regular expression processors and parsing texts conventionally use some variant of the dollar symbol.) The oracle does not need to be initialised, but the two global variables that index the input and the oracle are zeroed: `cc = co = 0`

We start the parse by calling the `parse` function for the start symbol `parse_S()`.

`parse_S()` remembers the entry values for the input and oracle indices before executing the clauses for the alternates in sequence.

Each clause begins by setting the global indices to these restart values before testing the input against the production using a nest of predicates each of which either call `match()` to test a terminal or call the relevant `parse_()` function. If none of the clauses succeed, then the `return false;` statement is executed.

By instrumenting the parse we can produce a trace of the function calls and terminal matches. Here is the output from one run.

```

Input: 'a x x @'
S() at rc = 0, cc = 0 'a'
  S() alternate 1 rc = 0, cc = 0 'a'
    At 0 'a' match b - reject
```

```

S() alternate 2 rc = 0, cc = 0 'a'
    At 0 'a' match a - accept
X() at rc = 2, cc = 2 'x'
    X() alternate 1 rc = 2, cc = 2 'x'
        At 2 'x' match x - accept
X() at rc = 4, cc = 4 'x'
    X() alternate 1 rc = 4, cc = 4 'x'
        At 4 'x' match x - accept
X() at rc = 6, cc = 6 '@'
    X() alternate 1 rc = 6, cc = 6 '@'
        At 6 '@' match x - reject
X() alternate 2 rc = 6, cc = 6 '@'
    At 6 '@' match @ - accept
Accepted
Oracle: 2 1 1 2

```

This shows us that a call to `S()` with `cc=0` tried alternates 1 and 2, before a call to `X()` with `cc=2` tries alternate 1 and then calls `X()` with `cc=4` which calls `X()` with `cc=6`. Alternate 1 cannot match `@` tp `x`, so alternate 2 is tried instead which must succeed because it is an ϵ -production.

All of the calls then unwind, and because the call to `S()` terminates with the current character `cc` pointing to the end of string marker, the string is **Accepted**.

The parser then prints out the oracle which (when combined with the grammar) encodes the successful derivation: we took alternates 2, 1, 1 and 2 as we went down through the nest of parse functions.

4.10 Engineering a complete Java parser

The parse functions in the previous section make use of auxilliary functions like `match()` and also require some initialisation code. In this section we shall look at how the generated parse functions are embedded into Java classes so as to make a standalone parser.

A minimal RDSOB parser comprises two classes in two source files:

1. `ARTRDSOBBase.java` which contains auxiliary methods.
2. `ARTGeneratedParser.java` which extends class `ATYRDSOBBase` with the parse member functions and a `main()` function which processes command line arguments.

Here is the contents of `ARTGeneratedParser.java` for a simple parse — in later sections we shall add more functions to support semantics processing and tree construction.

```

1 import java.io.File;
2 import java.io.PrintWriter;
3 import java.io.FileNotFoundException;
4 import java.util.Scanner;
5 import java.util.ArrayList;
6
7 class ARTGeneratedParser extends uk.ac.rhul.cs.csle.artRDSOB.ARTRDSOBBBase {
8     String input;
9     int cc, co, oracleLength, oracle[];
10    int ts, te;
11
12 ARTGeneratedParser() { oracleLength = 1000; oracle = new int[oracleLength]; cc = co = 0;}
13
14 String readInput(String filename) throws FileNotFoundException {
15     return new Scanner(new File(filename)).useDelimiter("\Z").next() + "\0";
16 }
17
18 void oracleSet(int i) {
19     if (co == oracleLength) {
20         int oracleLengthOld = oracleLength;
21         oracleLength += oracleLength / 2;
22         int newOracle[] = new int[oracleLength];
23         System.arraycopy(oracle, 0, newOracle, 0, oracleLengthOld);
24         oracle = newOracle;
25     }
26     oracle[co++] = i;
27 }
28
29 boolean match(String s) {
30     if (input.regionMatches(cc, s, 0, s.length())) {
31         cc += s.length();
32         builtIn_WHITESPACE();
33         return true;
34     }
35     return false;
36 }
37
38 boolean builtIn_WHITESPACE() {
39     while(Character.isWhitespace(input.charAt(cc)))
40         cc++;
41     return true;
42 }
43 }
```

The constructor sets the initial oracle length 1000, creates the array and zeroes the cc and co indices. Loading the input string is the responsibility of the

generated parse class `SBxyz`, though it makes use of function `readInput()` which automatically appends an end-of-string marker '`\0`'.

The `oracleSet(int i)` function resizes the oracle if necessary (adding 50% to its length at each resize operation) before loading the supplied alternate number `i` into the oracle and incrementing `co`, the current oracle index.

The `match(String s)` function checks to see whether a substring of `input` starting at character `cc` matches parameter `s`. If so, then a helper function `builtin_WHITESPACE()` is called to absorb any blank spaces and line ends after the string. This allows generated parsers to treat the strings `axx@` and `a x x @` as equivalent.

If the example grammar is in file `test.sb`, the parser generator generates this file `SBtest.java`.

```

1 import java.io.FileNotFoundException;
2
3 class ARTGeneratedParser extends uk.ac.rhul.cs.csle.artRDSOB.ARTRDSOBBBase {
4     boolean parse_S() {
5         int rc = cc, ro = co;
6
7         /* Nonterminal S, alternate 1 */
8         cc = rc; co = ro; oracleSet(1);
9         if (match("b")) { return true; }
10
11        /* Nonterminal S, alternate 2 */
12        cc = rc; co = ro; oracleSet(2);
13        if (match("a")) {
14            if (parse_X()) {
15                if (match("@")) { return true; }}}
16
17        return false;
18    }
19
20    boolean parse_X() {
21        int rc = cc, ro = co;
22
23        /* Nonterminal X, alternate 1 */
24        cc = rc; co = ro; oracleSet(1);
25        if (match("x")) {
26            if (parse_X()) { return true; }}
27
28        /* Nonterminal X, alternate 2 */
29        cc = rc; co = ro; oracleSet(2);
30        /* epsilon */ return true;
31    }
32}
```

```

33 SBtest(String filename) throws FileNotFoundException {
34     input = readInput(filename);
35
36     System.out.printf("Input: '%s'%n", input);
37     cc = co = 0; builtin_WHITESPACE();
38     if (!(parse_S() && input.charAt(cc) == '\0'))
39         { System.out.print("Rejected%n"); return; }
40
41     System.out.print("Accepted%n");
42     System.out.print("Oracle:");
43     for (int i = 0; i < co; i++) System.out.printf(" %d", oracle[i]);
44     System.out.printf("%n");
45 }
46
47 public static void main(String[] args) throws FileNotFoundException{
48     if (args.length < 1)
49         new SBtest("");
50     else
51         new SBtest(args[0]);
52 }
53 }
```

The `main()` function collects the name of a string file from the input and then instances `SBtest`, passing the filename as an argument to the constructor.

The constructor loads the `input` string from the supplied filename; prints it out; and then calls `builtin_WHITESPACE()` to consume any leading blanks in the input string.

If the start symbol's parse function consumes the entire string up to but not including the end-of-string marker, the string is accepted and the oracle printed out.

4.11 Using built in matchers

The RDSOB base class provides a set of builtin matchers which can be used to efficiently parse identifiers, numeric literals, strings and so on. It is quite easy to add new builtins as required.

The parser generator translates pseudo-terminals such as `&ID` into calls to the corresponding builtin matcher. In detail, `&XYZ` is trabslated to a call to `builtin_XYZ()`. The generator does not check the name of the builtin, so just by adding a new builtin member function to class `uk.ac.rhul.cs.csle.artRDSOB.ARTRDSOBBBase` we can extend the repertoire of builtin matchers.

Here is a slightly modified version of our test grammar.

¹ S ::= 'b' | 'a' X '@' .

`2| X ::= &ID X | # .`

It generates these parse functions.

```

1 boolean parse_S() {
2     int rc = cc, ro = co;
3
4     /* Nonterminal S, alternate 1 */
5     cc = rc; co = ro; ora cleSet(1);
6     if (match("b")) { return true; }
7
8     /* Nonterminal S, alternate 2 */
9     cc = rc; co = ro; oracleSet(2);
10    if (match("a")) {
11        if (parse_X()) {
12            if (match("@")) { return true; }}}return false;
16}
17
18 boolean parse_X() {
19     int rc = cc, ro = co;
20
21     /* Nonterminal X, alternate 1 */
22     cc = rc; co = ro; oracleSet(1);
23     if (builtIn_ID()) {
24         if (parse_X()) { return true; }}return true;
28}
```

which are identical to the previous versions except that the call to `match("x")` has been replaced by a call to `builtIn_ID()`.

Here is the source code for a set of builtins. Note how each of them calls `builtin_WHITESPACE()` after matching. Each builtin matcher remembers the start and end indices of the matched terminal in global variables `int ts` and `int te`.

```

1 boolean builtIn_ID() {
2     if (!Character.isJavaIdentifierStart(input.charAt(cc))) return false;
3     ts = cc++;
4     while (Character.isJavaIdentifierPart(input.charAt(cc)))
5         cc++;
6     te = cc;
```

```

7   builtIn_WHITESPACE();
8   return true;
9 }
10
11 boolean isxdigit(char c) {
12   if (Character.isDigit(c)) return true;
13   if (c >= 'a' && c <= 'f') return true;
14   if (c >= 'A' && c <= 'F') return true;
15   return false;
16 }
17
18 boolean builtIn_INTEGER() {
19   if (!Character.isDigit(input.charAt(cc))) return false;
20   ts = cc;
21   /* Check for hexadecimal introducer */
22   boolean hex = (input.charAt(cc) == '0' &&
23                 (input.charAt(cc + 1) == 'x' ||
24                  input.charAt(cc + 1) == 'X'));
25   if (hex) cc += 2; // Skip over hex introducer
26   /* Now collect decimal or hex digits */
27   while (hex ? isxdigit(input.charAt(cc)) :
28         Character.isDigit(input.charAt(cc)))
29     cc++;
30   te = cc;
31   builtIn_WHITESPACE();
32   return true;
33 }
34
35 boolean builtIn_REAL() {
36   if (!Character.isDigit(input.charAt(cc))) return false;
37   ts = cc;
38   while (Character.isDigit(input.charAt(cc)))
39     cc++;
40   if (input.charAt(cc) != '.')
41     return true;
42   cc++; // skip .
43   while (Character.isDigit(input.charAt(cc)))
44     cc++;
45   if (input.charAt(cc) == 'e' || input.charAt(cc) == 'E') {
46     cc++;
47     while (Character.isDigit(input.charAt(cc)))
48       cc++;
49   }
50   te = cc;
51   builtIn_WHITESPACE();
52   return true;
53 }
```

```

54
55 boolean builtIn_CHAR_SQ() {
56     if (input.charAt(cc) != '\'') return false;
57     cc++;
58     ts = cc;
59     if (input.charAt(cc) == '\\\\')
60         cc++;
61     cc++;
62     if (input.charAt(cc) != '\'') return false;
63     te = cc;
64     cc++; // skip past final delimiter
65     builtIn_WHITESPACE();
66     return true;
67 }
68
69 boolean builtIn_STRING_SQ() {
70     if (input.charAt(cc) != '\'') return false;
71     ts = cc + 1;
72     do {
73         if (input.charAt(cc) == '\\\\')
74             cc++;
75
76         cc++;
77     }
78     while (input.charAt(cc) != '\'');
79     te = cc;
80     cc++; // skip past final delimiter
81     builtIn_WHITESPACE();
82     return true;
83 }
84
85 boolean builtIn_STRING_DQ() {
86     if (input.charAt(cc) != '"') return false;
87     ts = cc + 1;
88     do {
89         if (input.charAt(cc) == '\\\\')
90             cc++;
91         cc++;
92     }
93     while (input.charAt(cc) != '"');
94     te = cc;
95     cc++; // skip past final delimiter
96     builtIn_WHITESPACE();
97     return true;
98 }
99
100 boolean builtIn_ACTION() {

```

```

101 if (!(input.charAt(cc) == '[' &&
102     input.charAt(cc + 1) == '*'))
103     return false;
104 cc += 2;
105 ts = cc;
106 while (true) {
107     if (input.charAt(cc) == 0)
108         break;
109     if (input.charAt(cc) == '*' && input.charAt(cc) == ']') {
110         cc += 2;
111         break;
112     }
113     cc++;
114 }
115 te = cc - 2;
116 builtIn_WHITESPACE();
117 return true;
118 }
```

4.12 Using attributes and inline semantics

Our generated RDSOB parsers will execute embedded semantic actions which may also use synthesized attributes. The current version does not support inherited attributes, but it is not hard to extend the parser generator to allow that. Only a single pass is made over the input string which significantly limits the kinds of behaviour which may be generated. However, the generated parsers can also generate explicit derivation trees which may be passed to a back end for arbitrary processing: we shall look at tree generation in the next section.

An embedded action is delimited by { } brackets and must be written in the implementation language for the generated parser (in our case Java, although an ANSI C++ versions exist for which actions must be written in C++).

Here is our example grammar extended with an action to report the location of matching x characters.

```

1 S ::= 'b' | 'a' X '@'.
2 X ::= 'x' [* System.out.printf("Matched an x at location %d%n", cc); *] X | #.
```

The generated parser running on the string a x x @ displays:

```

Input: 'a x x @ '
Accepted
Oracle: 2 1 1 2
```

```

Semantics phase
Matched an x at location 4
```

```
Matched an x at location 6
```

The parse is as before. After parsing is completed, a second pass is made during which the semantics are executed.

4.12.1 Attributes

Simply printing out messages showing where we are in a parse is interesting, but limited. If we want to perform useful computations, it turns out that we need to pass information *between* parse functions or, equivalently, around the derivation tree.

Recursive descent parsers provide a natural built-in mechanism for passing information around: we can use the parse function parameters to pass information down the derivation tree and the function return values to pass information up.

The formal underpinnings for this approach are part of the theory of *attribute grammars*. Attributes are classified as *synthesized* which means that move up the tree (like a return result) or *inherited* which means that they pass down the tree (like a parameter). In a general attribute grammar information can move round the derivation tree in arbitrary ways by making use of inherited and synthesized attributes, and the calculation of the final result requires an analysis of the dependency relationships between attribute definitions and their users. An *attribute evaluator* is a general tool for doing just that.

Two useful classes of attribute grammar are the *L-attributed* class in which attributes must be resolvable in a single top-down left to right pass and *S-attributed* grammars which may only contain synthesized attributes. Recursive descent parsers naturally support L-attributed grammars whilst bottom up parsing techniques such as LALR(1) (that is, Bison and YACC) naturally support S-attributed grammars. (In detail, tools often also make use of global attributes which extends their power a little.)

Here is a grammar that uses synthesized attributes to add up the number of 1's seen in a binary string:

```

1 S ::= 'b' | 'a' X:result [* System.out.printf("Result is %d\n", result); *] '@' .
2 X:int ::= '1' X:sum [* rv = sum + 1; *] |
3           '0' X:rv | # .

```

The language of this grammar is

```
{ b, a@, a1@, a0@, a11@, a10@, a01@, a00@, a111@, ...}.
```

The attributes and associated semantic actions implement a recursive function that runs along the string of 1's and 0's maintaining a count of the number of 1's seen.

The return value attribute `rv` is automatically defined for any parse function that has an associated type, and is used to carry the synthesized information back up the tree, or equivalently to pass it back to the calling function. At the top level, the accumulated value is printed out.

Sandbox decides on the type of attributes by looking at the type annotation for the left hand side of the associated nonterminal. In this case, since nonterminal `X` is declared as being of type `int`, the `sum` and `result` attributes will also be of type `int`.

The result of running this parser on the string `a101@` is

```
Input: 'a 1 0 1 @ '
Accepted
Oracle: 2 1 2 1 3

Semantics phase
Result is 2
```

4.12.2 A four function calculator

Let us now extend our example to a more general computing language: a four function calculator for integer constants of one, two or three digits. Warning: Sandbox parsers do not allow left recursion, so all of the operators have been implemented in right associative form, whereas they should really be left associative.

```

1 S ::= exprs:val [* System.out.printf("Final result: %d\n", val); *] .
2
3 exprs:int ::= add:val ';' [* System.out.printf("Result: %d\n", val); *] exprs:rv |
4     add:rv [* System.out.printf("Result: %d\n", rv); *] .
5
6 add:int ::= mul:l '+' add:r [* rv = l + r; *] |
7     mul:l '-' add:r [* rv = l - r; *] |
8     mul:rv .
9
10 mul:int ::= op:l '*' mul:r [* rv = l * r; *] |
11     op:l '/' mul:r [* rv = l / r; *] |
12     op:rv .
13
14 op:int ::= integer:rv |
15     '(' exprs:rv ')' .
16
17 integer:int ::= digit:hi digit:mid digit:lo
18     [* rv = hi*100 + mid*10 + lo; *] |
19     digit:mid digit:lo [* rv = mid*10 + lo; *] |
```

```

20      digit:rv .
21
22 digit:int ::= '0' [* rv = 0; *] |
23     '1' [* rv = 1; *] |
24     '2' [* rv = 2; *] |
25     '3' [* rv = 3; *] |
26     '4' [* rv = 4; *] |
27     '5' [* rv = 5; *] |
28     '6' [* rv = 6; *] |
29     '7' [* rv = 7; *] |
30     '8' [* rv = 8; *] |
31     '9' [* rv = 9; *] .

```

When the generated parser is run on the string $3 \cdot 4; 10; (7*2)+1+$ we get the following output

```

Input: '3 + 4; 10; (7*2)+1 '
Accepted
Oracle: 1 1 1 3 1 3 4 3 3 1 3 5 1 3 3 1 2 2 1 2 1 3 2 2 3 1 1 3
        8 3 1 3 3 3 3 1 3 2

Semantics phase
Result: 7
Result: 10
Result: 14
Result: 15
Final result: 15

```

4.13 Implementing inline semantics

RDSOB parsers explore the grammar in a way that may require tentative matches that are subsequently rejected. Whenever a parser backtracks, some decisions are being unmade.

As a result of this retry behaviour, we cannot simply execute inline semantics during the parse, even though when we design grammars and their semantic actions we tend to think of the action being executed as a side-effect of parsing. Instead, we need to complete the searching associated with parsing and only then run through the grammar the ‘correct’ way to execute the actions. This is the purpose of the oracle: during parsing we construct the oracle as we go, adjusting it as necessary when we backtrack. By the end of the parse we have a map of where the parser *should* have gone. We call the control data structure an oracle because it is as if we had a parser which instead of guessing where to go could simply ask an all-powerful oracle for advice.

To execute the semantics, we use a modified set of parse functions (the *semantics* functions) that (a) contains the embedded semantic actions and (b) look in the

oracle to see where to go rather than searching and backtracking.

Recall the attributed grammar that adds up the 1's in a string:

```

1 S ::= 'b' | 'a' X:result [* System.out.printf(" Result is %d\n", result); *] '@' .
2 X:int ::= '1' X:sum [* rv = sum + 1; *] |
3           '0' X:rv | # .

```

Here are the associated semantics functions.

```

1 void semantics_S() {
2     int result;
3     switch(oracle[co++]) {
4         case 1:
5             match("b");
6             break;
7
8         case 2:
9             match("a");
10            result = semantics_X();
11            System.out.printf("Result is %d%n", result);
12
13            match("@");
14            break;
15        }
16    }
17    int semantics_X() {
18        int rv = 0;
19        int sum;
20        switch(oracle[co++]) {
21            case 1:
22                match("1");
23                sum = semantics_X();
24                rv = sum + 1;
25
26                break;
27
28            case 2:
29                match("0");
30                rv = 0;
31
32                rv = semantics_X();
33                break;
34
35            case 3:
36                /* epsilon */
37                break;

```

```

38 |     }
39 |     return rv;
40 |

```

Note that the semantics functions here are void functions taking no parameters unless we declare a type for their associated nonterminals. Functions with a type T automatically have a local variable `rv` declared of type T which holds the return value; in addition the statement `return rv;` is inserted at the end of the corresponding function.

It is the user's responsibility to ensure that `rv` is loaded with a suitable value. There are two ways to get a value into `rv`: (i) by explicitly assigning to it using a semantic action as in the first alternate of nonterminal `X` and (ii) implicitly assigning to it by naming `rv` as the attribute receiving a synthesized result from a nonterminal, as in the second alternate of nonterminal `X`.

The overall control flow is *via* switch statements selecting on the current oracle index; as each element of the oracle is consumed, the index is incremented by one. The use of switch statements is fast compared to the sequential testing required in the parse functions.

Chapter 5

Reduction semantics

If we want to reason about systems which ultimately execute via an implementation language then we have to reason about the behaviour and correctness of that implementation language. If we want to be able to make provable statements about the correctness and completeness of the languages that we develop, then we must rely the formal correctness of the implementation language and of its own implementation. But how are we to establish formally the correctness of the implementation languages? We have a rather circular problem here.

When we use mathematics we typically try to establish relations and equations between mathematical objects, and then use substitution to propagate those relationships to derived objects. We should like to use those techniques to establish properties of the languages that we are developing, without having to rely on the semantics of some pre-existing programming language. If we can explain the execution of programs using only simple mathematical notions and symbol pushing games, then we have a way of talking about programs that is independent of their implementation on a real computer, that is a *formal semantics*. Even better, if we can find a way to *operationalise* the mathematical description of semantics, then we can in principle automatically generate interpreters from the formal semantics. If the automated construction process is sound, then our generated interpreters will faithfully meet their specification.

Now, in engineering terms formal specifications are still only as good as the person who wrote them: we can still write rubbish and so formal specification does not cure all ills. However, automatic generation certainly reduces the error rate, and the availability of very compact specifications allows us to share our design easily with other experts who can immediately see what we are doing, and can help refine our specification. The ideal situation would be for us to also be able to re-use parts of existing specifications, just as we re-use code in conventional software engineering by accessing libraries and API's.

5.1 The basic idea

The core idea is that we shall model program execution taking the tree form of the program and successively rewriting it into a new tree until no further rewrites are possible, at which point execution halts. We shall specify the kinds of rewrites that may be applied using a set of *rules* and an interpreter which will apply those rules to the tree: the set of rules together make up the *formal semantics* of the language.

For a pure functional language, the rewrites alone will completely model the language. Very few languages are purely functional though: real programs have *side effects* which may be as simple as outputting a sequence of characters or may involve complex manipulation of the contents of memory *via assignments*. Our rules will therefore allow us to specify side-effects that accumulate as the tree is repeatedly rewritten.

An important simplification is that the label of the root of a tree to be rewritten will be used to select which rule to use. If we did not make this restriction, then we would have to search all over a (potentially huge) tree to find putative rewrites, and for languages with side effects we would also need some mechanism for specifying the rewriting order. Our rule, then, will simply be that we must rewrite using a rule for the root node label, and if there is no such appropriate rule then execution will cease even if there are other possible rewrites elsewhere in the tree. This has the twin effects of establishing a rewrite order, and improving efficiency since we do not have to hunt around for rule matches.

We should be careful here though. Although only the root node can be used to select rules, there might still be multiple rules that can be activated, and we then need to consider how to process such specifications. One approach is to exploit this property in modelling *concurrency* and, where multiple rules are active, proceed with all of them at once, each effectively creating a separate thread of control. Alternatively, we may have some mechanism for prioritising the rules, say by checking them in the order that they appear in the specification and taking the first one.

We should also note that, although the root-node-first approach is much more efficient than the allowing rewrites anywhere, it is still a rather slow way to run a program. Depending on your application, it may be fast enough. In later chapters we shall consider program execution via computation of tree attributes which can provide good performance, but which is less tractable when we want to reason about properties of the programming language or of programs themselves. For ultimate performance, both approaches may be used to output *compiled* machine code (or its assembly language equivalent) which is then executed in the conventional way by a real computer.

5.2 Execution via substitution

We now need to formalise our approach into a *game* which can run as an automaton. We shall use ideas from mathematics – relations, equations and substitution – to describe a running computer program.

Program execution as a set of states with transitions between them representing the execution steps of a program. More formally, we developed the idea of a program execution trace as a series of steps that walk a *transition relation* over *configurations*. A configuration represents the state of a computer, and configuration Γ_1 is related to Γ_2 if and only if Γ_2 can appear immediately after Γ_1 in *some* execution of *some* program. Configurations always contain

a program term, and in addition we add entities that represent whatever side effects of program execution we need to record.

By ‘some execution of some program’ we mean any valid program step that you can imagine - it does not need to be useful or sensible, it just needs to be allowed by the language that we are writing a formal semantics for.

5.2.1 Configurations

When modelling a programming language, we begin by deciding on the configurations of that language. A configuration is a tuple of terms comprising at least a program term θ , and possibly including a store term σ which represents the values of program variables, an environment ρ (which holds information about the scope and location of program objects as a map from objects to values), an output list α , an input list β and a set of signals ν which are used to model exception handling. In the first part of this chapter we shall use configurations of the limited form $\langle \theta, \alpha \rangle$ comprising a program term θ and an output term α .

Execution starts with θ being equal to the whole program to be executed. We then pick one small part of it, such as the addition of two constant integers, that we could directly execute, and then rewrite the program to some new term θ_1 , replacing the addition with its result.

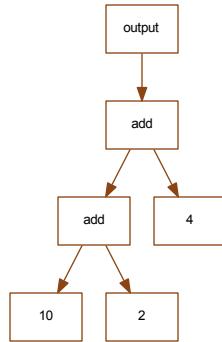
For example, this program fragment when interpreted will eventually result in the value 16 being output.

`1 | output(10+2+4)`

Conventional language compilers would typically convert this into machine-level instructions that add together the 10 and the 2, then add in the 4, and finally output the result. So as to avoid discussing the complexities of infix operators with their priorities and associativities, let us represent the program as a tree, or equivalently as a parenthesized term thus:

`1 | output(add(add(10, 2), 4))`

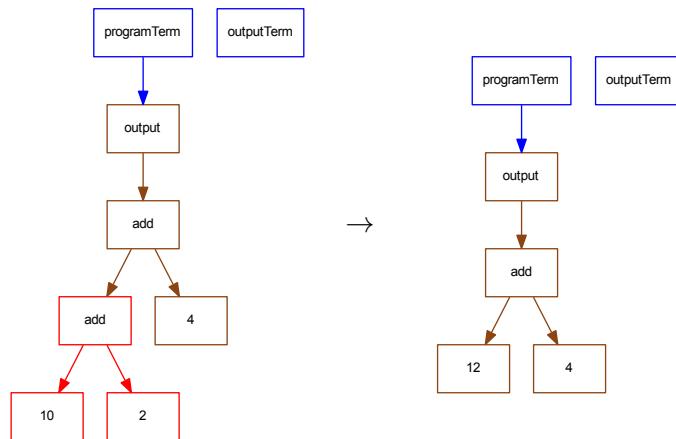
This sort of prefix functional term corresponds directly to a tree if we think of the written term as being the textual trace of a pre-order tree traversal, with parentheses being added to show when we pass down or up a tree edge:



The first computation step for this program reduces the expression to this simpler term

1 `output(add(12, 4))`

which we can show graphically as



We have highlighted the program part that is about to be rewritten in red. We call this part the *reducible expression* or *redex* for short. We also have blue nodes representing the various state components of a running program: in this case limited to the program term and a presently empty term intended to represent the list of outputs made by a program.

The full execution of the program is a sequence of three such steps, which we represent as tuples of the program term and the output.

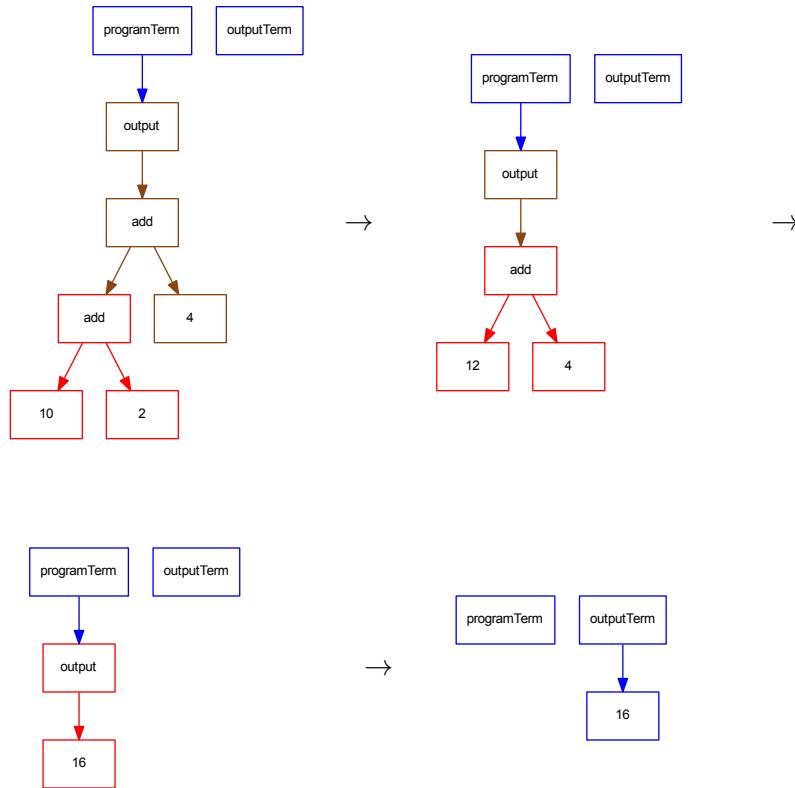
`<output(add(add(10, 2), 4)), []>`

```

⟨output(add(12,4)), [ ]⟩
⟨output(16), [ ]⟩
⟨, [ 16]⟩

```

As before, the subterm to be evaluated is in red, and we record any side effects of the computation in an appropriate semantic entity term. In this case, we have an output term which receives the element 16 as the output statement is reduced. Here is the graphical representation of this sequence of three transitions.



5.3 Avoiding empty terms – the special value `_done`

It turns out to be uncomfortable to have ‘empty’ terms. For instance, the output statement could be just the first component of longer program such as

`1 | output(10+2+4); output(6)`

Here, the ; symbol denotes a *sequencing* operation: a sequence action requires first the left hand side and then the right hand side to be evaluated. In prefix form, we might represent this as

```

1 seq(
2   output(add(add(10, 2),4)),
3   output(6)
4 )

```

and by using nested `seq()` instances we can construct sequences of arbitrary length.

Reducing the first `output()` statement to nothing would leave us with the curious term

```

1 seq(
2   ,
3   output(6)
4 )

```

and to proceed further we should need some notation for representing these sorts of missing or empty terms. To avoid this, we instead invent a special value `done` which represents the completion of a command.

```

1 seq(
2   __done,
3   output(6)
4 )

```

5.4 Term variables are metavariables

A *term variable* is a name which stands for an arbitrary term (tree): it is a sort of metavariable (as opposed to the program variables which are represented by elements of a program term). We shall write term variables in *italics* to distinguish them from actual term elements which we have been writing in sans-serif. In ART specifications, variable names start with a *single* underscore character.

The idea of a term variable is to allow us to speak generally about expressions with arbitrary subexpressions. For example, here is a rule describing how sequences containing the special value `done` may be rewritten.

A term describing the sequence of `__done` and then any other sub-program may be rewritten to just that sub-program. So, for instance, we can rewrite

```

1 seq(__done, output(6))

```

as

```

1 output(6)

```

This rather clumsy piece of English and its example can be more concisely, precisely and generally be expressed as follows.

If X is a term variable, then we can then say that a term of the form `seq(done, X)` can be rewritten as simply X where X is any valid term, or just

$$\text{seq}(_\text{done}, X) \rightarrow X$$

which as an ART specification would be written as

$$1 | \text{--- seq}(_\text{done}, _X) \rightarrow _X$$

We shall make extensive use of term variables as placeholders within trees.

5.5 Pretty-printed rules

Notice how the red machine-friendly version of the above rule only uses characters that are in the ASCII character set and available on a keyboard; the more human friendly typeset black version uses maths characters like \rightarrow and is less verbose. In ART, we always specify things in the red style, but ART can output pretty-printed typeset versions of the rules for use in L^AT_EX documents.

5.6 Pattern matching of terms

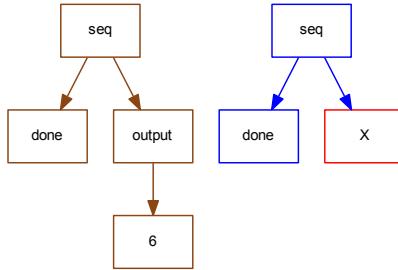
Since we are trying to build a formal game which will execute a program without human intervention, we need some syntax in which to write the rules of the game. As we have already seen, for this reduction of a sequence with `_done` as its left hand argument we write

$$\text{seq}(_\text{done}, X) \rightarrow X$$

This kind of rule is an unconditional rule: anywhere that we find a term that matches the *pattern* `seq(_done, X)` we can directly replace it with whatever the term variable X stands for.

A pattern is a term which may contain term variables. A term which has no term variables in it is called a *closed* term. *Pattern matching* is the process of comparing a closed term to a pattern to decide if they match and if they do, constructing a table of term variables showing what they represent. The relationship between a term variable and its corresponding subterm is called a *binding*; a set of bindings is called an *environment*.

We can represent this process visually as follows:



The sepia coloured closed term is to be matched against the blue pattern. Some of the leaves of the pattern term may be coloured red: these are nodes labeled with term variables. We perform the matching by recursively traversing both trees in tandem. If we arrive at a node which is blue in the pattern but for which the label does not match the label of the corresponding sepia node, then the pattern match fails. If we arrive at a node which is red in the pattern then we have found a term variable-labeled pattern node: we create a binding between that term variable and the corresponding sepia node (which of course represents the entire subtree rooted at that node). Otherwise we descend into the children and continue the recursive traversal.

We can encode this using a recursive function which takes an environment of bindings, a node from the closed term and a node from the pattern as follows:

```

1 match(M: set of term variables, E: environment, t: term, p: term)
2     returns environment OR bottom
3     if label(p) in M then add p |-> t to E
4     else if label(p) != label(t) then return bottom
5     else for ct in children(t), cp in children(p) do add match(E, ct, cp) to E
6     return E

```

Note that our unusual syntax for p in q , r in s do stands for sequential pairwise traversal of the two lists t and p . We initiate a pattern match by calling $\text{match}(M, \{\}, t, p)$ where M is the set of term variables in the pattern, t and p are the root nodes of the term and pattern trees respectively and $\{\}$ is an empty environment.

A pattern term may be arbitrarily deep, but in the version of pattern matching that we shall use term variables will always be the labels of leaf nodes. We shall also restrict ourselves to matching patterns against closed terms. It is easy to imagine more baroque pattern matching operations, but this will be sufficient for our style of semantics specification.

A further important restriction is that a term variable X may only appear at most once within a pattern. Again, one could imagine a version of pattern

matching in which the appearance of two instances of a term variable X meant that they must each match the same subtree, but we shall not allow this.

We shall write

$$\theta \triangleright \pi$$

for the operation of matching closed term θ against pattern π . The result of such a pattern match is either *failure* represented by \perp , or a set of bindings. So

$$\text{seq}(_\!_done, \text{output}(6)) \triangleright \text{seq}(_\!_done, X)$$

returns

$$\{X \mapsto \text{output}(6)\}$$

and

$$\text{seq}(_\!_done, \text{output}(6)) \triangleright \text{seq}(_\!_done, \text{output}(Y))$$

returns

$$\{Y \mapsto 6\}$$

whereas

$$\text{seq}(_\!_done, \text{output}(6)) \triangleright \text{seq}(X, _\!_done)$$

returns

$$\perp$$

because `output(6)` does not match `__done`.

An important special case of pattern matching uses a pattern which is itself a closed term: in such a case, the pattern matcher will return an empty environment if the two terms are identical, or \perp if they differ.

5.7 Pattern substitution

Pattern matching is a way to extract subtrees (subterms) from within closed terms. The bindings will associate term variable names with these subterms, which will themselves be closed (i.e. they will not contain nodes labeled with term variables).

Pattern substitution is the process by which we stitch subterms into a pattern to create a new closed term by substituting the bound subterms for term variables in the pattern

We shall write

$$\pi \triangleleft \rho$$

for the operation of replacing term variables in pattern π with their bound terms from the environment ρ . The result of such a substitution is a closed term; it is an error for π to contain a term variable that is not bound in ρ . So

`plus(X, 10) ↣ {X ↪ 6}`

returns

`plus(6, 10)`

Here is a recursive function to perform substitution

```

1 substitute(M: set of term variables, E: environment, t: term) returns term
2   if label(t) in M then return E.get(label(t)).deepCopy()
3   else {
4     ret = t.shallowCopy()
5     for ct in children(t)
6       t.addChild(substitute(M, E, ct))
7   return ret
8 }
```

5.8 Rules and rule schemas

We now have most of the machinery we need to construct our formal semantics game; we know how to decompose and compose terms (trees). There is one major gap though, and that is the *selection* of sub-phrases to rewrite. Of course, the ordering of these selections is important: for instance we know that $(x - y) - z$ is not in general the same as $x - (y - z)$ so the order in which we formally evaluate the subtractions will affect the final result.

For very simple languages, it might be practical to define their semantics by enumeration. Consider a language which allows a single expression to be output, and limits that expression to a single addition over numbers in the range 0–2. using configurations $\langle \theta, \alpha \rangle$ there are only nine possible programs that can be written in this language each of which we could evaluate directly using these nine rules:

$$\begin{aligned} &\langle \text{output}(\text{plus}(0, 0)), [] \rangle \rightarrow \langle \text{done}, [0] \rangle \\ &\langle \text{output}(\text{plus}(0, 1)), [] \rangle \rightarrow \langle \text{done}, [1] \rangle \\ &\langle \text{output}(\text{plus}(0, 2)), [] \rangle \rightarrow \langle \text{done}, [2] \rangle \\ &\langle \text{output}(\text{plus}(1, 0)), [] \rangle \rightarrow \langle \text{done}, [1] \rangle \\ &\langle \text{output}(\text{plus}(1, 1)), [] \rangle \rightarrow \langle \text{done}, [2] \rangle \\ &\langle \text{output}(\text{plus}(1, 2)), [] \rangle \rightarrow \langle \text{done}, [3] \rangle \\ &\langle \text{output}(\text{plus}(2, 0)), [] \rangle \rightarrow \langle \text{done}, [2] \rangle \\ &\langle \text{output}(\text{plus}(2, 1)), [] \rangle \rightarrow \langle \text{done}, [3] \rangle \end{aligned}$$

$$\langle \text{output}(\text{plus}(2, 2)), [] \rangle \rightarrow \langle \text{__done}, [4] \rangle$$

This is clearly not a very practical approach. What we need to do is to be able to express the pattern of additions more concisely. We call rules that have term variables in them *rule schemas* because they are really a compact way of generating a (possibly infinite) set of real rules.

In pseudo code, we might say something like this:

```

1 let x, y and z be term variables
2 if program term theta matches output(add(x,y)) with some alpha and
3   x is bound to an integer in the range 0–2 and
4   y is bound to an integer in the range 0–2 and
5   z is bound to the result of adding x and y together then
6     rewrite theta to __done and alpha to the substitution of z

```

More formally, using the notations we have developed we might say

```

if <math>\rho_1 = (\theta \triangleright \text{output}(\text{plusOp}(X, Y))), \alpha\rangle
and <math>\text{is012}(X) \triangleright \text{true}
and <math>\text{is012}(Y) \triangleright \text{true}
and <math>\rho_2 = ((\text{__add}(X, Y) \triangleleft \rho_1) \triangleright Z)
then <math>\theta \rightarrow \theta' = \langle \text{__done}, [Z \triangleleft \rho_2] \rangle

```

We have introduced a new mechanism here: simple functions that take terms and return terms which we write in `teletype font`. We can think of these as pre-existing mathematical functions whose definition is obvious, or if we are writing an interpreter then we might think of these as lookup tables, or calls to very small programs that compute results. The important thing is that these functions must be so small as to allow us to trivially check their correctness.

In this case we are using two new functions: `is012(x)` which returns a term *true* or a term *false* depending on whether *x* is in the set {0, 1, 2} or not; and `__add(x, y)` which returns a term labeled with the number formed from the addition of terms *x* and *y*. Notice how everything we are doing reduces to operations over terms: our functions are not returning values such as *true* or *false*; they are returning trees made up of a single node *labeled* with *true* or *false*. Notice also that we are using names such as `addOp` for the function that computes the operation of addition, as opposed to the term constructor `add` which is a tree label from a program term. You should think of `add` as the piece of syntax that requests an addition, and `addOp` as the name of the function (or perhaps even machine instruction) which will actually perform the addition. We shall follow this convention throughout: names with the *Op* suffix are used

for functions that perform computations, and they must not also appear as the names of a term (tree) element.

Even our formal version of the rule schema is rather a lot of writing. It will turn out that usually several of these rule schemas will be used together in a way that corresponds to inference in a logical system, and we use a special form of syntax that allows us to show derivations in that logical system as trees of rule schemas. A general inference rule looks like this:

$$\frac{C_1 \quad C_2 \quad \dots \quad C_n}{\langle \theta, \alpha \rangle \rightarrow \langle \theta', \alpha' \rangle}$$

The elements above the line (the C_i) are called *conditions*. Conditions can themselves be transitions although we have not yet encountered examples of that style. Conditions may also be simple matches against the return value of a function, in which case they are called *side conditions*. The single transition below the line is called the *conclusion*. You might read an inference rule in this style as:

if you have a configuration $\langle \theta, \alpha \rangle$,
 and C_1 succeeds and C_2 succeeds and ... and C_n succeeds
 then transition to configuration $\langle \theta', \alpha' \rangle$

so one reads this kind of rule by checking that the current configuration matches the left hand side of the conclusion, then by checking the conditions, and if everything succeeds rewriting the current configuration into the right hand side of the conclusion. We sometimes refer to this rather operational view of logical inference as ‘reading round the clock’.

The inference rule representation of our schema is

$$\frac{(\text{is012}(X) \triangleleft \rho_1) \triangleright \text{true} \quad (\text{is012}(Y) \triangleleft \rho_1) \triangleright \text{true} \quad \rho_2 = ((\text{addOp}(X, Y) \triangleleft \rho_1) \triangleright Z)}{\langle \rho_1 = (\theta \triangleright \text{output}(\text{plusOp}(X, Y)), \alpha \triangleright []) \rangle \rightarrow \langle \text{done}, [\alpha, Z \triangleleft \rho_2] \rangle}$$

We have been careful here to represent all of the pattern matching and substitution operations explicitly. In practice, it is understood that (i) each rule has its own set of term variables even if the same term variable name is used in multiple rules (that is, there is no communication of bindings from one rule to the next) and that (ii) as a rule is checked, a private environment is developed as we go round the clock, (iii) the first time we meet a term variable it is being used to create a binding, (iv) subsequent appearances of a term variable are to be substituted by its binding and (v) that the term in the left hand side of the conclusion is a pattern to be matched against θ in the current configuration.

This allows us to abbreviate our inference rule to:

$$\frac{\text{is012}(X) \triangleright \text{true} \quad \text{is012}(Y) \triangleright \text{true} \quad \text{addOp}(X, Y) \triangleright Z}{\langle \text{output}(\text{plusOp}(X, Y)), \alpha \triangleright \langle \text{done}, [\alpha, Z] \rangle \rangle}$$

and this is the style that we shall use in future.

5.9 The interpreting function F_{SOS}

Now that we have pattern matching and substitution operations along with notions of transitions and side conditions, we can think about a function which takes an input term and *interprets* it by looking through the rules for possible transitions.

5.9.1 Managing the local environment

When implementing F_{SOS} using a procedural language with assignment, there is a useful optimisation that we can apply. Recall that when we wrote out the full version of the inference rule we were careful to create new environments ρ_1, ρ_2, \dots each time we performed a pattern match. As we moved from the detailed version of a rule schema to the abbreviate form we noted that:

(iii) the first time we meet a term variable it is being used to create a binding

and

(iv) subsequent appearances of a term variable are to be substituted by its binding

As a result term variables will never be reused (that is a binding cannot subsequently be changed) and we can use a single environment to which bindings are added as we go round the clock. We shall call this mutable environment E .

5.9.2 Procedural pseudo-code for F_{SOS}

This is a procedural implementation of the rule application function F_{SOS} . The function takes a configuration made up of a program term and zero or more semantic entities and either returns a new configuration or \perp . It accesses a set of rule schemes R each of which has a conclusion and a set of conditions. In the pseudo code, we use the operators $|>$ and $<|$ for the pattern matching \triangleright and substitution \triangleleft operations.

```

1 let R be the set of rule schemas
2
3 FSOS(C: configuration) returns configuration OR bottom
4   for r in R
5     if C |> r.conclusion.lhs then
6       let E be an empty set of bindings
7       for c in r.conditions
8         if isSideCondition(c)
9           let res be (c.lhs <| E) |> c.rhs

```

```

10   if res = bottom then next r else add res to E
11   else
12     let T be c.lhs <| E
13     if isvalue(T) then return T
14     let res be Fsos(T) |> c.rhs
15     if res = bottom then next r else add res to E
16     return r.conclusion.rhs <| E
17   return bottom

```

The basic approach is just as we described in our ‘round-the-clock’ informal description of how to read an inference rule.

We start with a configuration, perhaps the initial program and an empty output list. We then scan through all of the rule schemas until we find one that matches the root node of our term. (As an aside, we can make this process more efficient by storing R as a map from constructor label L to subsets of R that have L as the root constructor of their conclusion’s left hand side; we have not used this optimisation here.)

We then create a new, empty environment called E and work our way across the conditions evaluating them; if they succeed we add their bindings into E , but if any fail we abort the processing for this rule and throw away E , seeking another rule whose conclusion left hand side matches our term.

Conditions can be either side conditions or transitions.

- ◊ For a side condition we call the left hand side function and then pattern match the result against the condition’s right hand side.
- ◊ If the condition is a transition, we first let T be the condition’s left hand side after substitution; if T is a value we return that, otherwise we recursively call $F_{SOS}(T)$ and pattern match the result against the condition’s right hand side.

There is an important technical detail here: a call to $F_{SOS}()$ may result in \perp but in line 12 we pattern match the result of the call. Now, the pattern match operator \triangleright is usually only defined over terms, but here we extend the definition so that an attempt to pattern match against \perp will yield \perp , and in that way the failure propagate up.

If all of the conditions succeed, then at line 14 F_{SOS} returns the right hand side of the conclusion after substitution against the final value of E .

If we exhaust the rules set R then we have arrived at a terminal configuration, that is one from which no further transitions may be made. For a correct program, this terminal transition will correspond to the program’s final version. If the rules were ill-formed, it would be possible to run out of applicable transitions prematurely: such a configuration is called a *stuck configuration* and

would be reported as an error by the interpreter which requires the rules to be changed. In detail, we nominate some program terms as *values*. If an evaluation terminates with a value term, then we have a normal execution. If an evaluation terminate with a non-value term, then we have a stuck execution. The `_done` constructor in our rule is an example of a value: it marks normal termination of commands. Numeric and other literals such as strings are also usually values; and we may indicate the successful completion of expression evaluations by reducing them to one of these values.

5.9.3 Program term rewrites - the outer interpreter

The `Fsos()` function only performs a single transition on the program term, so we usually need to wrap the initial call in an `interpret()` function which repeatedly applies `Fsos()` until we arrive at a terminal configuration; flagging an error if that configuration is not a value.

```

1 interpret(C: configuration) returns configuration
2   C' = Fsos(C)
3
4   while C' not bottom
5     C = C'
6     C' = Fsos(C)
7
8   if not isValue(C.theta) error('Stuck configuration')
9
10  return C

```

5.10 Structural Operational Semantics and execution traces

It is clear that in some sense the structure of the term to be executed specifies the execution order.

Using a technique due to Plotkin called *Structural Operational Semantics*, we shall specify similar inside-out steps often using repeated term traversals. The basic idea is to conditionally rewrite the term by isolating a subcomputation that can be performed immediately, and we ensure that the rewrites are done in the inside-out order by building inference rules that have the abstract syntax of our language embedded in the conclusions.

We shall now look in detail at some rules, and use ART's trace facility to see exactly what happens as the rewriter attempts to reduce a program to a value.

5.10.1 SOS rules for a subtraction language with output

As a first example, let us generalise the 0, 1, 2 addition language of the previous section to allow expressions involving arbitrarily nested additions of 32-bit integers with output.

We begin with a rule that performs addition for expressions such as $3 + 4$, that is where each operand is a single integer. The rule uses our style of abstract syntax, which would encode this example as `add(3, 4)`.

It is convenient to name rules for reference purposes by giving a label at the start of the rule comprising a `-` sign and an alphanumeric name. These names have no meaning in themselves but we often use names that indicate the purpose of the rule. If we do not use a label then ART generates a name of the form `Rn` when `n` is a rule number. The ART rewriter `!trace` facility will announce rules as they are used, and this is very helpful when debugging.

Configurations

We begin by deciding on a configuration for our language. Since we are only adding literal numbers we do not have program variables, and thus do not need a store σ or an environment ρ to act as a symbol table. However we do want to model the output, so we need a an output list which we shall call α .

In some specifications it is useful to have more than one relation, represented by different kinds of arrows each of which has its own configuration so we need to specify which relation we mean when we set up a configuration. When we specify a configuration for a particular relation, we write something of the form `!configuration -> x:y, ...` where `x:y` specifies a semantic entity with `x` is the root name of the variables we shall use to manipulate that entity and `y` is its type. This is a suitable directive for our language:

```

1 !configuration -> _alpha:_list
2 !trace 5

```

Here we have also set the trace level to 5 so as to show the variable bindings developed during a rule test.

Adding literals

Now we add a rule for addition which can only handle two literal numbers, and test it using a `!try`. Notice how the `!try` and both sides of the transition have exactly the same shape as the configuration: a program term followed by a literal empty `_list` for the `!try` and a program term followed by a variable for the rule.

```

1  -add
2  _n1 |> _int32(_) _n2 |> _int32(_)
3  ---
4  add(_n1, _n2), _alpha -> __add(_n1,_n2), _alpha
5
6 !try add(5,3), __list

```

The match expressions above the line (`_n1 |> _int32(_)` `_n2 |> _int32(_)`) are checking that both operands really are of type `_int32`. Without them, we might end up calling the `_add` function with invalid arguments, and that would give a fatal error that would stop the rewriter.

When we run this example at trace level 5 we get a blow-by-blow description of the rewriter's actions as it works its way round the rule, and eventually calls the `_add` Value function:

```

1 Step 1
2 Rewrite attempt 1: <add(5, 3), []> ->
3 -add _n1 |> _ _n2 |> _ --- <add(_n1, _n2), _alpha>-><__add(_n1, _n2), _alpha>
4 -add bindings after Theta match { _n1=5, _n2=3, _alpha=[] }
5 -add premise 1 _n1 |> _
6 -add bindings after premise 1 { _n1=5, _n2=3, _alpha=[] }
7 -add premise 2 _n2 |> _
8 -add bindings after premise 2 { _n1=5, _n2=3, _alpha=[] }
9 -add rewrites to <8, []>
10 Normal termination on <8, []> after 1 step and 1 rewrite attempt

```

Nested expressions

In general, we would like to handle nested expressions such as `add(3, add(4,5))`. However, our simple `-add` rule will not rewrite such expressions:

```

1 !try add(3, add(4,5)), __list

```

yields

```

1 Step 1
2 Rewrite attempt 1: <add(3, add(4, 5)), []> ->
3 -add _n1 |> _ _n2 |> _ --- <add(_n1, _n2), _alpha>-><__add(_n1, _n2), _alpha>
4 -add bindings after Theta match { _n1=3, _n2=add(4, 5), _alpha=[] }
5 -add premise 1 _n1 |> _
6 -add bindings after premise 1 { _n1=3, _n2=add(4, 5), _alpha=[] }
7 -add premise 2 _n2 |> _
8 -add premise 2 failed: seek another rule
9 Stuck on <add(3, add(4, 5)), []> after 1 step and 1 rewrite attempt

```

At line 5, we test the first *premise* (that is the leftmost match or transition expression). That succeeds, since `_n1` is bound to `_int32(3)` and that *does*

match `_int32(_)`. However, at line 7 we test the second premise and that fails because `_n2` is bound to `add(4, 5)` and that does *not* match `_int32(_)`. As a result, the whole rule fails. There are no other rules, so the rewriter reports that it is stuck.

What we need are *resolution rules* which can take the inner expression `add(3,4)` and rewrite it to a simpler form. Here is a suitable rule which will recursively call the rewriter on the right operand instead of just trying to match it against a type:

```

1 -addResolveRight
2 _n |> _int32(_) _E, _alpha -> _EP, _alphaP
3 ---
4 add(_n, _E), _alpha -> add(_n,_EP), _alphaP
```

Note how the first premise `_n |> _int32(_)` simply requires the left argument to be a 32-bit integer. The second premise `_E, _alpha -> _EP, _alphaP` requires the rewriter to take the second argument `_E` and separately *transition* it to some new `_EP` (expression primed). Note also that both sides of the transition have the same shape as the configuration: the original output list `_alpha` appears on the left hand side of the transition, but the resulting list is called `_alphaP`, and it is `_alphaP` that is used on the right hand side of the conclusion. We do this to ensure that any updates to the list that occur as side effects of the resolution are propagated to the final rewrite.

When we run `!try add(3, add(4,5)), _list` with rules `-add` and `-addResolveRight` we get quite a long trace. It is very important to work through this trace and ensure that everything makes sense: this is the first time we have seen the ‘transition above the line’ idiom: it is perhaps the most important idea in this chapter.

```

1 Step 1
2 Rewrite attempt 1: <add(3, add(4, 5)), []> ->
3 -add _1 |> _ _2 |> _ --- <add(_1, _2), _3>-><_add(_1, _2), _3>
4 -add bindings after Theta match { _1=3, _2=add(4, 5), _3=[] }
5 -add premise 1 _1 |> _
6 -add bindings after premise 1 { _1=3, _2=add(4, 5), _3=[] }
7 -add premise 2 _2 |> _
8 -add premise 2 failed: seek another rule
9 -addResolveRight _n |> _ <_E, _alpha>-><_EP, _alphaP> ---
10           <add(_n, _E), _alpha>-><add(_n, _EP), _alphaP>
11 -addResolveRight bindings after Theta match { _n=3, _E=add(4, 5), _alpha=[] }
12 -addResolveRight premise 1 _n |> _
13 -addResolveRight bindings after premise 1 { _n=3, _E=add(4, 5), _alpha=[] }
14 -addResolveRight premise 2 <_E, _alpha>-><_EP, _alphaP>
15   Rewrite attempt 2: <add(4, 5), []> ->
16   -add _1 |> _ _2 |> _ --- <add(_1, _2), _3>-><_add(_1, _2), _3>
17   -add bindings after Theta match { _1=4, _2=5, _3=[] }
18   -add premise 1 _1 |> _
```

```

19  -add bindings after premise 1 { _1=4, _2=5, _3=[] }
20  -add premise 2 _2 |> _
21  -add bindings after premise 2 { _1=4, _2=5, _3=[] }
22  -add rewrites to <9, []>
23  -addResolveRight bindings after premise 2 { _1=3, _2=add(4, 5), _3=[], _4=9, _5=[] }
24  -addResolveRight rewrites to <add(3, 9), []>
25 Step 2
26 Rewrite attempt 3: <add(3, 9), []> ->
27 -add _1 |> _2 |> _ --- <add(_1, _2), _3>-><_add(_1, _2), _3>
28 -add bindings after Theta match { _1=3, _2=9, _3=[] }
29 -add premise 1 _1 |> _
30 -add bindings after premise 1 { _1=3, _2=9, _3=[] }
31 -add premise 2 _2 |> _
32 -add bindings after premise 2 { _1=3, _2=9, _3=[] }
33 -add rewrites to <12, []>
34 Normal termination on <12, []> after 2 steps and 3 rewrite attempts

```

This trace breaks down into two steps. In Step 1 we rewrite the program term `add(3, add(4, 5))` to `add(3, 9)` by performing an inner rewrite on `add(4, 5)` to `9`, and in the second step we rewrite `add(3, 9)` to `12`. The second step is essentially the same as that for our earlier example on page 91 so we do not need to go through that again.

The first step is more interesting. In lines 2–8 we see the rewriter trying the `-add` rule first, and rejecting it because premise 2 fails. The rewriter then moves to the `-addResolveRight` rule. At lines 12–13, premise 1 is checked successfully. Line 15–22 show the rewriter recursing—at line 15 the rewriter announces that it is starting rewrite attempt 2, and lines 15–22 are indented to show that this is an *inner* rewrite. We see that the variable `E` has effectively chopped the subexpression `add(4, 5)` out of the original term, and the rewriter can now process that in isolation, starting with the `-add` rule which will accept this subexpression since it only has integer literal arguments. Once it has rewritten to 9, we can restart the outer rewrite attempt.

This example shows that our rewrite rules form an **inductive specification, which we can process in a natural way using recursion. We shall use this idea over and over again.**

The standard triad of rules for arity-2 (diadic) operations

Now, we are not quite finished with our rules for addition! We need a third rule which will handle compound expressions in the *left* argument. Here is the complete set of three rules that we need:

```

1  --add
2  _n1 |> _int32(_) _n2 |> _int32(_)
3  ---
4  add(_n1, _n2), _alpha -> _add(_n1,_n2), _alpha
5
6  --addResolveRight
7  _n |> _int32(_) _E, _alpha -> _EP, _alphaP
8  ---
9  add(_n, _E), _alpha -> add(_n,_EP), _alphaP
10
11 --addResolveLeft
12 _E1, _alpha -> _E1P, _alphaP
13 ---
14 add(_E1, _E2), _alpha -> add(_E1P,_E2), _alphaP

```

The new `--addResolveLeft` rule uses the same mechanism of a transition above the line to resolve the left argument. The right hand argument is simply carried across to the rewritten term. Note that only the first rule `--add` actually calls the addition function `_add`. The other two rules both rewrite `add(,)` terms to `add(,)` terms.

Here is a more compact trace (at level 3) showing all three rules working together to process `!try add(add(2,3), add(4,5)), __list`

```

1 Step 1
2 Rewrite attempt 1: <add(add(2, 3), add(4, 5)), []> ->
3   -add _n1 |> _ _n2 |> _ --- <add(_n1, _n2), _alpha>-><_add(_n1, _n2), _alpha>
4   -addResolveRight _n |> _ <_E, _alpha>-><_EP, _alphaP> ---
5     <add(_n, _E), _alpha>-><add(_n, _EP), _alphaP>
6   -addResolveLeft <_E1, _alpha>-><_E1P, _alphaP> ---
7     <add(_E1, _E2), _alpha>-><add(_E1P, _E2), _alphaP>
8   Rewrite attempt 2: <add(2, 3), []> ->
9     -add _n1 |> _ _n2 |> _ --- <add(_n1, _n2), _alpha>-><_add(_n1, _n2), _alpha>
10    -add rewrites to <5, []>
11   -addResolveLeft rewrites to <add(5, add(4, 5)), []>
12 Step 2
13 Rewrite attempt 3: <add(5, add(4, 5)), []> ->
14   -add _n1 |> _ _n2 |> _ --- <add(_n1, _n2), _alpha>-><_add(_n1, _n2), _alpha>
15   -addResolveRight _n |> _ <_E, _alpha>-><_EP, _alphaP> ---
16     <add(_n, _E), _alpha>-><add(_n, _EP), _alphaP>
17   Rewrite attempt 4: <add(4, 5), []> ->
18     -add _n1 |> _ _n2 |> _ --- <add(_n1, _n2), _alpha>-><_add(_n1, _n2), _alpha>
19     -add rewrites to <9, []>

```

```

20  -addResolveRight rewrites to <add(5, 9), []>
21 Step 3
22  Rewrite attempt 5: <add(5, 9), []> ->
23  -add _n1 |> _ _n2 |> _ --- <add(_n1, _n2), _alpha>-><__add(_n1, _n2), _alpha>
24  -add rewrites to <14, []>
25 Normal termination on <14, []> after 3 steps and 5 rewrite attempts

```

Step 1 rewrites `add(2,3)` to `5`, step 2 rewrites `add(4,5)` to `9` and step 3 rewrites `add(5,9)` to `14`.

A triad of rules in this particular order is the standard way in ART to manage arity-2 operations such as arithmetic, logic and relation functions.

It is important that the rules are listed in order from the most specific (rule `-add` which demands literal integer operands) to the most general. If we put `-addResolveLeft` first, then the rewriter will use it recurse down into the left hand operands until it finds a terminal, then rewrite to the same terminal. The recursion unwinds with each transition simply replicating its left hand side until we get back to the top, at which point the rewriter notices that nothing has changed and announces that it is stuck.

Handling output

The rules above make provision for an output list called α by having a configuration comprising a program term and a list, and by ensuring that any updates to the list are propagated across transitions. However, we have not yet done any output; here is the only rule we need to add output to our language:

```

1 -output
2 _E, _alpha -> _EP, _alphaP
3 ---
4 output(_E), _alpha -> __done, __put(_alphaP, 0, _EP)

```

Rule `-output` resolves its single argument in the same way as the addition resolution rules, by recursively reducing the argument to a value `_EP`. We also keep track of changes to the output with `_alphaP` just in case evaluating our argument itself makes changes to the output.

```

1 Step 1
2 Rewrite attempt 1: <output(add(3, 4)), []> ->
3 -output <_E, _alpha>-><_EP, _alphaP> --- <output(_E), _alpha>-><__done, __put(_alphaP, 0, _EP)>
4  Rewrite attempt 2: <add(3, 4), []> ->
5  -add _n1 |> _ _n2 |> _ --- <add(_n1, _n2), _alpha>-><__add(_n1, _n2), _alpha>
6  -add rewrites to <7, []>
7  -output rewrites to <__done, [7]>
8 Normal termination on <__done, [7]> after 1 step and 2 rewrite attempts

```

Sequencing

We would like to be able to execute a sequence of output statements so as to build an output list: at the moment we can only reduce a single `output()` and so our output list can never be more than one element long.

Now, some programming languages are *expression based* in which every phrase returns a value, but most (including Java and C) distinguish between *statements* and *procedures* which do not return a value and *expressions* and *functions* which do.

Our internal syntax is inherently expression based, but we use the special value `_done` to represent an expression which is really a statement, just as we use the pseudo-type `void` in Java to show that a method is a procedure that returns nothing rather than a function that returns a value.

We represent sequencing using the arity-2 constructor `seq`, so that a sequence of two `output` command would be written as `seq(output(...), output(...))` and a sequence of three outputs as `seq(output(...), seq(output(...), output(...)))`. The `seq` expressions nest in the same was as arithmetic expressions, and a one arity-2 constructor is enough to describe sequences of any length.

The whole point of a sequence is to perform the leftmost element before the rightmost. We impose this ordering in the same way that we controlled the evaluation order of the `add` arguments. The idea is that we have a resolution rule `-sequence` that reduces its left argument to `_done` and base rule `-sequenceDone` that only matches sequences that have already been reduced, and which rewrites itself to its own right argument:

```

1  -sequenceDone
2  -----
3  seq(_done, _C), _alpha -> _C, _alpha
4
5  -sequence
6  _C1, _alpha -> _C1P, _alphaP
7  -----
8  seq(_C1, _C2), _alpha -> seq(_C1P, _C2), _alphaP

```

Here is result of a level 2 trace using these rules (and the output rule) to reduce a sequence of two `output` statements:

```

1 !try seq(output(3),output(4),_list

```

```

1 Step 1
2   -sequence rewrites to <seq(_done, output(4)), [3]>
3 Step 2
4   -sequenceDone rewrites to <output(4), [3]>
5 Step 3

```

```

6 |   -output rewrites to <__done, [4, 3]>
7 | Normal termination on <__done, [4, 3]> after 3 steps and 6 rewrite attempts

```

We can see that the first step rewrites the left argument to `_done`, updating the output to `[3]` as a side effect. The second step then rewrites away the `_done` and the third step puts `4` onto the output list.

Here is a more detailed level 3 trace of the same `!try`:

```

1 Step 1
2 Rewrite attempt 1: <seq(output(3), output(4)), []> ->
3 -sequenceDone --- <seq(__done, _C), _alpha>-><_C, _alpha>
4 -sequenceDone Theta match failed: seek another rule
5 -sequence <_C1, _alpha>-><_C1P, _alphaP> --- <seq(_C1, _C2), _alpha>-><seq(_C1P, _C2), _alphaP>
6   Rewrite attempt 2: <output(3), []> ->
7     -output <_E, _alpha>-><_EP, _alphaP> --- <output(_E), _alpha>-><__done, __put(_alphaP, 0, _EP)>
8       Rewrite attempt 3: <3, []> ->
9         Terminal <3, []
10        -output rewrites to <__done, [3]>
11      -sequence rewrites to <seq(__done, output(4)), [3]>
12 Step 2
13   Rewrite attempt 4: <seq(__done, output(4)), [3]> ->
14     -sequenceDone --- <seq(__done, _C), _alpha>-><_C, _alpha>
15     -sequenceDone rewrites to <output(4), [3]>
16 Step 3
17   Rewrite attempt 5: <output(4), [3]> ->
18     -output <_E, _alpha>-><_EP, _alphaP> --- <output(_E), _alpha>-><__done, __put(_alphaP, 0, _EP)>
19       Rewrite attempt 6: <4, [3]> ->
20         Terminal <4, [3]>
21       -output rewrites to <__done, [4, 3]>
22 Normal termination on <__done, [4, 3]> after 3 steps and 6 rewrite attempts

```

5.10.2 Building rules for a full language

We now switch to building rules for a language that has assignment, but no output. We will have a single scope region, so we do not need separate symbol tables (represented by environments ρ) and can thus make do with just a store, so the configuration is:

```

1 !configuration -> _sig:_map

```

Assignment and dereferencing

Assignment is the process of creating (or updating) a binding in the store between an identifier and a value. Dereferencing is the process of replacing an identifier with its bound value in the store. We use the `_map` operations `_put` and `_get` to manipulate `_sig`.

```

1  --assign
2  _n |> _int32(_)
3  ---
4  assign(_X,_n),_sig -> _done,_put(_sig,_X,_n)
5
6  --assignR
7  _E,_sig -> _l,_sigP
8  ---
9  assign(_X,_E),_sig -> assign(_X,_l),_sigP
10
11 --deref
12 ---
13 deref(_R),_sig -> _get(_sig,_R),_sig

```

Control flow

For the GCD language, we need sequencing, conditionals and a basic loop:

```

1  --sequenceDone
2  --- seq(_done, _C),_sig -> _C,_sig
3
4  --sequence
5  _C1,_sig -> _C1P,_sigP
6  ---
7  seq(_C1,_C2),_sig -> seq(_C1P,_C2),_sigP
8
9  --ifTrue
10 --- if(true,_C1,_C2),_sig -> _C1,_sig
11
12 --ifFalse
13 --- if(false,_C1,_C2),_sig -> _C2,_sig
14
15 --ifResolve
16 _E,_sig -> _EP,_sigP
17 ---
18 if(_E,_C1,_C2),_sig -> if(_EP,_C1,_C2),_sigP
19
20 --while
21 ---
22 while(_E,_C),_sig -> if(_E,seq(_C,while(_E,_C)),_done),_sig

```

Note how the sequencing rules are deliberately slightly different to those on page pagerefseq: now we have a store `_sig` instead of an output list `alpha` and the labels and variable names (which have no semantic significance) are modified.

Arithmetic and relational operations

Our example only required greater-than, not-equal and subtraction so that is all we have implemented here, using the standard triad of rules. We show the `gt` rules in vertically-spaced style, and then a more compressed form for the other operations.

```

1  -gt
2  _n1 |> _int32(_) _n2 |> _int32(_)
3  -----
4  gt(_n1,_n2),_sig -> _gt(_n1,_n2),_sig
5
6  -gtR
7  _n |> _int32(_) _E2,_sig -> _l2,_sigP
8  -----
9  gt(_n,_E2),_sig -> gt(_n,_l2),_sigP
10
11 -gtL
12 _E1,_sig -> _l1,_sigP
13 -----
14 gt(_E1,_E2),_sig -> gt(_l1,_E2),_sigP
15
16 -ne _n1 |> _int32(_) _n2 |> _int32(_) ---- ne(_n1,_n2),_sig -> _ne(_n1,_n2),_sig
17 -neR _n |> _int32(_) _E2,_sig -> _l2,_sigP ---- ne(_n,_E2),_sig -> ne(_n,_l2),_sigP
18 -neL _E1,_sig -> _l1,_sigP ---- ne(_E1,_E2),_sig -> ne(_l1,_E2),_sigP
19
20 -sub _n1 |> _int32(_) _n2 |> _int32(_) ---- sub(_n1,_n2),_sig -> _sub(_n1,_n2),_sig
21 -subR _n |> _int32(_) _E2,_sig -> _l2,_sigP ---- sub(_n,_E2),_sig -> sub(_n,_l2),_sigP
22 -subL _E1,_sig -> _l1,_sigP ---- sub(_E1,_E2),_sig -> sub(_l1,_E2),_sigP

```

5.11 An eSOS specification for the GCD language

Finally, here is a complete set of rules for the GCD language, covering only those features required for the example on page 15, written in a compressed style.

```

1  -assign _n |> _int32(_) ---- assign(_X,_n),_sig -> _done,_put(_sig,_X,_n)
2  -assignR _E,_sig -> _l,_sigP ---- assign(_X,_E),_sig -> assign(_X,_l),_sigP
3
4  -deref ---- deref(_R),_sig -> _get(_sig,_R),_sig
5
6  -sequenceDone ---- seq(_done, _C),_sig -> _C,_sig
7  -sequence _C1,_sig -> _C1P,_sigP ---- seq(_C1,_C2),_sig -> seq(_C1P,_C2),_sigP
8
9  -ifTrue ---- if(true,_C1,_C2),_sig -> _C1,_sig
10 -ifFalse ---- if(false,_C1,_C2),_sig -> _C2,_sig

```

```

11 -ifResolve _E,.sig ->_EP,.sigP --- if(_E,_C1,_C2),.sig -> if(_EP,_C1,_C2),.sigP
12 -while --- while(_E,_C),.sig -> if(_E,seq(_C,while(_E,_C)),_done),.sig
13
14 -gt _n1 |> _int32(.) _n2 |> _int32(.) --- gt(_n1,_n2),.sig -> _gt(_n1,_n2),.sig
15 -gtR _n |> _int32(.) _E2,.sig -> _I2,.sigP --- gt(_n,_E2),.sig -> gt(_n,_I2),.sigP
16 -gtL _E1,.sig -> _I1,.sigP --- gt(_E1,_E2),.sig -> gt(_I1,_E2),.sigP
17
18 -ne _n1 |> _int32(.) _n2 |> _int32(.) --- ne(_n1,_n2),.sig -> _ne(_n1,_n2),.sig
19 -neR _n |> _int32(.) _E2,.sig -> _I2,.sigP --- ne(_n,_E2),.sig -> ne(_n,_I2),.sigP
20 -neL _E1,.sig -> _I1,.sigP --- ne(_E1,_E2),.sig -> ne(_I1,_E2),.sigP
21
22 -sub _n1 |> _int32(.) _n2 |> _int32(.) --- sub(_n1,_n2),.sig -> _sub(_n1,_n2),.sig
23 -subR _n |> _int32(.) _E2,.sig -> _I2,.sigP --- sub(_n,_E2),.sig -> sub(_n,_I2),.sigP
24 -subL _E1,.sig -> _I1,.sigP --- sub(_E1,_E2),.sig -> sub(_I1,_E2),.sigP
25

```

Here is a GIFT-annotated grammar which translates from our human-friendly syntax to terms suitable for use with these rules.

```

1 seq ::= statement^ ^
2 statement ::= assign^ ^ | while^ ^ | if^ ^
3 assign ::= &ID `:=` expression `;`^
4 while ::= `while`^ expression `do`^ statement
5 if ::= `if`^ expression `then`^ statement | `if`^ expression `then`^ statement `else`^ statement
6 expression ::= rels^ ^
7 rels ::= adds^ ^ | gt^ ^ | ne^ ^
8   gt ::= adds `>`^ adds
9   ne ::= adds `!=`^ adds
10  adds ::= operand^ ^ | sub^ ^ | add^ ^
11    add ::= adds `+`^ operand
12    sub ::= adds `−`^ operand
13  operand ::= _int32^ ^ | deref^ ^
14 _int32 ::= &INTEGER
15 deref ::= &ID

```

And here are two `!try`s that firstly exercise the rules using an explicit term, and secondly run a human-friendly string through the parser with the resulting term passed to the rewriter.

```

1 !try seq(assign(a,6), seq(assign(b,9), while(ne(deref(a), deref(b)),
2   if(gt(deref(a), deref(b)), assign(a, sub(deref(a), deref(b))), assign(b, sub(deref(b), deref(a))))))),_
3   __map
4 !try "a := 6; b := 9; while a != b do if a > b then a := a - b; else b := b - a;"
```

The two tries give identical results: in fact the term above was generated with the parser.

1 Normal termination on <__done, {a=3, b=3}> after 37 steps and 90 rewrite attempts
2 Normal termination on <__done, {a=3, b=3}> after 37 steps and 90 rewrite attempts

Chapter 6

Attribute interpreters

As we have seen, an SOS interpretation of a program starts with a derivation tree and progressively rewrites it until only a normal form remains whilst accumulating side effects in subsidiary semantic entities. This approach has the great benefit of compartmentalising the meanings of phrases in a way that naturally supports hierarchical composition of semantics, reflecting the nesting principle that is so widely used in programming languages. The nesting is primarily expressed in the hierarchy of string rewrite rules that forms the grammar; each rewrite rule only needs to express the meaning of an isolated syntactic phrase.

An *attribute interpreter* is an alternative way of developing a compositional semantics from a derivation tree. Instead of having a single current configuration which is rewritten at each step we associate data elements called *attributes* with each node of the derivation tree, and then write equations or small code fragments called *actions*. The derivation tree itself does not change during an attribute-based interpretation, and this means that an attribute interpreter may be much faster than a reduction-style interpreter.

attributes

actions

The general ideas behind attribute evaluators can be seen in several early compilers which implement semantics by traversing derivations and accumulating information in an *ad hoc* way. In 1968, Donald Knuth described a formalisation of this approach called an *Attribute Grammar (AG)* [Knu68, Knu71, Knu90]. The values of attributes in an Attribute Grammar are specified using a purely equational approach, and the equations for attributes at a particular derivation tree node must only depend on other attributes which are locally visible. For a subtree comprising some parent node and its children, attributes of the parent node whose equations depend only on attributes of the child nodes are called *synthesized* attributes. Attributes of child nodes which are updated by equations are called *inherited* attributes.

Attribute Grammar (AG)

A substantial research literature on Attribute Grammars exists, much of which focusses on procedures to evaluate attributes in an order that optimises evaluation time. One important special case (called an L-attributed grammar) allows evaluation in a single post-order traversal of the tree.

Many parser generator tools incorporate a modified notion of Attribute Grammars that we call an *Attribute-Action Grammar (AAG)*

Attribute-Action Grammar (AAG)

6.1 Attribute Grammars

6.1.1 Derivation traversers

There are a few applications for which the tree *is* the semantics, and no further processing is required. Consider, for instance, the task of deciding whether two student programs are identical up to variables names. The derivation tree over language tokens typically captures this information if we ignore the individual lexemes associated with identifier leaf nodes. In principle, we could output a textual form of the derivation tree and use an operating-system level utility to compare them with no other programming required.

We might be interested in software metrics of various kinds, for instance, such as the average number of statements in a method, or the maximum nesting depth of control flow statements. These sorts of applications require simple computations over the tree such as counting the number of nodes in a subtree, or counting the maximum depth of a tree. We might also want to write a tool that enforced some coding standards: for instance a software company might require all control flow statements in Java to have braces around their bodies even when they are single statements. These kinds of applications require straightforward tree traversals.

Now consider code refactoring. A common requirement is to rename all instances of a variable X , say, to Y in a program. A correct implementation must not simply change all of the X identifiers to Y because we may be using the same name for several different variables. For instance, many Java methods operating on strings might have an argument called `string`; a refactoring centred on one of those methods must leave the others untouched. Clearly we need a *semantics-aware* replacement which only updates instances which are within the same scope region. At this point, the derivation tree is no longer sufficient in itself: we additionally need some representation of the scope semantics of our language so that we can distinguish independent variables that happen to have the same name.

For a data-centric language, again the tree itself might be a near-sufficient representation. The derivation for an XML description of a document contains all of the information in the document, along with styling information, and one could build, say, a word processing application around routines which traversed the tree to compose an on-screen representation of the document, and which modified the tree in response to insertion, deletion and styling requests from a graphical user interface. The XML derivation tree is thus being directly used as the *internal form* for the word processor.

6.2 Attribute Grammars

It is natural to think of the leaves of a derivation tree as being associated with values: for instance an INTEGER token matching the string "0123" is associ-

ated with the value 123 and so on. When implementing arithmetic expressions, it is useful to think of these values as percolating up through the tree, being transformed by operators as we go.

Many early compilers used these sorts of ideas, and in the late 1960's Donald Knuth formalised these ideas by associating attributes with nonterminals in a grammar, such that every (a) instance in a derivation tree of some nonterminal X would have the same attribute set; (b) the values of attributes would be specified by equations; and that (c) if an attribute of X were defined in a production of X , then it must be defined in *all* productions of X .

Knuth distinguished between *inherited* and *synthesized* attributes. Conceptually, the use of inherited attributes causes information to be passed down the tree, and uses of synthesized attributes represent upwards data flow. It turns out that formally we can write equivalent specifications that use either only inherited or only synthesized attributes, but in practice it is convenient to use both. We have already seen that expression evaluation uses upwards propagation of values, and indeed expression evaluators typically use synthesized attributes. Context information, such as the declared types of variables often needs to be propagated down into sections of the tree, and inherited attributes are the appropriate means to do so.

It is natural to think of the leaves of a derivation tree as being associated with values: for instance an INTEGER token matching the string "0123" is associated with the value 123 and so on. When implementing arithmetic expressions, it is useful to think of these values as percolating up through the tree, being transformed by operators as we go.

Many early compilers used these sorts of ideas, and in the late 1960's Donald Knuth formalised these ideas by associating attributes with nonterminals in a grammar, such that every (a) instance in a derivation tree of some nonterminal X would have the same attribute set; (b) the values of attributes would be specified by equations; and that (c) if an attribute of X were defined in a production of X , then it must be defined in *all* productions of X .

Knuth distinguished between *inherited* and *synthesized* attributes. Conceptually, the use of inherited attributes causes information to be passed down the tree, and uses of synthesized attributes represent upwards data flow. It turns out that formally we can write equivalent specifications that use either only inherited or only synthesized attributes, but in practice it is convenient to use both. We have already seen that expression evaluation uses upwards propagation of values, and indeed expression evaluators typically use synthesized attributes. Context information, such as the declared types of variables often needs to be propagated down into sections of the tree, and inherited attributes are the appropriate means to do so.

6.2.1 The formal attribute grammar game

Let $\Gamma = (T, N, S, P)$ where T is a set of terminals, N is a set of nonterminals ($T \cap N = \emptyset$), $S \in N$ is the start nonterminal which must not appear on any RHS (and so S must not be recursive), and $P = (T \times (N \setminus S))^*$ is a set of productions.

Each symbol $X \in V$ has a finite set $A(X)$ of attributes partitioned into two disjoint sets, synthesized attributes $A_S(X)$ and inherited attributes $A_I(X)$.

The inherited attributes of the start symbol (elements of $A_I(S)$) and the synthesized attributes of terminal symbols (elements of $A_S(t \in T)$) are pre-initialised before attribute evaluation commences: they have constant values.

Annotate the CFG as follows: if Γ has m productions then let production p be

$$X_{p_0} \rightarrow x_{p_1} x_{p_2} \dots x_{p_{n_p}}, \quad n_p \geq 0, \quad X_{p_0} \in N, \quad X_{p_j} \in V, \quad 1 \leq j \leq n_p$$

A semantic rule is a function f_{pja} defined for all $1 \leq p \leq m, 0 \leq j \leq n_p$; if $j = 0$ then $a \in A_S(X_{p,0})$ and if $j > 0$ then $a \in A_I(X_{p,j})$.

The functions map $V_{a_1} \times V_{a_2} \times \dots \times V_{a_t}$ into V_a for some $t = t(p, j, a) \geq 0$

The ‘meaning’ of a string in $L(\Gamma)$ is the value of some distinguished attribute of S , along with any side effects of the evaluation.

6.2.2 Attribute grammars in practice

When we want to engineer a translator using attribute grammars we have to do two things: (a) consider the fragments of data that need to reside at each node (the attributes) and (b) the manner in which those attributes will be assigned values. Let us construct by example some concrete syntax for attribute grammars which incorporates the abstract syntax represented by the definitions in the previous section.

Consider a rule of the form

¹ $X ::= Y Z \quad Y.X.a = \text{add}(Y1.v, Y2.v) \quad Z1.v = 0$

The BNF syntax is as usual. The equation is written as an attribute $X.a$ followed by $=$ sign and then an expression involving other attributes. The scope of an equation is just a single production which means that the only grammar elements (and thus attributes) that may be referenced in an equation are the left hand side nonterminal, and the terminals and nonterminals on the right hand side. In Knuth’s definition, the LHS nonterminal has suffix zero, but typically in real tools we drop the suffix and just use the nonterminal name as here. The

right hand side instances are numbered in a single sequence; in tools we often maintain separate sequences for each unique nonterminal name as here.

One can tell syntactically whether an attribute is inherited or synthesized by examining the left hand side of its equation: if the LHS of an equation is an attribute of the left hand side of the rule, then in tree terms we are putting information into the parent node, and thus information is flowing up the tree and this must be a synthesized attribute. If the LHS of an equation references one of the right hand side production instances then we are putting information into one of the children nodes, and information is flowing down the tree (so this must be an inherited attribute). In the example above, $X.a$ is synthesized and $Z1.v$ is inherited.

It is perfectly possible to completely define a real translator or compiler using attribute grammars, and many tools exist to support this methodology. Pure attribute grammars are a declarative way of specifying language semantics using just BNF rewrite rules and equations, and it is the job of the attribute evaluator to find an efficient way to visit the tree nodes and perform the required computations.

6.2.3 Attribute grammar subclasses

A variety of attribute grammar subclasses have been defined, mostly in an attempt to ensure that equations may be evaluated in a single pass on-the-fly by near deterministic parser generators. For instance, the LR style of parsing used by Bison is bottom up, that is the derivation tree is constructed from the leaves upwards. If we are to do attribute evaluation at the same time, then we must restrict ourselves to equations that propagate upwards: hence all of the attributes must be synthesized (and equations are usually written at the end of the production to ensure that all values are available). Such attribute grammars are called *S-attributed*.

For top-down recursive descent parsers we can handle a broader class of attribute grammars. As with bottom up, the requirement is that attribute values be computable in an order which matches the construction order of the derivation tree. Such attribute grammars are called *L-attributed*: in an L-attributed grammar, in every production

$$X \rightarrow y_0 y_1 y_2 \dots y_k \dots y_n$$

every inherited attribute of y_k depends only on the attributes of $y_0 \dots y_{k-1}$ and the inherited attributes of X . This definition reflects the left-to-right construction order of the derivation tree.

6.3 Semantic actions in ART

Semantic actions and attributes in ART use the parser's implementation language to model the decalarative, equational attribute grammar formalism. Back

end languages for ART include C++ and Java, which are languages that do not enforce referential transparency. As a result, it is possible to write attribute grammar specifications in ART which are not equational: specifically, attributes in ART are procedural language variables to which we make assignments, and so in principle we can have several ‘equations’ in a rule all of which target the same attribute, which means that the value of an attribute is no longer a once-and-for-all thing, but instead may evolve during the parse.

From a formal point of view, this is very ugly. From a software engineer’s perspective, it is an opportunity to introduce efficiencies. Which camp you are in rather depends on your primary concerns.

Although ART attributes are in some senses more powerful than true AG attributes, the evaluator in ART is definitely less powerful than would be required for a true AG evaluator, as we shall see.

6.4 Syntax of attributes in ART

In ART, user attributes must be declared for each nonterminal. A rule such as

```
1 | X <value:String number:int> ::= 'x'
```

specifies that an instance of *X* has two attributes: one of type **String** called *value* and another of type **int** called *number*.

An action in ART is specified on the right hand side of the rule within curly braces { and }. Any syntactically and semantically valid fragment of Java may appear within the braces. It is important to understand that ART treats material within these braces as a simple string—ART does not understand the syntax or semantics of Java or any of the other backend languages, and so cannot test for errors within the string. If you write an action which is ill-formed, you will only find out when either (a) the compiler for the back end language attempts to process the output from ART or (b) when the evaluator actually runs. This can make debugging semantic actions somewhat challenging. This is in the nature of meta-programming: the ART specification is effectively a specification for a program that ART will write, so you are one step removed compared to the normal software engineering process.

So, for instance,

```
1 | X <value:String number:int> ::= 'x' { X.value = 3; }
```

is a valid ART specification which generates a compile-time-invalid piece of Java because the expression 3 is not type compatible with the attribute *value* which is of type **String**. Similarly, if an attribute was an array and an action tried to access an element which was out of range, then the error in the action would

not be picked up by either ART or the Java compiler, but instead generate a run-time exception.

6.4.1 Special attributes in ART

ART recognises two special attribute names: `leftExtent` and `rightExtent`. When the user declares attributes with those names and type `int` they are treated just as for ordinary attributes except that the parser initialises them with the start and end positions of the string matched by that nonterminal instance. As a result, one would not usually expect to find attribute equations in the actions that had either `leftExtent` or `rightExtent` on their left hand sides.

The use of these special attributes allows us to create attributes at the lowest level of the tree. In ART, attributes cannot be defined for builtin terminals or terminals that are created literally. However, we can wrap an instance of a terminal in a nonterminal, and then use these special attributes to extract the substring matched by some terminal. For instance, here is the definition of a nonterminal `INTEGER` which uses the `&INTEGER` lexical builtin matcher to match a substring, and then extracts a value using the `leftExtent` and `rightExtent` attributes

```
1 | INTEGER <leftExtent:int rightExtent:int v:int> ::= &INTEGER
2 | {INTEGER.v = artLexemeAsInteger(INTEGER.leftExtent, INTEGER.rightExtent);}
```

ART provides a set of methods for converting substrings of the input to values: `artLexemeAsInteger()`, `artLexemeAsDouble()` and so on.

6.5 Accessing user written code from actions in ART generated parsers

It would be cumbersome to have to put the entire functionality of a translator into semantic actions. Instead, we would like to parcel complex operations up into functions or class methods, and simply call them from the semantic actions.

Back end languages for ART vary in their requirements, but for the Java back-end we can imagine both wanting to access objects of classes outside of ART's generated parser class, and also the addition of members to the ART generated parser class itself. ART provides two mechanisms to help.

The `prelude{...}` declaration specifies that the material within the braces be copied into the generated code at the top of the file. This enables us to add, for instance, `import` declarations to the Java generated parser, and thus to access objects and static methods of other classes within our semantic actions.

The `support{...}` declaration specifies that material within the braces be

copied into the generated code within the ART generated parser class itself, allowing us to declare methods and variables which are visible throughout the generated parser's actions.

6.6 A naïve model of attribute evaluation

How would we build a (not very efficient) general attribute evaluator for ART? Let us begin by giving each node in the tree a unique *instance* number. Then each attribute may be uniquely named as (instance number, name). Make a set U which will contain the subset of attributes which are presently undefined. Make a map V from attribute names to attribute values which is initially empty.

We begin by handling the special attributes `leftExtent` and `rightExtent`. For each attribute in $u_i \in U$ with a name of the form $(k, \text{leftExtent})$ or $(k, \text{rightExtent})$, remove u_i from U and add an element to map V which maps $(k, \text{left}(\text{right})\text{Extent})$ to the first(last) index position of the substring matched by instance k .

Now, while U is nonempty, traverse the entire tree and examine, all of the equations for productions used in the derivation sequence and perform these actions

1. If the attribute on the LHS of the equation is not in U , then continue.
(The attribute has already been computed.)
2. If the attribute on the LHS is in U and any attribute on the RHS is in U , then continue. (The attribute is not ready to be computed.)
3. If the attribute (k, n) on the LHS is in U and no attribute on the RHS is in U , remove (k, n) from U and add an element to V mapping (k, n) to the result of computing the right hand side expression.

Recall that a well-formed attribute grammar must be (a) non-circular and (b) must have an equation in every production defining the value of all attributes in its RHS nonterminal. As a result, there must be some ordering over the equations that allows them to be resolved. This algorithm finds an ordering by brute force: it simply continually traverses the tree looking for so-far undefined LHS attributes whose right hand side attributes *are* defined, at which point it computes the new value and removes the attribute from the undefined set.

This algorithm is simple, but inefficient because in worst case we might only be able to compute one equation per entire pass of the tree. In practice, real general attribute evaluators perform *dependency analysis* on the equations to find much more efficient schedules.

6.7 The representation of attributes within ART generated parsers

ART provides an abstract class `ARTGLLAttributeBlock`. Inside ART, nonterminals are named $M.N$ where M is a module name and N is the name of a nonterminal defined in module M . The default module name is `ART` so in specifications with explicit module handing, a nonterminal called X by the user is called `ART_X` internally.

For each attributed nonterminal $M.X$, ART creates a concrete subclass of `ARTGLLAttributeBlock` called `ART_AT_M_N`, so for instance the ART rule

`1 X <p: int q:double> ::= 'x'`

in module M generates the class

```

1 public static class ART_AT_M_X extends ART_GLLAttributeBlock {
2     protected double q;
3     protected int p;
4 }
```

A separate instance of this class is created for each instance of nonterminal $M.X$ in the derivation. Each instance effectively has two names within the attribute evaluator: $M.X$ for left hand side attributes and $M.X_k$ for right hand side instances, where k is an integer. When we write an action like `M_X.v = 3;` we mean, locate the attribute block for my left hand side which is called `M_X` and then access the field called `v`. When we write an action like `M_X.v = M_X1.v` we are asking for the `v` value from the attribute block for the first instance of $M.X$ on the right hand side of our rule to be copied to the left hand side instance.

6.8 The ART RD attribute evaluator

Whilst we could implement an attribute evaluator based on the general model above, it would be inefficient. Instead we implement *syntax directed translation*.

Rather than seeking a schedule which resolves all of the data dependencies in the attribute grammar, we instead assert a particular schedule and require the writer of the attribute grammar to not write equations which violate its constraints. We say that an AG specification is *admissible* if it may be computed by our predefined schedule, and *inadmissible* otherwise.

The ART attribute evaluator correctly evaluates *L-attributed* grammars. In an L-attributed grammar, in every production

$$X \rightarrow y_0 y_1 y_2 \dots y_k \dots y_n$$

every inherited attribute of y_k depends only on the attributes of $y_0 \dots y_{k-1}$ and the inherited attributes of X .

There is quite a strong parallel here with parsing: a general parsing algorithm such as GLL (the algorithm ART implements) can handle any specification, but with the risk of poor performance on some grammars. A non-backtracking Recursive Descent parser, on the other hand, can only handle deterministic LL(1) grammars (or ordered grammars which are nearly LL(1)) but will run in linear time.

Our evaluator is essentially a recursive descent evaluator. It will only traverse the tree once. As long as the equations may be fully resolved in a single pass, all will be well. The ART evaluator is limited to attribute schemes that are essentially the L-attributed schemes. However, we can do a lot with such schemes, and the evaluation time is linear in the size of the tree.

In detail, the ART evaluator recurses over the datastructure constructed by the GLL parser. This is not a single derivation, but a (potentially infinite) set of derivation trees embedded within a structure called a Shared Packed Parse Forest (SPPF). However, prior to starting the evaluator, we will have marked some parts of the SPPF as suppressed, and some parts as selected, and the net effect is that the evaluator can assume that it is recursing over a single derivation tree.

As the evaluator enters a node labeled X , it creates the attribute block for each nonterminal child below it (corresponding to the nonterminal instances in the derivation step $X \Rightarrow \alpha$ encoded in this height-1 sub-tree). These newly-created attribute blocks are assigned to variables with names like $Y1$ and $Z2$ corresponding to the first and second instances of Y and Z in some production like $X ::= Y \ Z \ Z$

The evaluator is a nest of functions, one for each nonterminal, in a way that is isomorphic with our OSBTRD parser functions. In our example above, the evaluator function for X will be called and make attribute blocks for the children $Y1$, $Z1$ and $Z2$. It will then call the evaluator function for Y passing block $Y1$ as an argument. The evaluator functions all take a single parameter block whose name is the same as that of the nonterminal. By this means, the block for Y allocated in X is called $Y1$ in the evaluator for X but called Y in the evaluator for Y .

Just like an RD parser, the evaluator functions call each other in the same order as instances are encountered within the grammar, and the semantic actions are inserted directly into the evaluator functions.

6.9 Higher order attributes

There is a well-developed theory of higher order attributes, which are attributes that represent parts of derivation trees rather than simple values. There are essentially two classes of HO attributes: attributes which capture part of an existing derivation tree, and attributes which contain new pieces of tree which can be used to extend a derivation tree from the parser. In ART, we support

the former, but not yet the latter. This means that the shape and labelling of a derivation tree in ART cannot be modified by an attribute grammar: only the attribute values associated with tree nodes can be modified. In a later section we shall describe ART's GIFT operators which do allow trees to be modified. In the present implementation, the attribute evaluator works on the full derivation and completes evaluation before the GIFT rewriter changes the tree.

ART's notion of higher order attributes requires only two things: a way of marking tree nodes as having a higher-order attribute associated with them, and a way to allow the user to activate the evaluator function under the control of semantic actions.

The first is achieved by adding an annotation < to any right hand side instance of a nonterminal in a grammar rule. The second is achieved by providing a method `artEvaluate()` which takes as an argument a higher order attribute.

When the evaluator function arrives at a node with a higher-order attribute, it does not descend into it (although it will construct the attribute block for it). The idea is that instead of automatically evaluating a subtree, the outer evaluator will ignore it, but the user may specify semantic actions to trigger its evaluation on demand.

Why is this useful? Well one application is to allow our recursive evaluator to interpret flow-control constructs. Consider an `if` statement. It comprises a predicate, and a statement which is only to be executed if the predicate is true. We can specify this as follows:

```

1 ifStatement ::= 'if' e0 'then' statement<
2   { if (e0.v != 0) artEvaluate(ifStatement.statement1, statement1); } ;

```

The < character after the instance of `statement` creates a higher order attribute called `statement1` in the attribute block for `ifStatement`. ART will also have created an attribute block called `statement1`. The evaluator will automatically descend into the subtree for `e0`, but will not descend into the subtree for `statement`: instead it loads a reference to the subtree for this instance of `statement` into the attribute `statement1` in `ifStatement`.

In the action, we look at the result that was computed within `e0`, and if it is not zero (signifying false) we call the evaluator on the the subtree root node held in the attribute `ifStatement.statement1` and pass in parameter block `statement1`. This effectively emulates what would have happened automatically if we had left off the < annotation, but under the control of the result of `e0`. Hence the evaluation order of the tree is being dictated by the attributes and semantic actions themselves! This is exactly the sense in which our attributes are higher order. However, we can only traverse bits of tree that were built by the parser: we cannot make new tree elements and call the evaluator on them. Full higher order attributes do allow that. We call our restricted form *delayed* attributes so as to distinguish them from the more general technique.

```

!global variables:_map

statements ::= statement | statement statements

statement ::=

  ID ':=' subExpr ';' statement.v = __put(variables, ID1.v, subExpr1.v)
  | 'if' relExpr statement 'else' statement
    statement.v = relExpr.v ? statement1.v !! statement2.v
  | 'while' relExpr statement statement.v = relExpr.v @ statement.v !! __done

relExpr ::=

  subExpr           relExpr.v = subExpr1.v
  | relExpr '>' subExpr relExpr.v = __gt(relExpr.v, subExpr.v)
  | relExpr '!=>' subExpr relExpr.v = __ne(relExpr.v, subExpr.v)

subExpr ::=

  operand           subExpr.v = operand.v
  | subExpr '-' operand subExpr.v = __sub(subExpr.v, operand.v)

operand ::=

  ID                 operand.v = __get(variables, ID.v)
  | INTEGER          operand.v = INTEGER1.v
  | '(' subExpr ')' operand.v = subExpr1.v

```

Figure 6.1 An attribute-action specification for MiniGCD

We can use these delayed attributes to build interpreters for languages with conditionals, loops and function calls, as we shall see in the laboratory exercises.

6.10 An attribute-action specification for MiniGCD

6.11 Exercises

Chapter 7

Software language design and pragmatics

7.1 The semantic facets of programming languages

We shall group our discussion of semantic components of high level languages into five *facets*: (a) values, types and expressions; (b) storage, assignment and commands; (c) identifiers, bindings and scope; (d) control flow and (e) abstraction mechanisms, both procedural and data.

7.1.1 Values, types and expressions

The ultimate purpose of a program is to compute *values*; those values might be numeric solutions to equations, textual outputs, visual effects on a screen or movements of a robot arm. At machine level, all of these correspond to patterns of binary digits in memory, but programming languages provide a range of abstractions which enable us to reason more effectively about program execution. Grouping the available values by class of abstraction naturally leads to the notion of *type*.

7.1.2 Storage, assignment and commands

In a von Neumann view of computing, values are stored in reusable cells, and in that world, storage is also fundamental. Most program texts are dominated by identifiers which stand for, for instance, constant values, locations in store,

7.1.3 Identifiers, scope and binding

Nested scope rules

7.1.4 Control flow

D-structures

Concurrency

Jumps

Exceptions

7.1.5 Procedural and data abstraction

Procedures

Higher order functions

Abstract data types, classes and packages

Generics

Appendix A

ART user manual

A.1 Downloading and running ART for the first time

1. ART is written in Java; therefore an up-to-date Java installation is required. At the time of writing, the UK Oracle download page for Java is at

<https://www.oracle.com/uk/java/technologies/downloads/>

Select and install the appropriate version for your operating system.

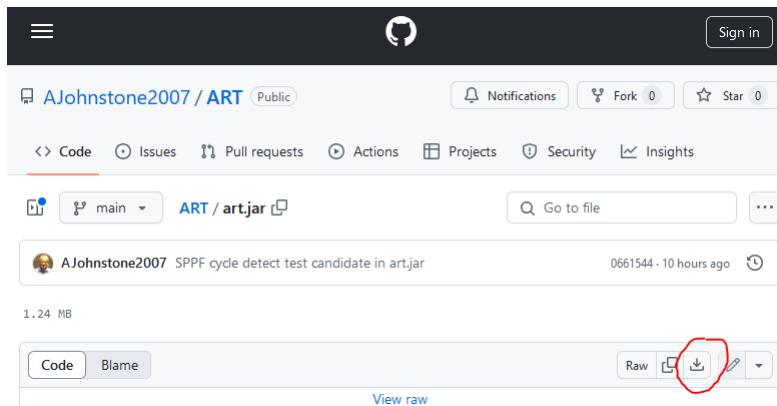
Other Java implementations are available and locatable via search engines.

2. Make a work directory, the location of which we shall call *artwork*.

Download the *art.jar* file by opening a Web browser on:

<https://github.com/AJohnstone2007/ART/blob/main/art.jar>

Click the GitHub download button (circled in red below) to download a copy of *art.jar* to your work directory *artwork*.



3. Test the download and your Java installation by opening a command window, changing your directory to *artwork* and typing the command

`java -jar art.jar`

The expected output is a version number, a build timestamp and summary usage information which will look like this:

ART 5_0_241 2024-11-01 08:12:44

Usage:

...

4. Instead of the ART message you may see something like this:

```
Class has been compiled by a more recent version of the Java Environment
(class file version xy.0), this version of the Java Runtime only recognizes
class file versions up to pq.0.
```

This means that your Java installation is for an old version of Java, and you will need to install a current version: see step 1.

5. The official ART repository is at

<https://github.com/AJohnstone2007/ART>

It includes the latest version of this document at

<https://github.com/AJohnstone2007/ART/blob/main/doc/slewa.pdf>

and the source code under

<https://github.com/AJohnstone2007/ART/tree/main/src>

A.2 IDE and graphical component installation

The `art.jar` file includes an Integrated Development Environment (IDE) and support for building languages that display 2D and 3D graphics. Graphical support requires JavaFX, and the IDE in addition uses the RichTextFX Java component for text editing.

1. JavaFX is no longer part of the main Java installation, and must be separately installed via the page at

<https://gluonhq.com/products/javafx/>

2. The RichTextFX component repository is at

<https://github.com/FXMisc/RichTextFX>

ART looks for a copy of the ‘fat jar’ (which contains all RichTextFx’s dependencies) when it starts up. This can be directly downloaded to your `artwork` directory from

<https://github.com/AJohnstone2007/ART/blob/main/richtextfx.jar>

3. Running ART with JavaFX requires a very long command line because of the need to specify class and module paths.

The Windows batch script `art.bat` contains this line:

```
java --module-path %jfxHome%\lib --add-modules javafx.controls
      -cp .;%artHome%\art.jar;%artHome%\richtextfx.jar
      uk.ac.rhul.cs.csle.art.ART %*
```

where `%jfxhome%` is the name of an environment variable bound to the location of the Java FX modules and `%artHome%` is the location of `art.jar`.

A.3 Command line interface

A.4 Integrated Development Environment

January 2025: the IDE needs further development before it is ready for serious work. Please use the command line interface.

A.5 ART script language

An ART script is built from four kinds of elements:

1. Rewrite rules, of the form *premises* \dashrightarrow *conclusion*
2. Context Free Grammar Rules, of the form *identifier* $::=$ *cfgExpression*
3. Choose rules, of the form *slotSet* $>$ *slotSet* or *slotSet* \gg *slotSet*
4. Directives, which begin with an exclamation mark !

The script language is free format, and arbitrary whitespace or comments may appear before and after each script language token. The ART script language specification is available from the repository at

<https://github.com/AJohnstone2007/ART/blob/main/src/uk/ac/rhul/cs/csle/art/script/scriptSpecification.art>

A.6 Lexical structure

1. **identifier** zzz
2. **signed integer**
3. **signed real**
4. **string literal**
5. **character literal**
6. **filename**

A.7 Abbreviations

```
(* ART front end abbreviations and their internal forms: demonstrate the cooked and the raw)
!print "A term with no abbreviations"
!print term a(b,c)

!print "** __bool"
!print term true
!prinraw term true
```

```
!print term false
!prinraw term false

!print "** __char"
!print term `a
!prinraw term `a

!print "** __intAP"
!print term £1234
!prinraw term £1234

!print "** __int32"
!print term 1234
!prinraw term 1234

!print "** __realAP"
!print term £1234.0
!prinraw term £1234.0

!print "** __real64"
!print term 1234.0
!prinraw term 1234.0

!print "** __array of size 3"
!print term [3 | a,b,c ]
!prinraw term [ 3 | a,b,c]

!print "** __list"
!print term [ a,b,c ]
!prinraw term [ a,b,c]

!print "** __set"
!print term { a,b,c }
!prinraw term { a,b,c }

!print "** __map"
!print term { a=p, b=q, c=r }
!prinraw term { a=p, b=q, c=r }
```

A.8 Rewrite rules

A.9 Context free grammar rules

A.10 Choose rules

A.11 Directives

1. **!include**
2. **!whitespace**
3. **!paraterminal**
4. **!lexer**
5. **!parser**
6. **!interpreter**
7. **!start**
8. **!configuration**
9. **!clear**
10. **!trace**
11. **!print**
12. **!show**
13. **!prompt**
14. **!try**
15. **!nop**

A.12 Lexical builtins

Lexical builtins are hardcoded recognisers for certain classes of substring which may be used as shorthands for common lexical patterns on the right hand side of Context Free Grammar rules. Builtin names begin with an ampersand & character.

1. **&CHAR_BQ** ‘C
2. **&ID** AlphanumericIdentifier
3. **&INTEGER** 123
4. **&REAL** 12.3

5. **&STRING_BRACE** {A string delimited by braces}
6. **&STRING_BRACE_NEST** {A string {with nested instances} delimited by braces}
7. **&STRING_DOLLAR** \$A string delimited by dollar signs\$
8. **&STRING_DQ** "A string delimited by double quotes"
9. **&STRING_PLAIN_SQ** 'A string delimited by single quotes with no escapes'
10. **&STRING_SQ** 'A string delimited by single quotes'

The following lexical builtins can only appear as an argument to the `!whitespace` directive. They are discarded by the lexer, and will never appear in a lexicalisation (and so it would be an error for them to appear within a Context Free Grammar rule).
11. **&SIMPLE_WHITESPACE**
12. **&COMMENT_BLOCK_C** /* a C-style block comment */
13. **&COMMENT_LINE_C** // a C-style line comment
14. **&COMMENT_NEST_ART** (* An ART style comment (* nestable *) *)

A.13 The ART value system

ART provides several builtin types and operations which may be used instead of rewrite rules to perform more efficient basic arithmetic and collection operations.

A.13.1 Types

A.13.2 Operations

A.14 ART value plugins

Revised: 24 January 2025		Type	bottom	done	empty	quote	blob	adpprod	adtsum	proc	bool	char	int32	real32	real64	string	array	list	set	map	Return type
Operation	Literals	bottom	done	empty	quote	blob	<1>	<1>	\00	true/false	'x'	£32	32	£32.00	"xyz"	[N ...]	[...]	{...}	{->}		
Equal	eq	V,V	V,V	V,V	V,V	V,V	V,V	V,V	V,V	V,V	V,V	V,V	V,V	V,V	V,V	V,V	V,V	V,V	V,V	B	
Not equal	ne	V,V	V,V	V,V	V,V	V,V	V,V	V,V	V,V	V,V	V,V	V,V	V,V	V,V	V,V	V,V	V,V	V,V	V,V	B	
Greater than	gt																				B
Less than	lt																				B
Greater than or equal	ge																				B
Less than or equal	le																				B
compare: return -1, 0 or +1	comp																				N
Logical/bitwise not	not																				B
Logical/bitwise and	and																				B
Logical/bitwise or	or																				B
Logical/bitwise exclusive or	xor																				B
Shift	shift																				N
Shift with sign extension	sshift																				N
Rotate	rot																				N
neg																					N
add	Add																				N
subtract	Subtract																				N
multiply	Multiply																				N
divide	Divide																				N
remainder after division	mod																				N
Exponentiate	exp																				N
Cardinality	card																				N
Insert element	put																				N
Lookup element	get																				N
Remove	remove																				N
Concatenate	cat																				N
Prefix	prefix																				N
Suffix	suffix																				N
Set union	unite																				N
Set intersection	intersect																				N
Set difference	diff																				N
New value from deep copy	cast																				N
		Bottom	Done	Empty	Quote	Blob				Boolean	Character	Biginteger	Integer	BigDecimal	Double	String	ArrayList	List	LHSMap		
Raw term operations		Return																			
Arity of T	termArity	T																			
T without children	termRoot	T																			
Call user plugin	plugin	A,A,...	A																		
Collection constructors																					
a	s	m																			

Figure A.1 ART Value system: types, operations and signatures

Constructor	Rôle	Operations
<code>--bottom</code>	Match failure	<code>--eq</code> <code>--ne</code>
<code>--done</code>		
<code>--empty</code>		
<code>--quote</code>		
<code>--blob(<i>N</i>)</code>		
<code>--proc(<i>M,B</i>)</code>		
<code>--adtprod(<i>V,N</i>)</code>		
<code>--adtsum(<i>V,N</i>)</code>		
<code>--bool(<i>V</i>)</code>		
<code>--char(<i>V</i>)</code>		
<code>--intAP(<i>V</i>)</code>		
<code>--int32(<i>V</i>)</code>		
<code>--realAP(<i>V</i>)</code>		
<code>--real64(<i>V</i>)</code>		
<code>--string(<i>V</i>)</code>		
<code>--array(<i>N</i>,<code>--a(...)</code>)</code>		
<code>--list(<code>--l(...)</code>)</code>		
<code>--set(<code>--s(...)</code>)</code>		
<code>--map(<code>--m(...)</code>)</code>		
<code>--map(<code>--map(...)</code>, <code>--m(...)</code>)</code>		

Table A.1 ART value types and allowed operations

Constructor	Returns	Action
<code>__eq(L, R)</code>	<code>_bool</code>	value of L equal to value of R
<code>__ne(L, R)</code>	<code>_bool</code>	value of L not equal to value of R
<code>__gt(L, R)</code>	<code>_bool</code>	value of L greater than value of R
<code>__lt(L, R)</code>	<code>_bool</code>	value of L less than value of R
<code>__ge(L, R)</code>	<code>_bool</code>	value of L greater than or equal to value of R
<code>__le(L, R)</code>	<code>_bool</code>	value of L less than or equal to value of R
<code>__comp(L, R)</code>	<code>_int32</code>	if $L < R$ then -1 else if $L > R$ then $+1$ else 0
<code>__not(L)</code>	$T(L)$	Logical or bitwise inversion
<code>__and(L, R)</code>	$T(L)$	Logical or bitwise conjunction
<code>__or(L, R)</code>	$T(L)$	Logical or bitwise disjunction
<code>__xor(L, R)</code>	$T(L)$	Logical or bitwise exclusive OR
<code>__shift(L, R)</code>	$T(L)$	Left shift L by R bits
<code>__sshift(L, R)</code>	$T(L)$	Right shift L by R bits, propagating zeroes
<code>__rot(L, R)</code>	$T(L)$	Right shift L by R bits, propagating sign bit
<code>__neg(L)</code>		
<code>__add(L, R)</code>		
<code>__sub(L, R)</code>		
<code>__mul(L, R)</code>		
<code>__div(L, R)</code>		
<code>__mod(L, R)</code>		
<code>__exp(L, R)</code>		
<code>__card(L)</code>		
<code>__put(L, K, V)</code>		
<code>__get(L, R)</code>		
<code>__remove(L, K)</code>		
<code>__cat(L, R)</code>		
<code>__prefix(L, R)</code>		
<code>__suffix(L, K)</code>		
<code>__unite(L, R)</code>		
<code>__intersect(L, R)</code>		
<code>__diff(L, R)</code>		
<code>__cast(L, R)</code>		

Table A.2 ART value operations

Appendix B

The Royal Holloway course

This chapter is for students studying Software Language Engineering at Royal Holloway where we approach the material in a particular order designed to allow students to complete their projects within the footprint of a one semester course.

Holloway students start with internal syntax and reduction semantics and only then learn about external syntax parsers and the use of GIFT operators to generate terms in their chosen internal syntax, before moving on to attribute-action systems. After studying that core material we look at topics in lexicalisation and ambiguity management.

Experienced readers will note that this is a ‘semantics-first’ approach: we encourage students to first enumerate the features of their language as a set of signatures, then write reduction rules to interpret those signatures, and only then to consider the external appearance of phrases in their language. We justify and expand on this general approach in Chapter [7](#).

For readers who are not following the Holloway course: if you are using ART just as a parser, or have a particular interest in one or other approach to semantics you may want to take a different route.

B.1 Learning outcomes

After working through these notes, you will

1. know how to use *Context Free Grammar* rules to define programming language syntax;
2. be able to use grammar idioms and the *GIFT annotations* to create *derivation trees* with useful properties;
3. understand how to write *reduction semantics* rules that the rewriter uses to interpret programs;
4. be able to write *attribute-action rules* that may provide more efficient implementations;
5. understand the types and operations of a *Value system*, and how to use a *plugin* to connect to Java classes; and
6. be able to recognise ambiguity in language specifications and progressively eliminate it using *choosers*.

B.2 Assessment

Your command of these learning outcomes will be assessed *via* a personal project and an invigilated examination.

The focus of this course is on the constructs of general purpose programming languages, but to motivate the project work we offer three project variants to construct a *Domain Specific Language* for

1. music generation *via* the Java MIDI interface;
2. image processing using the two-dimensional (2D) features of Java FX; and
3. Computer Aided Design for 3D printing using an extended form of the three-dimensional (3D) features of Java FX.

You must commit to one of these topics by the end of the week 4 lab. To help you decide, please read the background material on each of these in Appendices ??, ?? and ??.

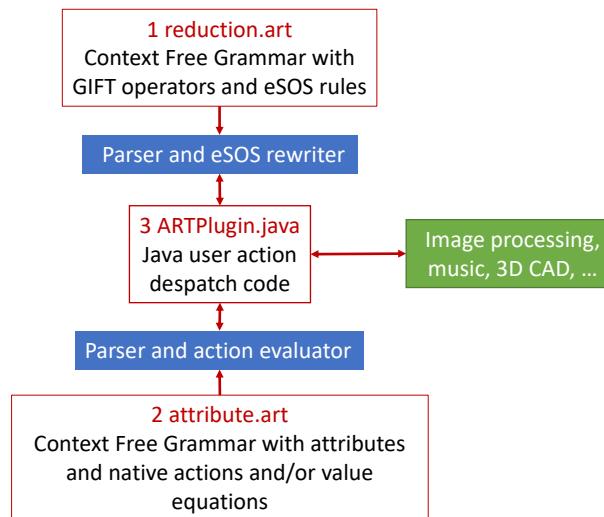
The project work is to be submitted at the end of the course, and requires the following four deliverables.

1. **reduction.art** - a reduction semantics interpreter specified by eSOS rules and a context free grammar annotated with GIFT operators.
2. **attribute.art** - an attribute-action interpreter specified by a context free grammar annotated with attributes and actions written as Value expressions.

3. `ARTPlugin.java` - a plugin which connects your interpreters to your chosen domain specific actions

It is important to remember that this course is about language design and implementation: a project that has a trivial language but a very rich plugin will not score as highly as a rich programming language with a simple backend. Don't be tempted to write a fancy set of domain specific Java code but neglect the core programming language deliverables; on the other hand your plugin should have real functionality and not just report calls as the example.

Here is a diagram showing how the three deliverables interact with the ART system:



The blue boxes above represent language interpretation mechanisms that are built into the ART tool; their behaviour is specified by the rules that you write in `reduction.art` and `attribute.art`. The green box represents arbitrary Java code that you can connect to via a despatch routine in `ARTPlugin.java`.

After each weekly lab session, you will submit a snapshot of your work. These snapshots will not be assessed, but will be automatically analysed so that your progress can be tracked. If you are falling behind then you *must* ask for help.

B.2.1 Marking scheme

The invigilated, unseen examination accounts for 50% of the marks on this course. The remaining 50% of the marks come from the three project deliverables.

Project marking scheme

There are two submission points, A and B shown in blue and green respectively in this table:

Deliverable	Basic	Extension	Beyond	Totals
Internal	2.5			2.5
reduction.art	5	5	5	15
ARTplugin.java	5	5	5	15
attribute.art	5	5	5	15
Examples			2.5	2.5
Blue - submission A				50
Green - submission B				

For each deliverable, the markers look for (a) basic work, which means ‘filling-in’ the missing components of the examples, (b) extension work which means bringing material into your project which is in the examples of this book and (c) ‘beyond’ work which means working out things for which examples are not given in the book.

A complete set of basic language features corresponds to a pass mark of around 40%. Satisfactorily implementing extension features found elsewhere in the course materials will yield marks in the 40-70% range. Adding in features that you have thought about and implemented yourself will help you achieve marks above 70%.

What does this mean in practice? Consider the different facets of language design.

For the expressions facet: the project templates include a few integer-only arithmetic operators. For expressions, a basic pass mark will be achieved if you expand the example to include all of the simple integer operators that are present in Java. An extension-level mark will be achieved if you add in more complex operations such as `++` and `-`, extend the operations to floating point numbers where appropriate and use booleans rather than integers for predicates. A ‘beyond’ mark would require you to also implement arbitrary precision arithmetic with a set of casting operations to allow only sensible conversions - hence some type checking would be required.

For the control flow facet: the project templates only support a basic while loop and an if-then-else statement. A basic pass mark will be achieved if you expand the project template to include other kinds of loop. An extension level mark will be achieved if you add a case (switch) statement and procedure call. A ‘beyond’ mark would require more subtle forms of control flow such as exceptions, lambdas or co-routines.

For the typing facet:

B.3 Teaching week by week

So as to help you prepare your project submission in a timely manner, we approach the learning outcomes in this order.

Week	Chapter	Lecture A	Lecture B
1	1&B	Software language processors	The course
2	2&3	Models of program execution	Rewriting
3	5	Reduction semantics	Reduction semantics
4	5	Reduction semantics	Reduction semantics
5	4	Parsing	Parsing
6	4	Parsing	Parsing
7	6	Attribute action systems	Attribute action systems
8	6	Attribute action systems	Attribute action systems
9	7	Pragmatics	Pragmatics
10	ARTInt	ART internals	ART internals
11		(Contingency)	(Contingency)

Week	Section	Lab sessions and project submission
1	C.1	Lab 1 - OpenSCAD
2	C.2	Lab 2 - rewriting
3	C.3	Lab 3 - eSOS 1
4	C.4	Lab 4 - eSOS 2
5	C.5	Lab 5 - parsing and GIFT
6		Part A project support
7	C.6	Lab 6 - attribute action systems
8	C.7	Lab 7 - delayed attributes and control flow
9		Part B project support
10		Part B project support
11		No lab

Project part A submission in week 7

Project part B submission in week 11

B.4 Protocol for asking questions

Please ask for help over email only.

I will need to be able to reproduce your issue on my machine. Please send a single email, the body of which should contain a concise explanation of your concern with copies of the relevant deliverables as text attachments.

Do not send scripts as screenshots since I cannot run those; you must send scripts as text attachments to your emails.

B.5 Project work week by week

This is suggested timeline for your project work

Weeks 1–2 Learn about the basics of languages from a user's and an implementor's point of view. Decide on a domain of interest—Music, Image Processing or 3D CAD.

Week 3 Enumerate the features of your language and make a list of types and signatures that you need. There should be a signature for every operation in your language: see the lab exercise C.3.1 and C.3.2 to understand how to do this in the case of the GCD language. Decide on a configuration for your language: see lab exercise C.3.3. Extend the GCD language to have a full set of arithmetic, logic and relational operations: see section C.3.8.

Week 4 Extend your rules to use your particular configuration: see section C.4.2 as a starting point.

Week 5 Take the example plugin and extend to a *test* plugin which recognises the back end calls that you need. Write rules for signatures that use the backend and test them using your test plugin, as in the Section C.5 lab exercises. Write a context free grammar for some suitable external syntax, perhaps by extending the existing GCD syntax or by writing something new following the process from Section C.5.)

Weeks 6–7 Add real backend functionality to your plugin. Complete your part A submission by producing test programs that fully exercise your external syntax and backend functions.

Weeks 8 Rework the example GCD attribute grammar to (a) accept the same syntax as your reduction interpreter and (b) interface with your plugin following the approach given in the lab exercises in Sections C.6 and C.7

Weeks 9–11 Enhance your two implementations with more sophisticated *language* (not backend) features such as those suggested in the marking scheme above.

Appendix C

Lab exercises

C.1 Solid modelling with OpenSCAD

This laboratory session aims to help you understand the users' view of Domain Specific Languages. We shall look at the domain of solid modelling languages which may be used, amongst other things, to create objects suitable for 3D printing.

Whilst you are working, keep in mind the following questions: (i) is OpenSCAD easy to get started with? (ii) Is the syntax easy to use? (iii) Are there places where the syntax could be simpler? (iv) Can you pass 3D objects as arguments?

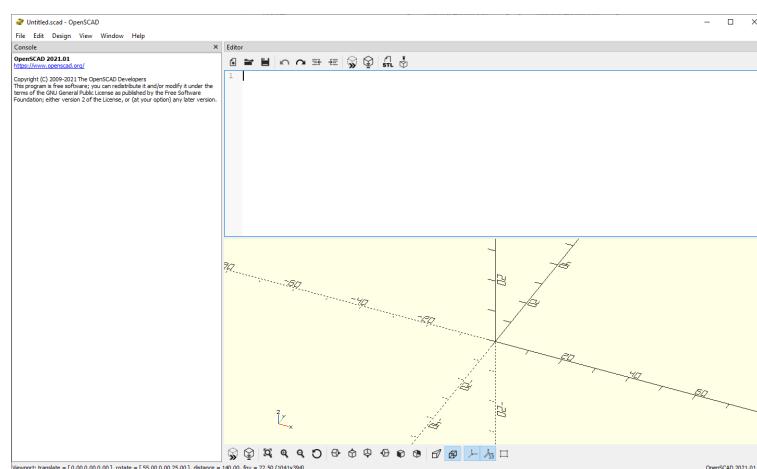
Install OpenSCAD for your system by accessing the download page at

<https://openscad.org/downloads.html>

When you first run OpenSCAD you will see something like this box:



Click on the **New** button and the OpenSCAD environment will open:



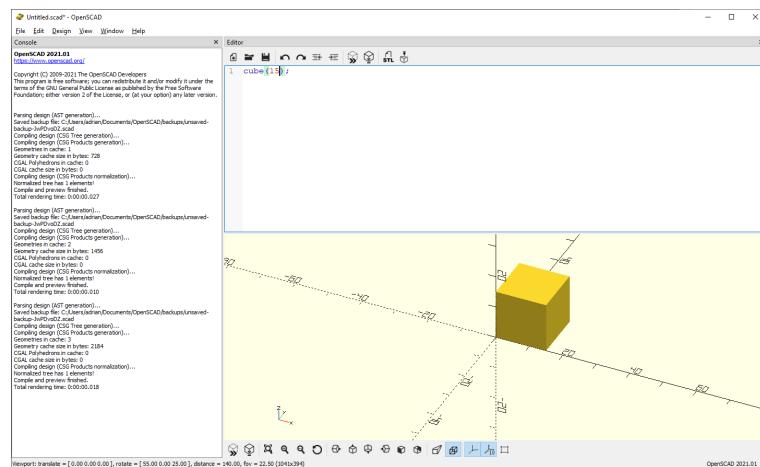
You see three windows: to the left a console and to the right a text editor above a graphics display window.

C.1.1 A first object

We begin by creating a cube. Use the text editor window to enter this command:

```
1 | cube(15);
```

You *view* the code by pressing the **F5** function key. This specification asks for a cube with 15mm edges. You should see this:



C.1.2 Changing the view

Left-click on the displayed object and drag the mouse pointer around. You will find that you can rotate the object around the origin. You can pan the view by right-dragging, and zoom in and out by pinching, using the mouse wheel, or by clicking on the magnifying glass buttons.

There is a Help entry on the menu bar which will open a browser window on some online documentation and tutorials.

C.1.3 Changing size and colour

The numeric argument sets the length of the cube's edge. Change it to 25 and re-view with **F5**. The cube will become larger.

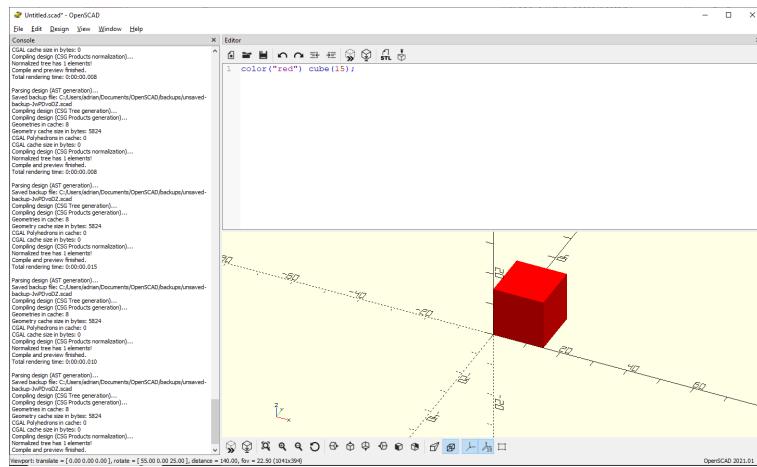
An interesting feature of the text editor is the ability to interactively change the value of a quantity, and immediately see its effect. Position the cursor between

the 5 and the) in the text editor, and then press **[ALT-UPARROW]**. The argument will increment to 16, and the cube will be redrawn. **[ALT-UPARROW]** decrements, and you may use **[ALT-RIGHTARROW]** to add decimal places to a quantity.

You can set the colour of objects by preceding them with a `color()` modifier (note the American spelling!). Try:

```
1 color("red") cube(15);
```

The output should look like:

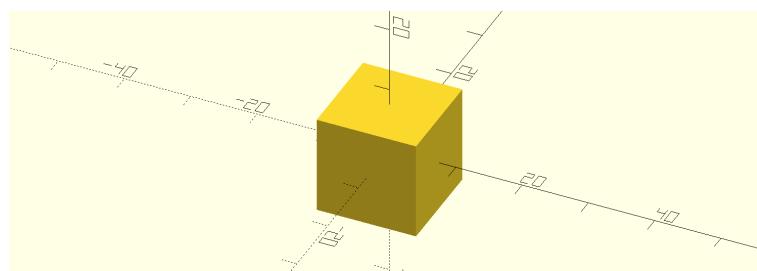


C.1.4 Where is the centre?

Rotations are specified around the origin, and as a result it is useful to ensure that all objects are created at the origin so that they can be re-oriented before being moved to their final position. Some objects like spheres are centred by default, but cubes are not. However, we can add an argument to force centring (again, note American spelling):

```
1 cube(15, center=true);
```

Notice how the cube is centred in all three axes: the coordinate origin is at the centre of the cube.

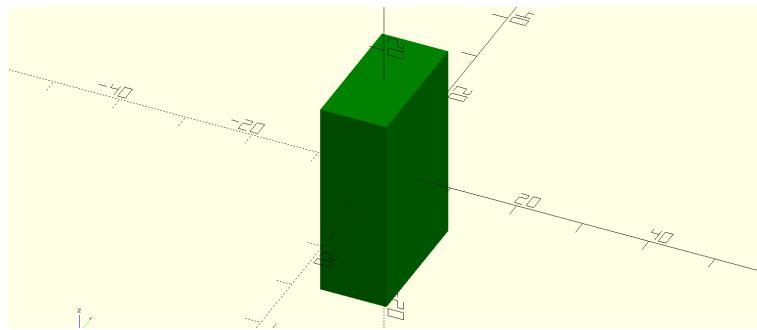


C.1.5 Cubes are really cuboids

The `cube` primitive can be used to specify cuboids, that is objects with varying x, y and z edge lengths.

```
1 | color("green") cube([10,20,30], center=true);
```

The sequence of three numbers within square brackets is a *vector* and may be used to specify the three coordinates. Actually, just using a single integer, say 13, in this context is taken to be shorthand for the vector [13, 13, 13].

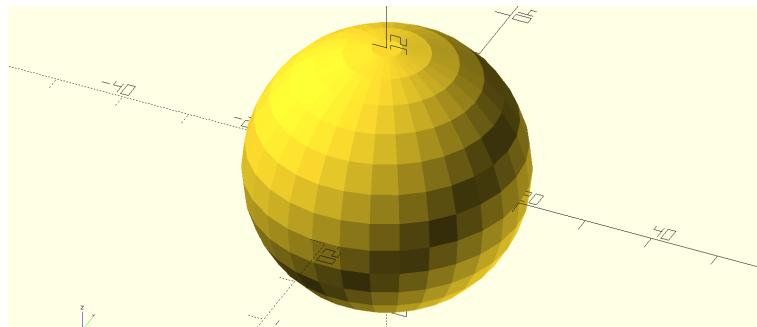


C.1.6 Spheres

OpenSCAD uses the *triangle mesh* method of representing objects: in reality all of the objects are represented as collections of flat triangles.

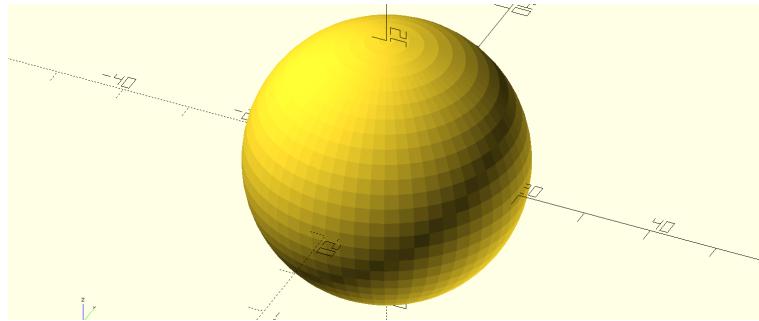
We can model spheres as polyhedra with sufficient faces to make the surface look smooth. The default for a sphere is only 30, which looks blocky if the sphere is large. The `sphere()` primitive constructs a spherical mesh:

```
1 | sphere(20);
```



We can increase the number of segments in a circle (and by extension facets around a sphere) with the special argument `$fn`

```
1| sphere(20, $fn=100);
```



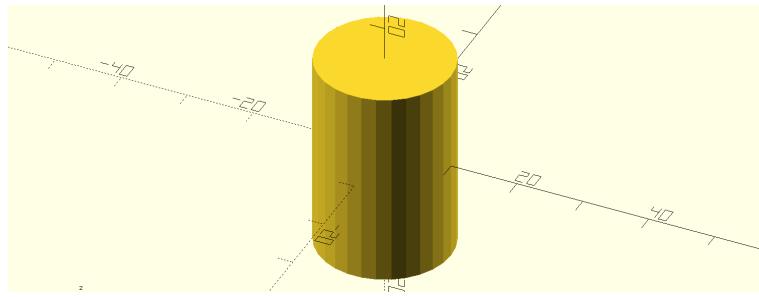
Try raising the value of `$fn` to 120, and then to 360. You will see that the sphere gets much smoother.

This example also shows the use of a language feature called *named arguments* which may be omitted (in which case a default value applies) or they can be explicitly specified. Contrast this with, say, Java method arguments which are strictly positional.

C.1.7 Cylinders and polyhedral bars

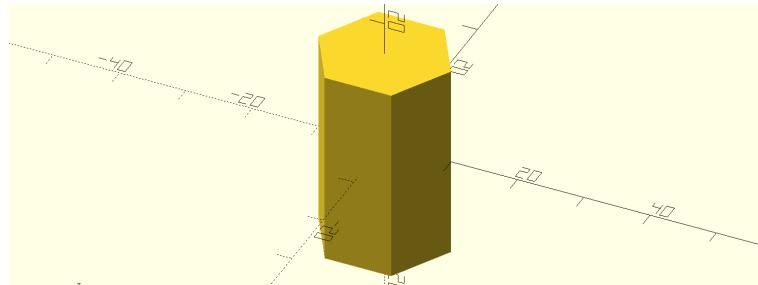
The `cylinder()` primitive takes a height `h` and a radius `r`:

```
1| cylinder(h=30, r = 10, center=true);
```



Now really, the cylinder primitive is just extruding a polygon. As before, we can use the `$fn` argument to make a cylinder smoother. We can also go the other way to make simpler objects. For instance, if we want to make an hexagonal bar, we can set it to six:

```
1| cylinder(h=30, r = 10, center=true, $fn=6);
```



We could in fact make many kinds of box this way too. Is there a `cylinder()` equivalent for every `cube()`?

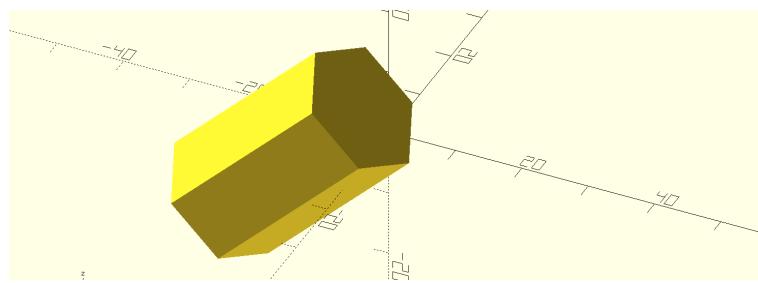
C.1.8 Translation and rotation

So far, we have made objects at the coordinate origin. We can move an object in space using the `translate()` operation which takes as an argument a vector: that is, the usual three values within square brackets corresponding to the x , y and z displacements.

For instance, changing the previous example to

```
1 translate([-10,-10,0]) rotate([0,45,0]) cylinder(h=30, r=10, center=true, $fn=6);
```

moves the hexagonal rod so that its centre is at $(x, y, z) = (-10, -10, 0)$ and so that is tilted 45 degrees around the y axis.



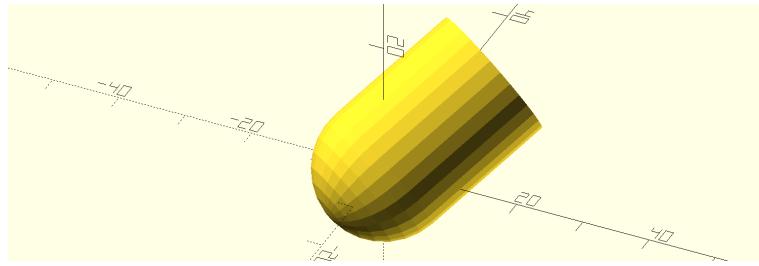
Rotations always occur around axes, so you will get very different effects if you rotate and then translate (as here) as opposed to translating and then rotating. The `translate` and `rotate` operations work like prefix operators, so they are effectively executed right-to-left. Experiment with changing the order of operations, and using the `ALT-ARROW` text editor mechanism to rotate and move things interactively.

C.1.9 Grouping objects

Typically we want to rotate and translate groups of objects together, which we can do by *uniting* them into a single mesh

```

1  rotate([-45,45,0])
2  union() {
3    cylinder(h=20, r = 10);
4    sphere(10);
5 }
```

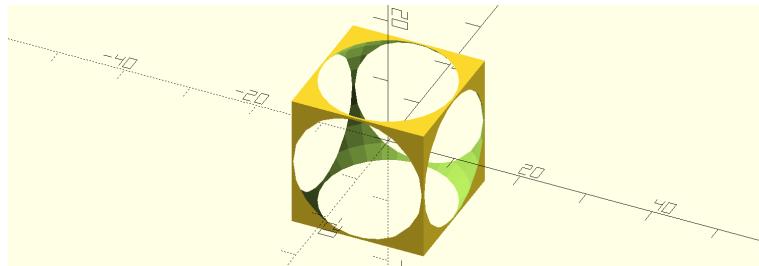


C.1.10 Computational solid geometry

Computational Solid Geometry (CSG) is a technique for making new objects from old via the operations of `union`, `difference` and `intersection`. We met `union` in the previous example. Here the equivalent examples for difference and intersection.

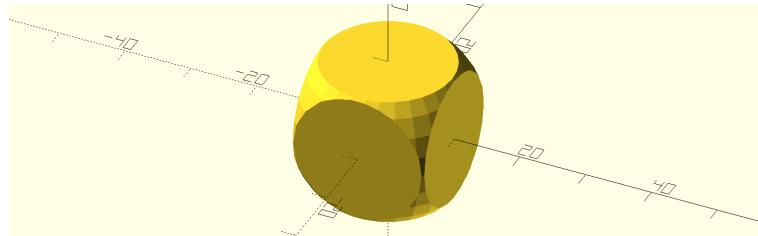
```

1  difference() {
2    cube(20, center=true);
3    sphere(14);
4 }
```



```

1  intersection() {
2    cube(20, center=true);
3    sphere(14);
4 }
```



The `difference()` operation is widely used to make holes in objects by subtracting a cylinder from them.

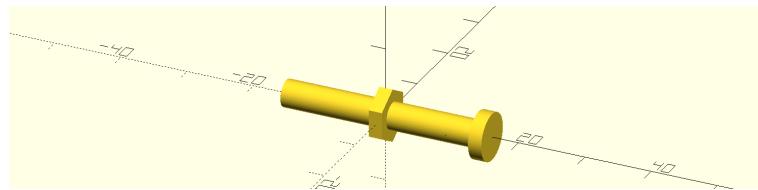
C.1.11 Using modules to structure a design

When programming, we use functions, procedures and methods to break our task down into manageable small pieces. In OpenSCAD, the equivalent construct is the *module* which wraps some 3D object descriptions up together with an implicit union, and allows arguments to *parameterise* the resulting objects.

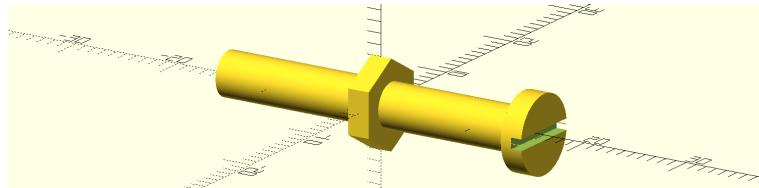
In this example, we make a rough model of an hexagonal nut and a cylindrical bolt which are combined together in the `main()` function. The `bolt()` function takes an argument `length` which specifies how long the bolt should be.

```

1 module nut() { cylinder(h=2, r = 4, $fn=6, center=true); }
2
3 module bolt(length) {
4     cylinder(h=length, r=2, $fn=200, center=true);
5     translate([0,0,length/2]) cylinder(r=3.5, h=2, $fn=200, center=true);
6 }
7
8 rotate([0,90,0]) union() {
9     nut();
10    bolt(30);
11 }
```



See if you can add a screw cut to the head of the bolt so that it looks like this:

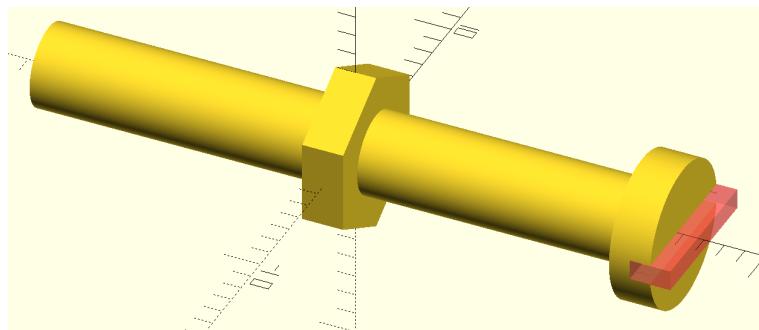


Hint: you need to subtract a cuboid from the larger of the two cylinder above.

C.1.12 The # 'ghost' modifier

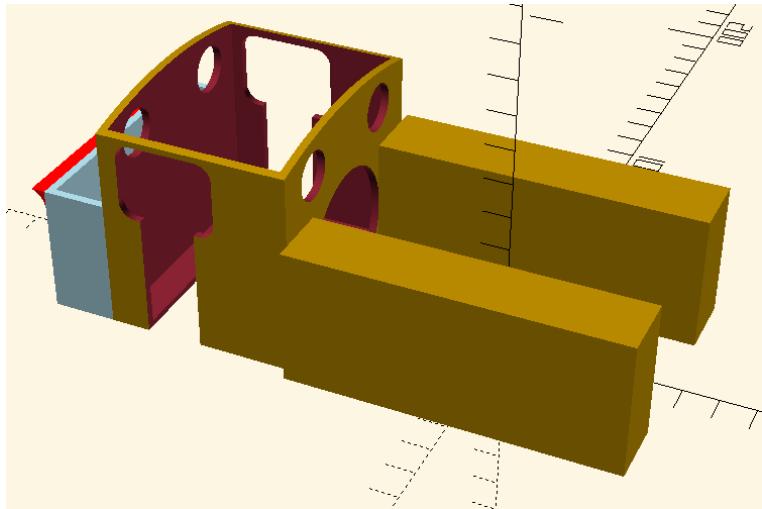
When trying to use the difference and intersection commands, it can be hard to visualise what is going on as some of the elements will be invisible.

If you place a `#` character in front of an invisible object it will be displayed in a ghostly form. So, for instance, when I was adding the screw slot above, I put a `#` character in front of the cuboid I was using to make the slot so that I could see its relative position to the screw head:



C.1.13 Your exercise: the J69

This is the body section from a simple steam engine, suitable for 3D printing:



and here is a painted model:

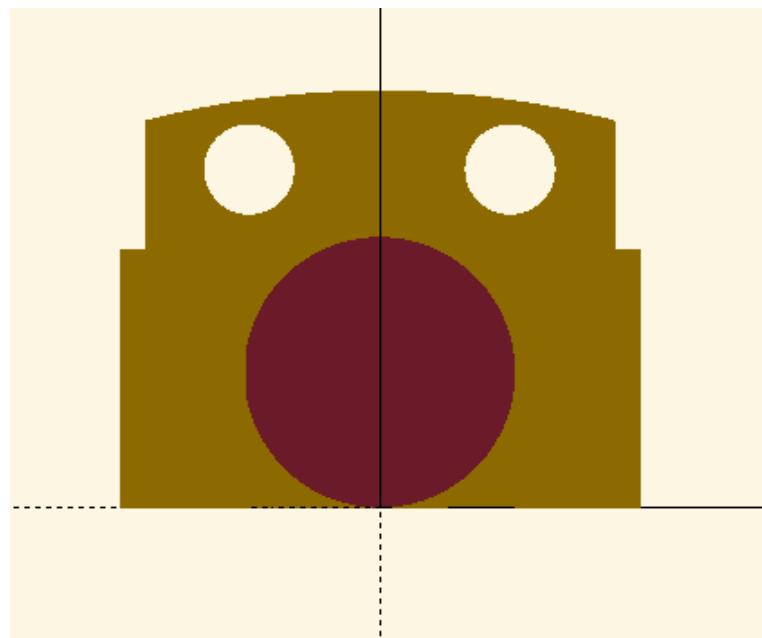


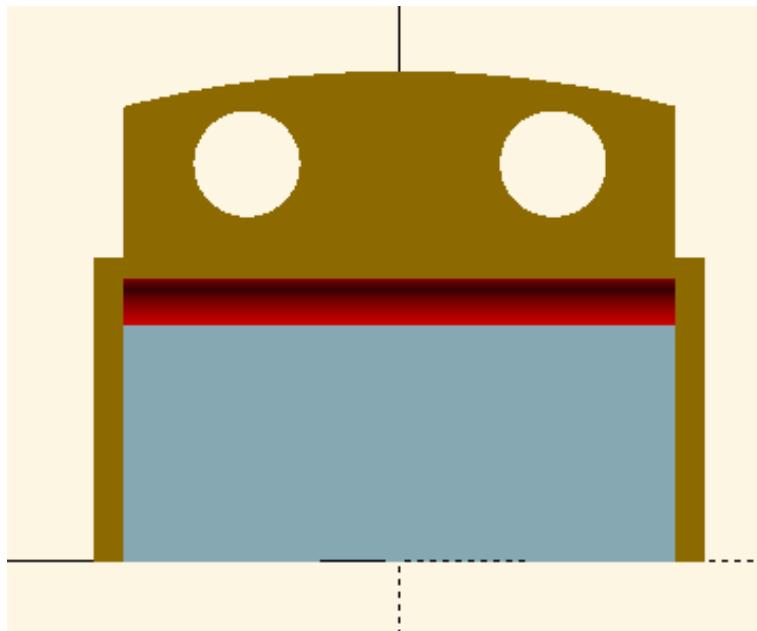
It is based on a prototype, real world engine called the J69 class. You can read more about them at <https://www.lner.info/locos/J/j67j69.php>. There is just one real example left, shown here at Bressingham:



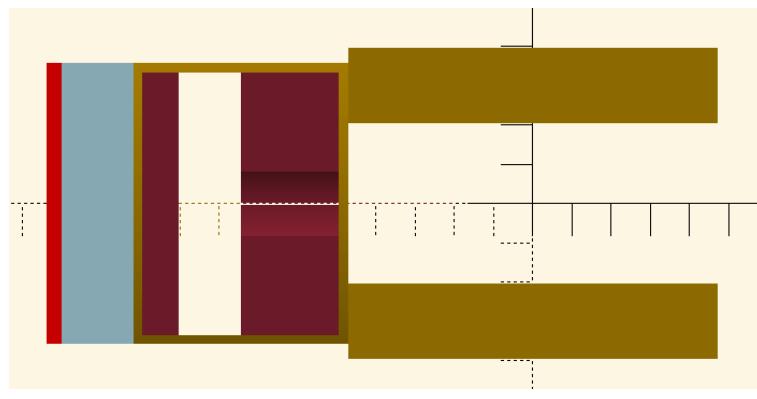
Your task now is to build your own version of the simple steam engine body, and if you are so inclined to embellish it so that it captures more of the prototype. Begin by defining the tanks, which are simple cuboids, and then make the cab as a cuboid which has another cuboid subtracted from it to make a hollow box. You can then subtract cylinders from it to make the round windows.

Here are reference views of the model from various angles.

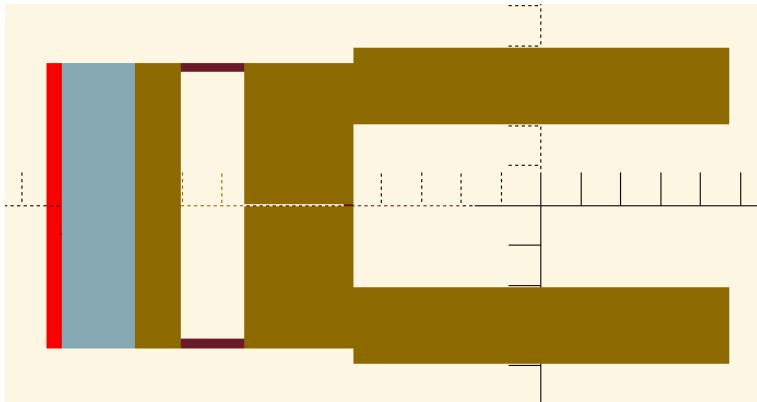




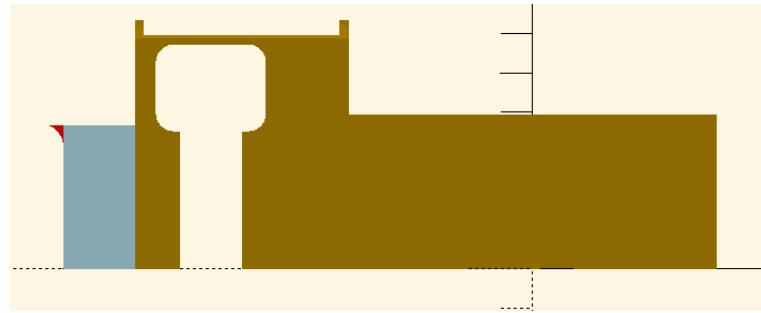
Rear:



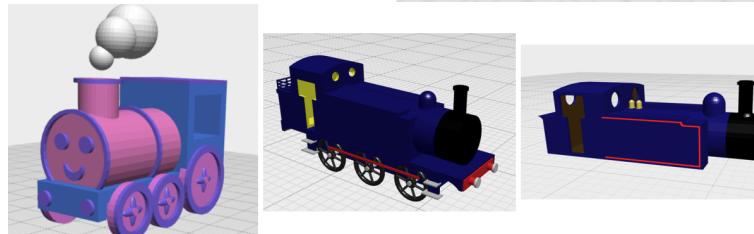
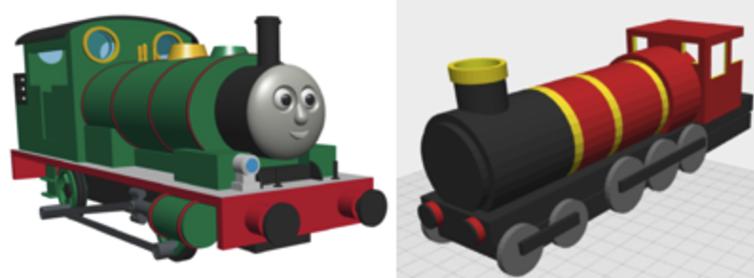
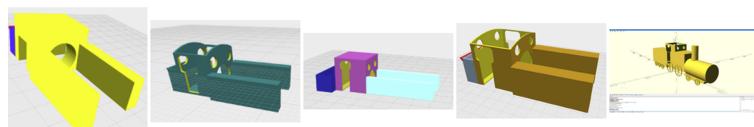
Above:



Below:



When you have something you are pleased with, take a screenshot and then submit your work via Moodle. For inspiration, here are some examples from previous years.



C.2 Simple rewrites

This laboratory session introduces you to the ART rewrite engine through *simple* rewrites—that is, rewrite rules that have no premises (i.e. nothing ‘above the line’). In logic, these kinds of rules are called *axioms* because they exist as independent statements of fact that are not derived from any other facts. In our programming-style view of rewrite rules, these axioms can be viewed as rewrites that *must* be applied whenever the left hand side of their conclusion matches the top of the current term, Θ .

C.2.1 Download and test art.jar

Turn to section A.1 and follow the instructions there. After completion, you will have a new directory (whose name and location is your choice, but which we shall refer to as *artwork*). The directory will be empty except for a copy of *art.jar*.

C.2.2 A quick test

Open a command line window and change directory to *artwork*. At the command line, type

```
java -jar art.jar !print version
```

You should receive a response from ART similar to this:

```
Attached to ARTDefaultPlugin
Interpreter set to eSOS
ART 5_0_385 2025-01-29 06:20:22
```

Of course, the version number, date and time will probably be later than shown here.

If you are running under Unix, Linux or MacOS then your command line shell may intercept ! characters and treat them specially... If you do not get the response above, try typing the command with an escape character before the ! like this:

```
java -jar art.jar \\!print version
```

C.2.3 Make a work file

Use your preferred text editor to make a new working file called **reduction.art**. Eventually this file will contain the full specification of the reduction semantics

for your language although in this lab we are just going to use it as a sandbox for experiments. When you have completed these exercises you *must* submit your final version of `reduction.art` via the Lab 2 Moodle page submission box so that we can track your progress, and give you support if you need it.

Add one line to the start of your new file containing

```
!print version
```

Save your file, and then at the command line type

```
java -jar art.jar reduction.art
```

You should see the same response as before.

```
Attached to ARTDefaultPlugin
Interpreter set to eSOS
ART 5_0_385 2025-01-29 06:20:22
```

In what follows, the idea is for you to copy each example into your reduction.art file, and then run it with the command

`java -jar art.jar reduction.art`

As you start each new example, please just comment out the parts you are no longer working on so that the final reduction.art that you submit shows all of the work that you have done. ART block comments are delimited (*like this*).

C.2.4 A first rewrite

In ART, rewrite rules are written as an optional sequence of conditions, followed by three minus signs `---` followed by a *transition* comprising two terms separated by a *relation* symbol. In this lab we are using axioms which have no conditions, so each rule will just start with `---`. We shall also limit ourselves to the *small step transition* `->`.

Our simplest example, then, looks like this

```
1 --- a->b
2 !try a
```

The `!try` directive initialises the current term Θ to be the term you specify (just `a` in this case), and then passes it to the rewrite engine for processing.

The rewrite engine repeatedly attempts to apply its known rules to Θ until either it reaches a *terminal* or it gets stuck, because it has a Θ which is not terminal but for which no rules match. Each rewrite attempt is called a *step*.

When we run the example above, we get this trace output.

```
*** Warning: in relation -> constructor b has no rule definitions
Step 1
  Rewrite attempt 1: a ->
  -R1 --- a->b
  -R1 rewrites to b
Step 2
  Rewrite attempt 2: b ->
  No rules in relation -> for constructor b
  Stuck on b after 2 steps and 2 rewrite attempts
```

Now this has worked as we can see that the **a** has been rewritten to **b** but something is not right because we are getting a warning message about missing rule definitions, and the rewriter says it is stuck. The reason for this unpleasantness is that the rewriter does not know that it is supposed to stop when it gets to **b**

We fix this by declaring that **b** is a *terminal* for which there are no rules, and at which the rewriter can terminate.

C.2.5 Tidying things up

The **!terminal** directive takes a transition arrow and a comma-delimited list of terminal terms. We need to specify a particular transition because sometimes we want different transition systems to have different terminals.

Change your example by adding the appropriate terminal declaration:

```
1 !terminal -> b
2 --- a->b
3 !try a
```

Try running it, and you should see this trace:

```
Step 1
  Rewrite attempt 1: a ->
  -R1 --- a->b
  -R1 rewrites to b
Normal termination on b after 1 step and 1 rewrite attempt
```

That is much more satisfactory. Once the **b** appears, the rewriter knows that it has reached a terminal and so it stops trying to step and terminates. We also get rid of the warning message about there being no rules for **b**.

C.2.6 Adding test outcomes to !try

We use `!try` directives to test our rules, and like any good testing system we would like to be able to supply an expected outcome for the test, and have the system give us a summary of how many tests have passed and failed. In ART, you can simply add a test term as part of each `!try` like this:

```

1 !terminal -> b
2 --- a->b
3 !try a = b
4 !try b = b
5 !try a = a

```

Here we have three `!try`s which will be performed in order. The third one is deliberately incorrect; we get this trace:

```

Step 1
  Rewrite attempt 1: a ->
    -R1 --- a->b
    -R1 rewrites to b
  Normal termination on b after 1 step and 1 rewrite attempt
  *** Successful test
Step 1
  Rewrite attempt 1: b ->
    Terminal b
  Normal termination on b after 1 step and 1 rewrite attempt
  *** Successful test
Step 1
  Rewrite attempt 1: a ->
    -R1 --- a->b
    -R1 rewrites to b
  Normal termination on b after 1 step and 1 rewrite attempt
  *** Failed test: expected a
  Successful tests: 2; failed tests 1

```

This is how we do unit testing and regression testing in ART.

C.2.7 Changing the trace level

The rewriter tells you what it is doing, which is helpful for debugging but can be a bit overwhelming for long runs. ART has a global *trace level* which controls how verbose these messages are. The default is trace level 3, which tells you every rule that is being tested. Trace level 2 restricts that to just listing the successful rewrite rule. Trace level 1 further restricts the output to just giving the terminating term, and trace level 0 suppresses all of the rewriter messages.

There are higher trace levels too which we shall use in the next lab to look in detail at the processing of conditions.

You set the trace level by writing (for instance)

```
1 | !trace 2
```

This will remain in force until the next `!trace` directive, or until the end of the script.

C.2.8 Comments and console messages

ART scripts can contain comments in two forms: line comments that begin with `//` and continue to the next line break, and block comments delimited by `(*...*)`. So we can write things like

```
1 (* Lab 2 sandbox *)
2 !trace 2
3 !terminal -> b // set b as a terminal under -> transitions
4 ---- a->b
5 !try a = b // first unit test
```

The `!print "a string"` directive allows you to output a message whilst a script is running.

C.2.9 Using variables: swapping

ART rewrite rules may contain *variables* which are bound to subterms as part of the matching process. Variable names all begin with a *single* underscore character.

A simple example is to take an arity-2 term `swap` and rewrite it so that the arguments are swapped over:

```
1 !terminal -> swapped
2 ----swap(_X,_Y) -> swapped(_Y,_X)
3 !try swap(a,b) = swapped(b,a)
```

Step 1

Rewrite attempt 1: `swap(a, b) ->`
`-R1 --- swap(_X, _Y)->swapped(_Y, _X)`
`-R1 rewrites to swapped(b, a)`

Normal termination on `swapped(b, a)` after 1 step and 1 rewrite attempt

```
*** Successful test
Successful tests: 1; failed tests 0
```

If you increase the trace level to 5, then the rewriter will show you the variable bindings as it goes along.

```
Trace level set to 5
Step 1
  Rewrite attempt 1: swap(a, b) ->
    -R1 --- swap(_X, _Y)->swapped(_Y, _X)
    -R1 bindings after Theta match { _X=a, _Y=b }
    -R1 rewrites to swapped(b, a)
Normal termination on swapped(b, a) after 1 step and 1 rewrite attempt
*** Successful test
Successful tests: 1; failed tests 0
```

C.2.10 A logic inverter

We shall now develop some rewrite rules that simulate basic logic gates. The idea is to encode the truth table of the logic function in a set of rewrite rules.

The simplest logic gate is the *inverter* or NOT-gate. It has one input and one output, if the input is true then the output is false, otherwise the output is true.

We shall use T and F to denote true and false respectively. Add these lines to your file

```

1 !terminal -> F,T
2   --- invert(T) -> F
3   --- invert(F) -> T
4
5 !try invert(T) = F
6 !try invert(F) = T
```

You should get the following trace:

```
Step 1
  Rewrite attempt 1: invert(T) ->
    -R1 --- invert(T)->F
    -R1 rewrites to F
Normal termination on F after 1 step and 1 rewrite attempt
*** Successful test
Step 1
  Rewrite attempt 1: invert(F) ->
```

```

-R1 --- invert(T)->F
-R1 Theta match failed: seek another rule
-R2 --- invert(F)->T
-R2 rewrites to T
Normal termination on T after 1 step and 1 rewrite attempt
*** Successful test
Successful tests: 2; failed tests 0

```

This is what we want, but notice how inverting **T** to **F** needs less work than inverting **F** to **T** because we have to search through the rules until we find one that matches.

C.2.11 An AND gate

We can now extend the approach to two-input gates. These rules embody the truth table for the AND function:

```

1 !terminal -> F,T
2 ---and(F,F) -> F
3 ---and(F,T) -> F
4 ---and(T,F) -> F
5 ---and(T,T) -> T
6
7 !try and(T,F) = F
8 !try and(T,T) = T

```

Try these out. Note again how the order of the rules affects the length of the trace.

C.2.12 Rewrite rules are looked at in order

We can quite often use a dirty trick to reduce the amount of searching that the rewrite engine has to do, and to reduce the number of rules that we have to write. ART has a special ‘wildcard’ variable written as **_** (a single underscore) which will match anything. We can rewrite our **and** rules like this:

```

1 !terminal -> F,T
2 ---and(T,T) -> T
3 ---and(_,_) -> F
4 !try and(T,F) = F
5 !try and(T,T) = T

```

How does this work? Well the first rule handle the only row of the truth table that yields T (when both inputs are also T). All of the other truth table rows

yield false. As long as we look for `and(T,T)` first, we can use the wildcards to handle all the other cases.

Of course, if we get the order of the rules wrong, then bad things happen. If we swap the rules

```

1 !terminal -> F,T
2 ---and(_,_) -> F
3 ---and(T,T) -> T
4 !try and(T,F) = F
5 !try and(T,T) = T

```

then we get this trace:

```

Step 1
  Rewrite attempt 1: and(T, F) ->
    -R1 --- and(_, _) -> F
    -R1 rewrites to F
  Normal termination on F after 1 step and 1 rewrite attempt
  *** Successful test
Step 1
  Rewrite attempt 1: and(T, T) ->
    -R1 --- and(_, _) -> F
    -R1 rewrites to F
  Normal termination on F after 1 step and 1 rewrite attempt
  *** Failed test: expected T
  Successful tests: 1; failed tests 1\

```

The rule is that **we must proceed from the particular to the general**.

Now, this technique is useful for improving interpreter efficiency, but exploiting the ordering goes against the declarative styles that theorem provers and code verifiers demand. So if you are working with a formal verification process, you may need to write out rules that are order-independent (and sometimes that is hard).

C.2.13 Your first exercise: OR, NOR, NAND and XOR

Using the style developed in the previous two sections, implement and write tests for the two input functions OR, NOR, NAND and XOR. In case you can't remember the truth tables for these functions, see the Wikipedia page at

https://en.wikipedia.org/wiki/Logic_gate

C.2.14 Addition using ‘beads’

Now that we have basic logic gates, we could in principle extend to full adders and indeed any binary digital hardware system. However we can use an alternative approach to implement arithmetic more directly. You’ll recall that the Roman numbering system does not have a denotation for zero, and that their non-positional notation makes arithmetic directly in Roman numerals difficult. For everyday calculations, the Romans used an abacus: essentially a set of wires strung with beads: by moving beads in and out of the viewing windows we can represent different numbers. We shall use terms and rewrites to model an abacus.

We shall make use of two terminals `bead` and `zero`. `bead` has arity one and `zero` has arity zero. The idea is that the number n is represented by a nest of n `bead` terms wrapped around a `zero` term:

```
0 zero
1 bead(zero)
2 bead(bead(zero))
3 bead(bead(bead(zero)))
```

and so on. We can increment (add one to) a number by adding an extra outer bead, and we can decrement (subtract one) by stripping the outer bead off.

Now, the addition operator takes two numbers represented in this way. We write rules that increment the first operand and decrement the second operand until the second operand is zero. Effectively the second operand counts down to zero, and the first counts up in lockstep.

```
1 !terminal -> bead, zero
2 --- add(_X, zero) -> _X
3 --- add(_X, bead(_Y)) -> add(bead(_X),_Y)
```

Use these `!try` commands like these to exercise these rules, and understand how they work.

```
1 !try add(zero,zero) = zero
2 !try add(bead(bead(bead(bead(zero)))), zero) = bead(bead(bead(bead(zero))))
3 !try add(bead(bead(bead(bead(zero)))), bead(bead(zero))) = bead(bead(bead(bead(bead(bead(zero))))))
```

What can you say about the number of steps that will be executed when adding p to q ? Can you add rules to reduce the number of steps?

C.2.15 Subtraction using beads

For addition, we incremented the left operand whilst decrementing the right operand to zero. For subtraction, we can decrement the left operand whilst counting the right operand down. We do this by stripping the outer bead from both operands.

```

1 --- sub(_X, zero) -> _X
2 --- sub(bead(_X), bead(_Y)) -> sub(_X,_Y)

```

Here are some !trys that exercise subtraction.

```

1 !try sub(bead(bead(bead(bead(zero)))), bead(bead(zero))) = bead(bead(zero))
2 !try sub(bead(bead(zero)), bead(bead(zero))) = zero
3 !try sub(bead(bead(zero)), zero) = bead(bead(zero))

```

Now, there is a problem. Our subtraction operation only works with natural numbers: we have no way of representing and manipulating negative numbers. If we try a subtraction that yields a negative integer, the rules get stuck:

```

1 !try sub(bead(bead(zero)), bead(bead(bead(zero))))

```

gives us this trace:

Step 1

```

Rewrite attempt 1: sub(bead(bead(zero)), bead(bead(bead(zero)))) ->
-R1 --- sub(bead(_X), bead(_Y))->sub(_X, _Y)
-R1 rewrites to sub(bead(zero), bead(bead(bead(zero))))

```

Step 2

```

Rewrite attempt 2: sub(bead(zero), bead(bead(bead(zero)))) ->
-R1 --- sub(bead(_X), bead(_Y))->sub(_X, _Y)
-R1 rewrites to sub(zero, bead(bead(zero)))

```

Step 3

```

Rewrite attempt 3: sub(zero, bead(bead(zero))) ->
-R1 --- sub(bead(_X), bead(_Y))->sub(_X, _Y)
-R1 Theta match failed: seek another rule

```

```

-R2 --- sub(_X, zero)->_X

```

```

-R2 Theta match failed: seek another rule

```

Stuck on sub(zero, bead(bead(zero))) after 3 steps and 3 rewrite attempts

As soon as the left hand operand reaches zero, we get stuck.

There is a kind of arithmetic called *saturation arithmetic* which sets lower and upper bounds (l, u) on numbers, typically zero and some integer such as 255 or 65,536. You can read more at https://en.wikipedia.org/wiki/Saturation_arithmetic

In saturation arithmetic, $l - 1 = l$ and $u + 1 = u$. We can model this by adding bounds rules. Here is subtraction again, but with a rule that sets a lower bound of zero.

```

1 | --- sub(bead(_X), bead(_Y)) -> sub(_X,_Y)
2 | --- sub(_X, zero) -> _X
3 | --- sub(zero, _Y) -> zero

```

You will now get normal termination, but of course any negative results are jammed to zero. This kind of arithmetic is useful in signal processing (both image processing and audio processing) and some specialised signal processing devices such as the Analog Devices ADSP 2105 processor offer machine code instruction for saturated arithmetic.

What rule would you need to put a upper bound of ten on the addition operation?

C.2.16 Your second exercise: arithmetic with negative numbers

As we have seen, the subtraction rules above really only work with the naturals, and once the left hand operand gets down to zero we get stuck. Your second (quite hard) exercise is to find a number representation and some rules that allow subtraction involving negative quantities, i.e. to extend our `sub` operation from the Natural numbers to the Integers.

Here's a hint: represent each Integer n as a pair of Natural numbers (p, q) where $n = p - q$. You may (and actually you may not) find this CS Stack exchange page useful: <https://cs.stackexchange.com/questions/2272>

Try and write rules for both addition and subtraction.

Can you do multiplication and division?

(Health warning: I have not worked out answers for these myself yet, and these puzzles can keep you awake at night... Have a play and learn, but don't be concerned if you can't find a perfect solution.)

C.2.17 The ART value system - a preview

It is pretty clear that doing arithmetic in element rewrites is hard work, both in design terms but perhaps more importantly in terms of the quantity of rewriting that is required: we do not really want to have to represent the number 1,000,000 with a million-and-one nested terms.

In practice, we use these ideas to convince ourselves that rewrite systems *could* handle all of the operations we see on a modern processor, and then we cheat a bit and provide a set of special terms which when substituted into other terms magically perform complicated calculations. In reality these special operations are escape hatches from the rewriting world into conventional computation, and they act both to help structure our rules around well-understood operations, and to massively improve efficiency.

In ART, these special functions all have names starting with two underscores, like `__add` and `__sub`. We also provide a set of types, and some shorthands when writing terms. So, for instance, here is a much easier way to do addition:

```

1 --- add(_X, _Y) -> __add(_X, _Y)
2 !try add(3,4) = 7

```

Note carefully that the `add` constructor is *not the same thing* as the `__add` built-in function. It is easy to get confused.

When you run this example, you should see:

```

Step 1
Rewrite attempt 1: add(3, 4) ->
-R1 --- add(_X, _Y)->__add(_X, _Y)
-R1 rewrites to 7
Normal termination on 7 after 1 step and 1 rewrite attempt
*** Successful test

```

The addition gets done in one magical step, rather than the countdown sequence that we saw earlier.

We shall look at ART's *value system* in much more detail later. For now, you might like to glance at Section A.13 to see a list of these builtin types and operations.

Please do not forget to submit your reduction.art file to Moodle when you have finished, and by the deadline.

C.3 Starting a language

This laboratory session introduces you to the discipline of language design.

The work is in two parts: firstly we shall develop the *internal syntax* of your project language which comprises a set of primitive types, a set of *signatures* and a *configuration*. We will then look at the efficient implementation of arithmetic and logic operations using ART's builtin types and operations (rather than the bead-based rules we looked at in Lab 2) and how to compose them into compound expressions.

You have a blank sheet of paper, so where do you begin?

Often, folk try to replicate the ‘look’ of their favourite language but with some personal tweaks. Of course, studying existing languages can be a fertile source of inspiration but extending such a language can be very challenging, and many projects founder in what Fred Brooks famously called *the tar pit of complexity*. We do not recommend this approach because fitting new capabilities into an existing language is hard, especially if you have no experience of building a language from the ground up.

Instead, we recommend a ‘semantics first’ approach in which the designer begins by simply listing the features of the language, independent of the syntax. We begin by thinking about primitive (atomic) types, the basic arithmetic and logical operations and then turn to control flow, data type constructors and special domain-specific operations.

For each operation, we write a *signature* comprising a name for the operation and then a sequence of operands with optional type constraints.

So, for instance, subtraction over 32-bit integers might have this signature:

`1 sub(_L: _int32, _R: _int32):_int32`

This says that there is an operation called `sub` which takes two operand called `_L` and `_R` (for left and right) each of which must be of type `_int32`, and the operation returns a value which will also be of type `_int32`. The type constraints are just that—things that must be checked—and are optional. If they are omitted then that element of the signature will accept anything.

Similarly, the ‘not equal to’ operation might have this signature:

`1 neq(_L, _R):_bool`

taking two unconstrained operands and returning a boolean.

The arguments to a signature can be code as well as data. For instance we might represent an `if-then-else` feature as

```
1 if(_P:_bool, _S1, _S2):_bool
```

Here `_P` is a *predicate* return a boolean result, and `_S1` and `_S2` are unconstrained terms that might represent simple numbers, expressions or statements. We can use this to represent these three Java-like features:

1. A selection expression `a > b ? 3 : 4`; represented as `if(gt(a,b),3,4)`
2. An `if-then` statement `if (a > b) return 3;` as `if(gt(a,b),3,_done)`
3. An `if-then-else` statement `if (a > b) return 3; else return 4;` represented as `if(gt(a,b),3,4)`

Notice how cases 1 and 3 have the same internal syntax but different external syntax. (Actually we have cheated a little as the Java `return` may need some special handling.)

Once we have a plausible set of signatures, we start to write the rules that will ‘explain’ their semantics using rewrite rules. Often one finds that a few new ‘helper’ signatures will be needed, especially for features such as a `switch` (or `case`) statement that have a sequence of sub-operations.

We can test individual features and features in combination by trying different test terms. Once we have a complete and consistent set of rules that process our test cases to give the expected results (and do not get stuck!) then we can write some context free rules to create an external syntax parser which outputs expressions built from our semantic signatures.

Like all software engineering processes, we end up iterating the design as issues arise during development and testing, but this separation of concerns between the semantics and the design of syntax (the ‘look’) of the language avoids a problem that often arises otherwise: syntactic constraints which create special cases in the semantics that complicate the rules. We seek uniformity and orthogonality in the semantic rules; then we use the syntax and static typing rules to outlaw dubious combinations.

As an example, in the original C language the array indexing operation is rather general; one can write something like `a[3]` (which is conventional) but also `101[3]`, that is the base of the array may be an explicit number. This eases the design of a translator, because the indexing operation may have a signature like `index(_base:_int32, _offset_int32)` where any expression that yields a 32-bit integer may appear as the base. Now, in some contexts this is a feature; in others a bug. Allowing a numeric base means that we can directly address any specific location in memory and are not just limited to areas that have been allocated by the translator and given an identifying name. For some embedded applications, accessing specific memory locations is a necessary feature. On the other hand, unfettered access to memory on multi-process systems allows

malignant code access to information which should be secured, and that is at the root of much malware. Therefore we would expect a compiler for such a system to use syntax and static semantic checks to filter out such idioms before the dynamic semantic rules are activated.

Similarly, in the selection examples above we would expect the parser to enforce the requirement that selection *expressions* must always have both a THEN and an ELSE clause so that the expression always returns some value, whereas a selection *statement* is valid without the else clause, as that simply means ‘do-nothing’.

C.3.1 Exercise: primitive type selection

Review the type structure of the Java language by reading sections 4.1 and 4.2 of the Java Language Specification at <https://docs.oracle.com/javase/specs/jls/se23/html/jls-4.html#jls-4.1>.

Pocket calculators do not distinguish between floating point and integer numbers, but Java and other languages provide a plethora of different integer and floating point data types. Why?

Review the Java `BigDecimal` and `BigInteger` classes, documented at

<https://docs.oracle.com/javase/8/docs/api/java/math/BigDecimal.html>
and

<https://docs.oracle.com/javase/8/docs/api/java/math/BigInteger.html>

How do they work, and what are their advantages and disadvantages?

ANSI C-89 does not have a Boolean type. How are the results of predicates represented?

Examine the Java GCD program on page 15 and the ‘compact’ version that follows. What data types are in use?

For the three domains of music, image processing and 3D CAD, what data types do we need, and which might be desirable?

What will you use in your language, and why?

C.3.2 Exercise: enumerate the operations and types in the GCD language

Refer back to the GCD program on page 15. Do not think about syntax at the moment, but just try to list the primitive types and operations in use in that example, and write a signature for each of them. Do not overlook control flow, assignment to variables and dereferencing of variables.

C.3.3 Deciding on the configuration

We model programming languages as transitions between states of a *configuration* which contains at the very least a program term. For anything other than a truly pure functional language (and these are rare beasts) the configuration will also contain some other *semantic entities* whose function is to capture side effects such as assignment and output as we rewrite the program term.

Once you have a clear idea of the operations that your language will provide, you will be able to decide which semantic entities they need, and hence what your language's configuration will be. We represent configurations as *tuples* of semantic entities, written, for instance, as $\langle \theta, \sigma, \alpha \rangle$. Every language needs at least a program term: the configuration for a pure functional language would be $\langle \theta \rangle$. We shall now consider other kinds of semantic entity.

The store and environment - σ and ρ

A language with assignment needs at least a store.

In more detail, a language where all variables are global can only have one binding to each variable name. In languages with multiple scope regions, the same variable name can refer to different variables (and thus different bindings) in each scope region. In either case we need a store (usually called Sigma, or σ) which represents main memory. If we also want multiple scope regions, then we need a *symbol table* or *environment* (usually called Rho or ρ) to be associated with each scope region.

When just using σ we implement it as a map from identifiers to values. When using environments, we have single map σ from location identifiers (often natural numbers, as they are in the hardware) to values, and then each scope region has environment which is a map from identifiers to location identifiers. This two-level representation allows the same identifier to appear in different scopes bound to different locations in the store, and indeed for multiple identifiers to bind to the same store location, as is required for reference based languages such as Java.

If your language is to have a procedure-like construct (e.g. methods in Java, functions in C and modules in OpenSCAD) then multiple scopes are almost mandatory otherwise you cannot model local variables.

Output and input - α and β

Many languages need output α and input β . Whilst prototyping our language, we usually just represent these semantic entities as linked lists. In reality, input and output are rather fraught aspects of programming as illustrated by the scale of the [java.io](#) and [java.nio](#) packages: see

<https://docs.oracle.com/javase/8/docs/api/java/io/package-summary.html> and
<https://docs.oracle.com/javase/8/docs/api/java/nio/package-summary.html>

And then, of course, there is input and output associated with Graphical User Interface systems such as Java Swing and Java FX. All of these libraries are complex to the point of being overwhelming, but fortunately we do not need to consider them as they are all written *in* the host programming language and are not features *of* the language. However there are languages such as Pascal and BASIC that do have simple I/O built into them. Some languages, such as the OpenSCAD language from Lab 1 do not really have input at all, and only very limited output facilities.

Exceptions and signals - ν

An *exception* breaks the normal flow of control. Typically, we *raise* an exception which then causes nested functions to return until we encounter an *exception handler* that is given control. Operating system *signals* and hardware *interrupts* are similar events which are generated outside of the running user program.

It is perfectly possible to model exceptions in our rewrite rules, but when starting out it is probably best to delay thinking about exceptions until the core of the language is stable, so we shall not discuss exceptions in this lab.

Domain specific entities

It *may* (and may not) be useful to have semantic entities representing state for the domain specific entities that your programming language is manipulating. A common approach is to keep all such entities on the Java side of the system, and use anonymous numeric handles to refer to them on the rewrite side. An alternative is to use the builtin type `_blob` to hold references to arbitrary binary objects. We shall consider this more fully when we look at the ART plugin conventions in a later lab.

C.3.4 Exercise: decide on your configuration

For a general purpose language that will act as the base for your project, what configuration should you use?

At this point in the design process, you have decided on the ***internal syntax*** of your language, the notation that defines the state space and operation repertoire of your language.

You may need to revisit and extend your design, but now we are ready to begin implementing a prototype rewrite-based interpreter.

C.3.5 Values and their types in ART

ART's rewriter is very simple minded: it looks for matches between closed terms and patterns which may contain variables, and creates an environment of bindings for those variables that may be substituted into another pattern to make a new closed term. It is the very simplicity of this matching operation that makes it a suitable basis for defining language semantics since it acts a sort of lowest-common-denominator description of computation.

The matching process only consider the shape of the terms (trees) and the labels on the nodes which are treated as simple text strings without interpretation. So the label `123` is just a string of characters. At the level of the programming language we are modelling we might want that label to represent an integer, or a real number or the a programming language string. How do we differentiate these cases?

The idea is that a value should be a term in the form `type(payload)`. The `type` label can be anything you like, but ART has built-in support for a set of type names that are recognised during substitution, all of which start with two underscores. So, for instance, a 32-bit integer 123 is represented as `_int32(123)`, a Java double as `_real64(123)` and a string as `_string(123)`. You can also have `_intAP(123)` which represents an arbitrary precision integer (like a BigInteger in Java) and `_realAP(123)` which represents an arbitrary precision real number.

Shorthands in the ART front end

It becomes quite tedious to have to write `_int32(123)` and `_real64(123.0)` so the ART scripting front end *automatically* converts certain forms to the built in type equivalents.

You can see this process in action by asking ART print out the raw form of terms:

```

1 !try 123.0
2 !print derivation
3 !prinraw derivation
4
5 !try 321

```

```

6 | !print derivation
7 | !printraw derivation

```

In detail the `!try` directive loads the current derivation term with its argument. The `!print` directive prints out the ‘cooked’ (abbreviated) version of this term; the `!printraw` directive shows explicitly what is going on under the hood.

```

1 current derivation term: [1383]
2 123.0
3 current raw derivation term: [1383]
4 trTopTuple(_real64(123.0))
5 current derivation term: [1392]
6 321
7 current raw derivation term: [1392]
8 trTopTuple(_int32(321))

```

The `trTopTuple` constructor is used to tie together the elements of a configuration tuple.

C.3.6 Implementing arithmetic operations with the ART Value system

The bead based rules from Lab 2 are an interesting parlour game, and show that pure rewriting can perform arithmetic, but they are not at all practical for our present purpose since the size of the term representing a number is proportional to that number, and for anything other than very simple examples, vast amounts of storage and rewriting are required.

The ART builtin types are accompanied by a fixed set of builtin operations, the names of which also start with two underscores. When the rewriter is making a substitution, if it sees a builtin operation it converts its arguments to their equivalent Java partner types, then performs the operation as a Java expression, and then converts the result it back to a term.

The complete set of builtin types and operations is summarised in the table on page 123. The blue entries are disallowed combinations: for instance it makes no sense to divide strings by one another. Allowed operations have a grey background, and the label tells you the expected type of the arguments.

We use these functions by creating our own signatures for operations such as subtraction which then rewrite to the builtin operation. As a side effect of substituting the bindings, the builtin operation is evaluated;

```

1 ---- subtract(_L, _R) -> _sub(_L, _R)
2 !try subtract(12,7)

```

```

1 Step 1
2 Rewrite attempt 1: subtract(12, 7) ->
3   -R1 --- subtract(_L, _R)->_sub(_L, _R)
4   -R1 rewrites to 5
5 Normal termination on 5 after 1 step and 1 rewrite attempt

```

Now, the example above is *fragile* in that there is no check that the arguments to `_sub` are valid, and that can cause ART to stop with a fatal error:

```

1 --- subtract(_L, _R) -> _sub(_L, _R)
2 !try subtract(12,7.0)

```

yields

```

1 Step 1
2 Rewrite attempt 1: subtract(12, 7.0) ->
3   -R1 --- subtract(_L, _R)->_sub(_L, _R)
4 *** Fatal: Term 7.0 failed type check type against _int32

```

Note that this is not the same thing as getting stuck: the rewriter has had to completely stop in an uncontrolled way because it is not able to evaluate the `_sub` operation. The problem is that we are trying to subtract the real number `7.0` from the integer `12`. The builtin operations require the arguments to be of the same type, so you can add and subtract pairs of reals or integers, but cannot do mixed mode arithmetic.

How can we avoid the fatal error? We add conditions above the line. Recall that the match operation is represented by the `|>` symbol. In ART's script language we write that as `|>`, and use it to check the type of both of the operands:

```

1 !trace 5
2 _L |> _int32(_) _R |> _int32(_)
3 ---
4 subtract(_L, _R) -> _sub(_L, _R)
5
6 !try subtract(12,7.0)
7 !try subtract(12,7)

```

The way to interpret this rule is to start at bottom left, with the left hand side of the conclusion. That must match the current term for this rule to be 'triggered'. As part of the match, `_L` and `_R` will be bound to terms. We then check the conditions above the line: both `_L` and `_R` must match `_int32(_)` that is the root must be `_int32` and there must be exactly one child. If the matches pass, then the rewrite occurs. If not, then the rule will be discarded and the rewriter will look for another. If it can't find one then it will announce that

it is stuck; which is what we want here because we are trying to avoid that irrecoverable fatal type error.

```

1 *** !try(subtract(12, 7.0))
2 Step 1
3 Rewrite attempt 1: subtract(12, 7.0) ->
4 -R1 _L |> _R |> _ --- subtract(_L, _R)->_sub(_L, _R)
5 -R1 bindings after Theta match { _L=12, _R=7.0 }
6 -R1 premise 1 _L |> _
7 -R1 bindings after premise 1 { _L=12, _R=7.0 }
8 -R1 premise 2 _R |> _
9 -R1 premise 2 failed: seek another rule
10 Stuck on subtract(12, 7.0) after 1 step and 1 rewrite attempt
11 *** !try(subtract(12, 7))
12 Step 1
13 Rewrite attempt 1: subtract(12, 7) ->
14 -R1 _L |> _R |> _ --- subtract(_L, _R)->_sub(_L, _R)
15 -R1 bindings after Theta match { _L=12, _R=7 }
16 -R1 premise 1 _L |> _
17 -R1 bindings after premise 1 { _L=12, _R=7 }
18 -R1 premise 2 _R |> _
19 -R1 bindings after premise 2 { _L=12, _R=7 }
20 -R1 rewrites to 5
21 Normal termination on 5 after 1 step and 1 rewrite attempt

```

Note how the mixed mode `!try subtract(12,7.0)` sticks, but the well-formed `!try subtract(12,7)` passes both type checks and rewrites successfully.

We raised the `!trace` level to 5 so that you can see the bindings and the individual premises being evaluated. This is very whilst during debugging.

C.3.7 Nested expressions

We now have an efficient way of getting the rewriter to perform *single* arithmetic operations, but of course we really want to be able to evaluate nested terms: we might represent the Java expression `5−2−1` as `sub(sub(5,2),1)`. The left hand operand of the outer `sub` is not an integer, so the rule above will not pass its type checks:

```

1 !trace 5
2 _L |> _int32(_) _R |> _int32(_)
3 ---
4 sub(_L, _R) -> _sub(_L, _R)
5
6 !try sub(sub(5,2),1)

```

yields

```

1 Step 1
2 Rewrite attempt 1: sub(sub(5, 2), 1) ->
3   -R1 _L |> _R |> _ --- sub(_L, _R)->_sub(_L, _R)
4   -R1 bindings after Theta match { _L=sub(5, 2), _R=1 }
5   -R1 premise 1 _L |> _
6   -R1 premise 1 failed: seek another rule
7 Stuck on sub(sub(5, 2), 1) after 1 step and 1 rewrite attempt

```

(Notice by the way that we changed the constructor from `subtract` to `sub` - it is just a text string and we can use whatever we like as long as the `!try` argument has a matching rule.)

The trick is to somehow get the rewriter to process the `_L` argument separately so as to reduce it to an integer. We achieve this by making a second rule that has a *transition* above the line. This has the effect of recursively calling the rewriter on the argument until it is reduced to a value. We call this kind of rule a *resolution rule*. In this specification, we have named the first rule `subBase` and the second rule `subResolve`. These rule names make it easier to read the trace output, which as you can see gets quite long.

```

1 !trace 5
2 - subBase _L |> _int32(_) _R |> _int32(_)
3 ---
4 sub(_L, _R) -> _sub(_L, _R)
5
6 -subResolve _L -> _LP
7 ---
8 sub(_L, _R) -> sub(_LP, _R)
9
10 !try sub(sub(5,2),1)

```

The order of the rules is important. We need the rewriter to first check the base rule, and only if that fails can we go on to try and resolve sub-expressions.

```

1 Step 1
2 Rewrite attempt 1: sub(sub(5, 2), 1) ->
3   -subBase _L |> _R |> _ --- sub(_L, _R)->_sub(_L, _R)
4   -subBase bindings after Theta match { _L=sub(5, 2), _R=1 }
5   -subBase premise 1 _L |> _
6   -subBase premise 1 failed: seek another rule
7   -subResolve _L->_LP --- sub(_L, _R)->sub(_LP, _R)
8   -subResolve bindings after Theta match { _L=sub(5, 2), _R=1 }
9   -subResolve premise 1 _L->_LP
10  Rewrite attempt 2: sub(5, 2) ->

```

```

11  --subBase _L |> _ _R |> _ ---- sub(_L, _R)-->_sub(_L, _R)
12  --subBase bindings after Theta match { _L=5, _R=2 }
13  --subBase premise 1 _L |> _
14  --subBase bindings after premise 1 { _L=5, _R=2 }
15  --subBase premise 2 _R |> _
16  --subBase bindings after premise 2 { _L=5, _R=2 }
17  --subBase rewrites to 3
18  --subResolve bindings after premise 1 { _L=sub(5, 2), _R=3, _3=1 }
19  --subResolve rewrites to sub(3, 1)

20 Step 2
21 Rewrite attempt 3: sub(3, 1) ->
22  --subBase _L |> _ _R |> _ ---- sub(_L, _R)-->_sub(_L, _R)
23  --subBase bindings after Theta match { _L=3, _R=1 }
24  --subBase premise 1 _L |> _
25  --subBase bindings after premise 1 { _L=3, _R=1 }
26  --subBase premise 2 _R |> _
27  --subBase bindings after premise 2 { _L=3, _R=1 }
28  --subBase rewrites to 2
29 Normal termination on 2 after 2 steps and 3 rewrite attempts

```

What happens if the order of the rules is reversed, and why?

The beauty of this approach is that the resolution will recursively compute any depth of nested expression. Here is a trace of `!try sub(sub(sub(5,2),1),1)` with `!trace 3` so as to reduce the volume of output, whilst still showing the individual rule activations.

```

1 Trace level set to 3
2 Step 1
3 Rewrite attempt 1: sub(sub(sub(5, 2), 1), 1) ->
4  --subBase _L |> _ _R |> _ ---- sub(_L, _R)-->_sub(_L, _R)
5  --subResolve _L->_LP --- sub(_L, _R)-->sub(_LP, _R)
6    Rewrite attempt 2: sub(sub(5, 2), 1) ->
7      --subBase _L |> _ _R |> _ ---- sub(_L, _R)-->_sub(_L, _R)
8      --subResolve _L->_LP --- sub(_L, _R)-->sub(_LP, _R)
9        Rewrite attempt 3: sub(5, 2) ->
10          --subBase _L |> _ _R |> _ ---- sub(_L, _R)-->_sub(_L, _R)
11          --subBase rewrites to 3
12          --subResolve rewrites to sub(3, 1)
13          --subResolve rewrites to sub(sub(3, 1), 1)

14 Step 2
15 Rewrite attempt 4: sub(sub(3, 1), 1) ->
16  --subBase _L |> _ _R |> _ ---- sub(_L, _R)-->_sub(_L, _R)
17  --subResolve _L->_LP --- sub(_L, _R)-->sub(_LP, _R)
18    Rewrite attempt 5: sub(3, 1) ->
19      --subBase _L |> _ _R |> _ ---- sub(_L, _R)-->_sub(_L, _R)
20      --subBase rewrites to 2

```

```

21  —subResolve rewrites to sub(2, 1)
22 Step 3
23  Rewrite attempt 6: sub(2, 1) ->
24  —subBase _L |> _R |> _ --- sub(_L, _R)->_sub(_L, _R)
25  —subBase rewrites to 1
26 Normal termination on 1 after 3 steps and 6 rewrite attempts

```

Now, we have not quite yet finished because our resolution rule only resolves left operands. We need to add a third rule which handles the right operand, and we need to ensure that the left operand will always be evaluated first.

```

1 !trace 3
2 — subBase _L |> _int32(_) _R |> _int32(_)
3 ---
4 sub(_L, _R) -> _sub(_L, _R)
5
6 —subRight _L |> _int32(_) _R -> _RP
7 ---
8 sub(_L, _R) -> sub(_L, _RP)
9
10 —subLeft _L -> _LP
11 ---
12 sub(_L, _R) -> sub(_LP, _R)
13
14 !try sub(sub(7,1),sub(4,2))

```

When processing the term `sub(sub(7,1),sub(4,2))`, at the first step the left operand will be rewritten to an integer, then the right operand and finally the two resolved subexpressions will be handled by the base rule.

```

1 Step 1
2  Rewrite attempt 1: sub(sub(7, 1), sub(4, 2)) ->
3  —subBase _L |> _R |> _ --- sub(_L, _R)->_sub(_L, _R)
4  —subRight _L |> _R->_RP --- sub(_L, _R)->sub(_L, _RP)
5  —subLeft _L->_LP --- sub(_L, _R)->sub(_LP, _R)
6  Rewrite attempt 2: sub(7, 1) ->
7  —subBase _L |> _R |> _ --- sub(_L, _R)->_sub(_L, _R)
8  —subBase rewrites to 6
9  —subLeft rewrites to sub(6, sub(4, 2))
10 Step 2
11  Rewrite attempt 3: sub(6, sub(4, 2)) ->
12  —subBase _L |> _R |> _ --- sub(_L, _R)->_sub(_L, _R)
13  —subRight _L |> _R->_RP --- sub(_L, _R)->sub(_L, _RP)
14  Rewrite attempt 4: sub(4, 2) ->
15  —subBase _L |> _R |> _ --- sub(_L, _R)->_sub(_L, _R)
16  —subBase rewrites to 2

```

```

17  --subRight rewrites to sub(6, 2)
18 Step 3
19  Rewrite attempt 5: sub(6, 2) ->
20  --subBase _L |> _R |> _ --- sub(_L, _R)->_sub(_L, _R)
21  --subBase rewrites to 4
22 Normal termination on 4 after 3 steps and 5 rewrite attempts

```

Why are we so concerned about ordering? Well, some languages allow side effects in expressions and we need them to appear in some well-defined order (usually left to right). Examples of operators that have side effects include `++` and `--` in C and Java, and the stream output operator `>>` in C++.

C.3.8 Your first project task - fill in the operator gaps

The triad of rules (`xxxBase`, `xxxRight` and `xxxLeft` in that order) is the standard way of implementing arity operations in ART. Referring back to the Value table on page 123 you will find six predicate operators (such as `_ne` and `_gt`), six arithmetic operators (`_add` to `_exp`) and a variety of logic and shift operations. For each operation/type combination that you want in your language, you will need to create this triad of rules. That is a lot of rules - nearly 60 in fact if you want to use all of the basic operations. However each triad has the same pattern, and if you make your own constructor name the same as the builtin name without the underscores, then a little use of search and replace will enable you to build the rule set quickly.

For this week's submission, save your old `reduction.art` and make a new empty file. Then add in the necessary rules to support basic arithmetic and comparisons.

Using your rules, show a reduction trace for a term corresponding to

$$(7 - 3) * 2 > 3 + 4 + 5$$

Please do not forget to submit this week's `reduction.art` file to Moodle when you have finished, and by the deadline.

C.4 Making a full language, and project demos

This laboratory session is an opportunity for you to exercise the examples that you saw in Chapter 5. You will then hybridise the two examples to create a version of the GCD language that also supports an output list. Finally, we shall look at the use of a *plugin* to give our languages access to general Java code.

C.4.1 Go through chapter 5 examples

Begin by looking through the contents of the ART examples directory at

<https://github.com/AJohnstone2007/ART/tree/main/doc/examples>

and fetch a copy of `slewaChapter5Examples.art` from <https://github.com/AJohnstone2007/ART/blob/main/doc/examples/slewaChapter5Examples.art> which has many of the rules from Chapter 5 already setup for you.

Your task now is simply to *play* with the rules. Run different `!trys` with the rules, and change the rules around. Find out what happens if you change the order of the rules in a triad and try and understand what is happening. Work methodically through the chapter, learning by doing.

C.4.2 GCD with output

Now take a copy of `gcdReduction.art` from <https://github.com/AJohnstone2007/ART/blob/main/doc/examples/gcdReduction.art> and rename it `reduction.art`; this is the file that you will submit for this week's lab work.

Your task is to extend this example so that it has an output command, and change the test program so that it writes the final GCD result on to the output list.

You will need to change the configuration, and that means every transition will have to be changed... Search and replace is your friend.

C.4.3 Project setup and plugins

We shall demonstrate how to connect ART to your own Java code. The first example is in

<https://github.com/AJohnstone2007/ART/tree/main/doc/project/music>

Please do not forget to submit this week's reduction.art file with your extended GCD language to Moodle when you have finished, and by the deadline.

C.5 Context Free Grammar Rules and GIFT rewrites

This laboratory session introduces you to the

Please do not forget to submit this week's reduction.art file to Moodle when you have finished, and by the deadline.

C.6 SOBRD parsing and attribute evaluation

This laboratory session introduces you to the

Please do not forget to submit this week's reduction.art file to Moodle when you have finished, and by the deadline.

C.7 Control flow with delayed attributes

This laboratory session introduces you to the

Please do not forget to submit this week's reduction.art file to Moodle when you have finished, and by the deadline.

Glossary

abstraction the hiding of unnecessary detail, page 7

actions , page 103

Algebraic Data Types , page 8

assignment model , page 17

attribute grammar , page 12

Attribute Grammar (AG) , page 103

Attribute-Action Grammar (AAG) , page 103

attributes , page 103

back end , page 2

binding *s*, page 31

bindings , page 17

closed , page 30

compiler , page 2

confluent , page 29

Context Free Grammar , page 8

context-free terminal , page 34

control flow , page 16

derivation , page 2

disambiguation rule , page 10

dynamic property , page 26

Executable semantics , page 22

external syntax ., page 2

formal semantics , page 6
front end , page 2
grammar , page 2
internal syntax , page 2
interpreter , page 2
lexemes , page 2
match context , page 31
middle end , page 2
multilexer , page 11
multiparser , page 11
nonterminal , page 34
normal forms ., page 20
open expression , page 30
pattern , page 30
phases , page 1
Program Counter , page 16
redex , page 20
reduction step , page 18
reduction trace , page 18
refactoring , page 26
relation symbol , page 36
rewrite rule , page 26
rewrite schema , page 26
rewriter , page 27
rule schema , page 28
semantic entities ., page 20
semantics the meaning of a language, page 5
single-pass compiler , page 6
SOS , page 12

state the set of values maintained by a program, page 5
static property , page 26
store , page 16
substitute , page 31
substitution , page 30
substitution model , page 17
symbol table , page 2
syntax the written form of a language, page 5
term rewriting , page 12
transition , page 36

Bibliography

- [BG05] Joshua Bloch and Neal Gafter. *Java Puzzlers*. Addison Wesley, 2005.
- [DDH72] O. J. Dahl, E. W. Dijkstra, and C. A. R. Hoare, editors. *Structured programming*. Academic Press Ltd., GBR, 1972.
- [KD14] Donald. E. Knuth and Edgar G. Daylight. *Algorithmic Barriers Falling*. Lonely Scholar, 2014.
- [Knu68] Donald E. Knuth. Semantics of Context-Free Languages. *Mathematical Systems Theory*, 2(2):127–145, 1968.
- [Knu71] Donald E. Knuth. Semantics of Context-Free Languages: Correction. *Mathematical Systems Theory*, 5(2):95–96, 1971.
- [Knu90] Donald E. Knuth. The genesis of attribute grammars. In P. De-
ransart and M. Jourdan, editors, *Attribute Grammars and their Applications*, pages 1–12, Berlin, Heidelberg, 1990. Springer Berlin
Heidelberg.
- [MHMT97] Robin Milner, Robert Harper, David MacQueen, and Mads Tofte. *The Definition of Standard ML (Revised Edition)*. The MIT Press, 1997.
- [Plo04] Gordon D. Plotkin. A structural approach to operational semantics. *J. Log. Algebraic Methods Program.*, 60-61:17–139, 2004.
- [SJW23] Elizabeth Scott, Adrian Johnstone, and Robert Walsh. Multiple input parsing and lexical analysis. *ACM Trans. Program. Lang. Syst.*, 45(3), July 2023.